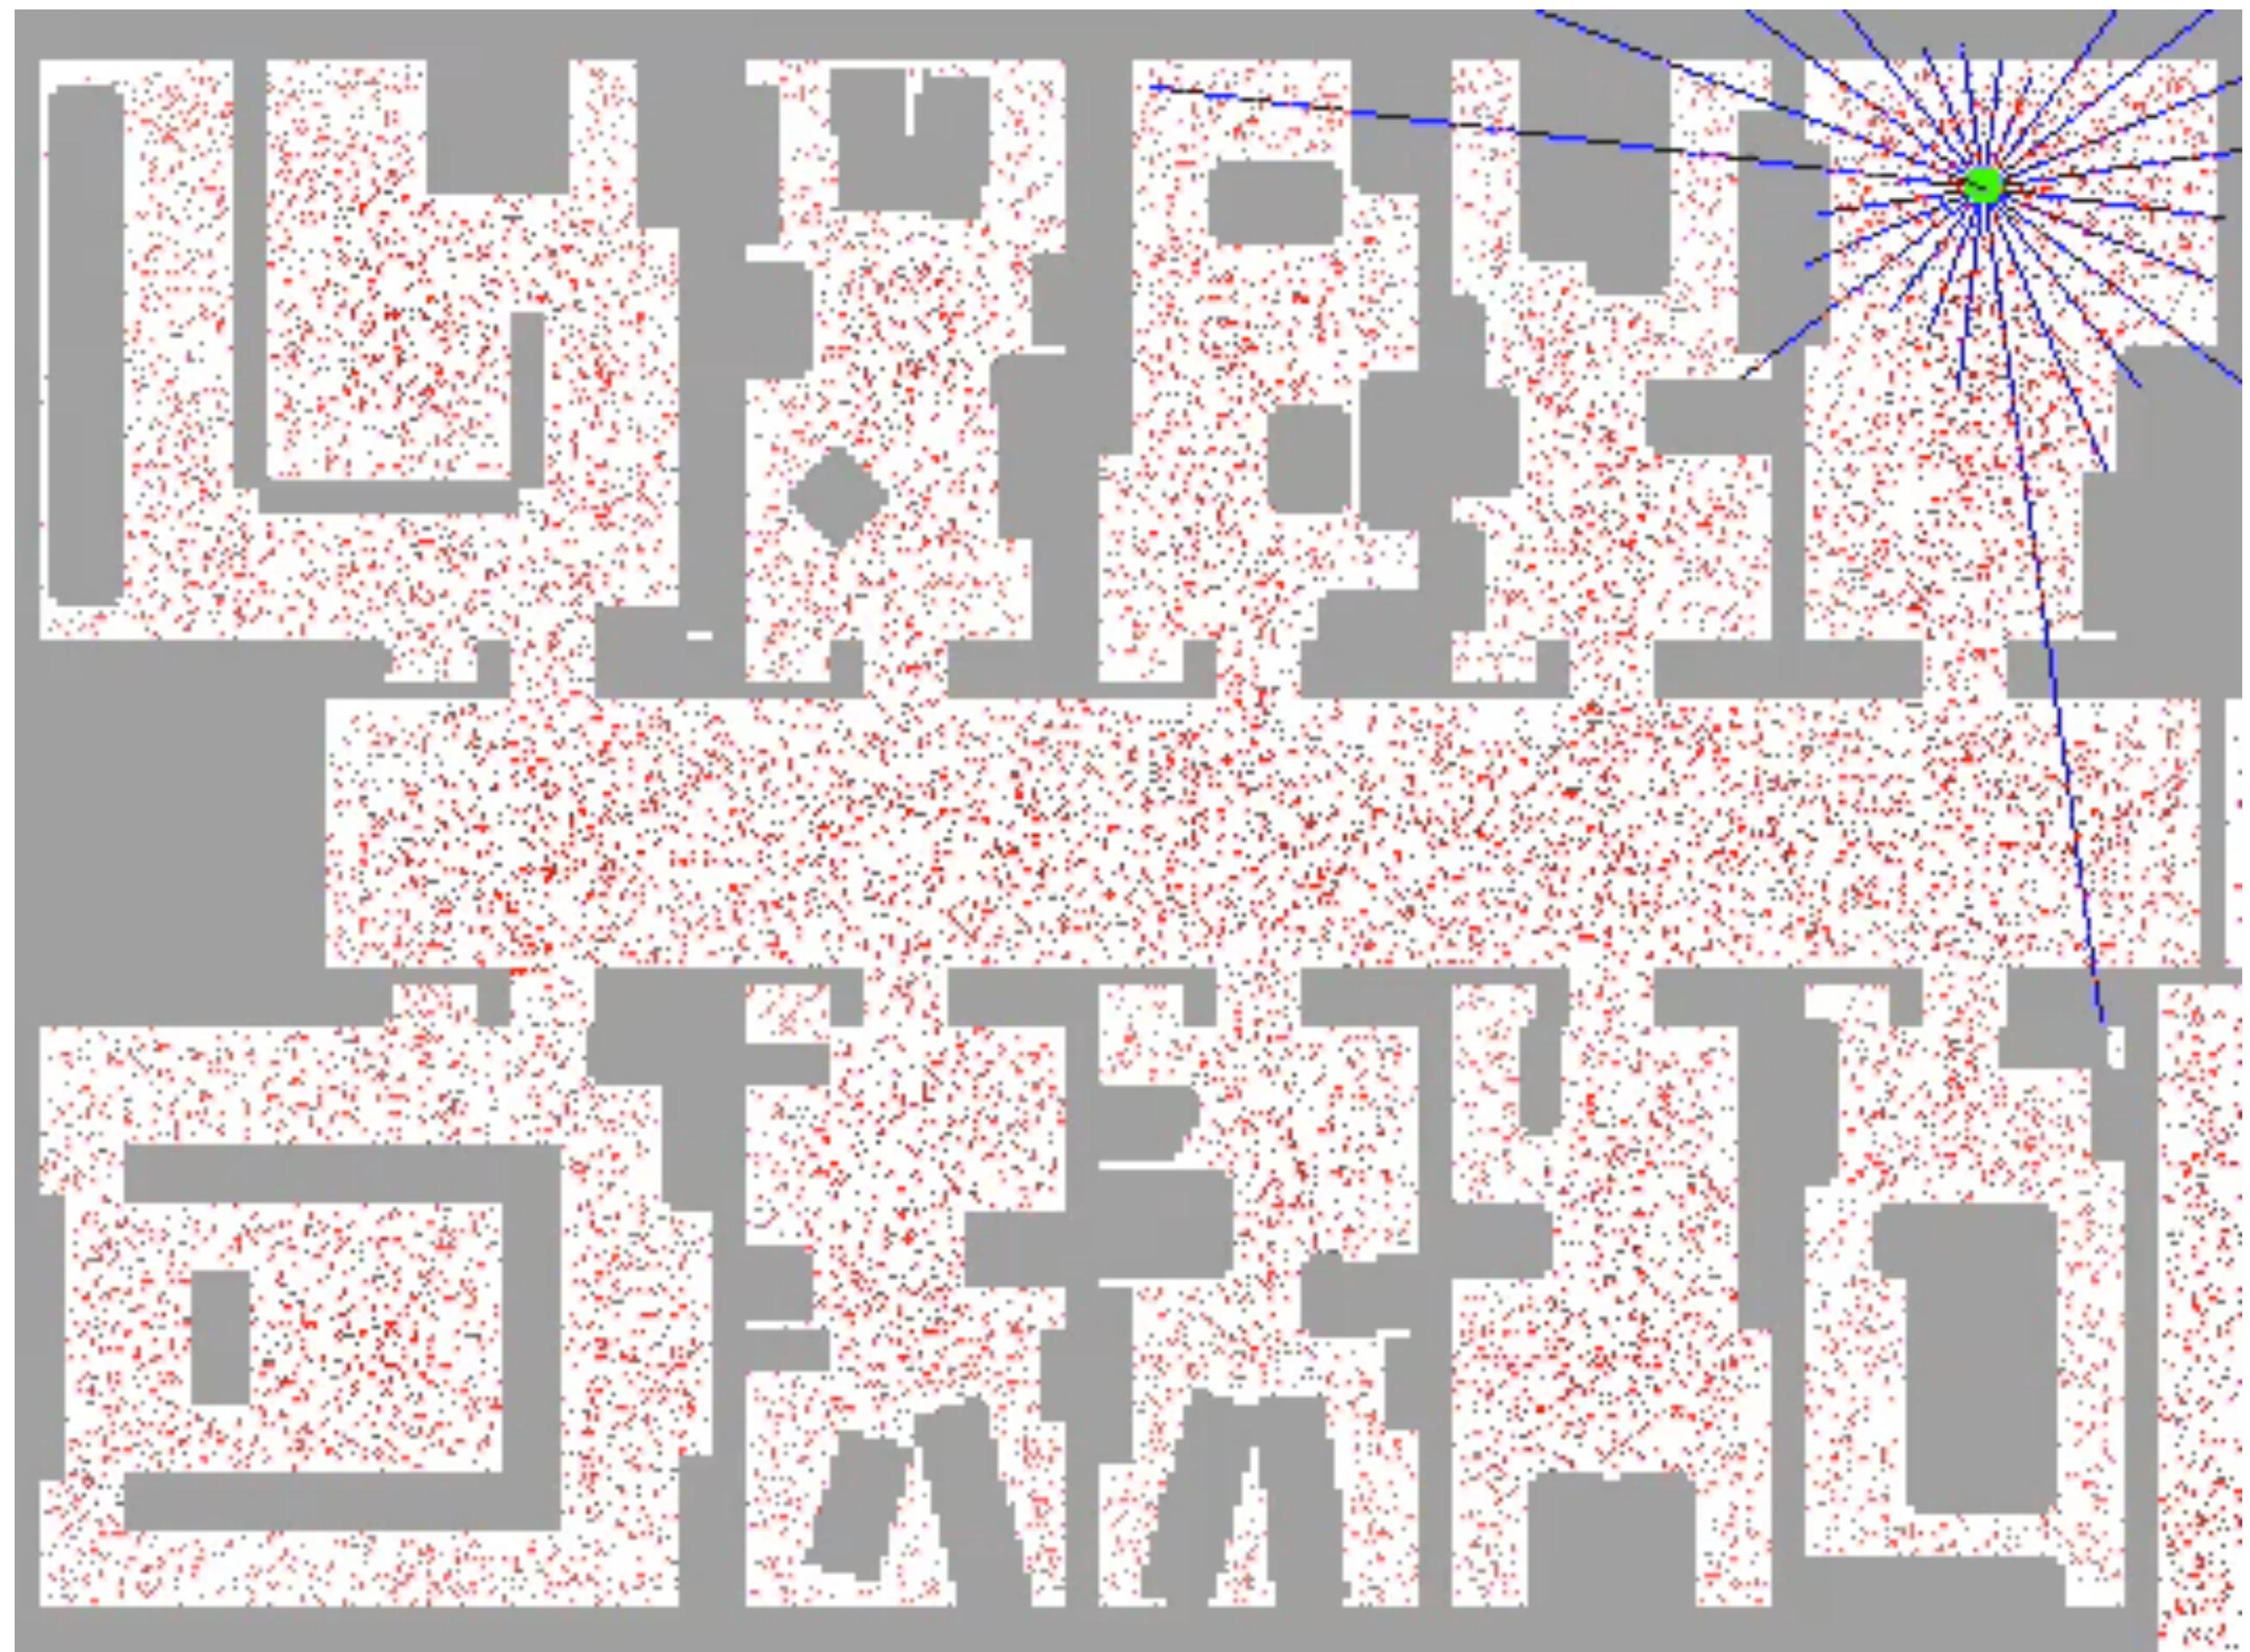


Lecture 20

Mobile Robotics - IV -

Particle Filter



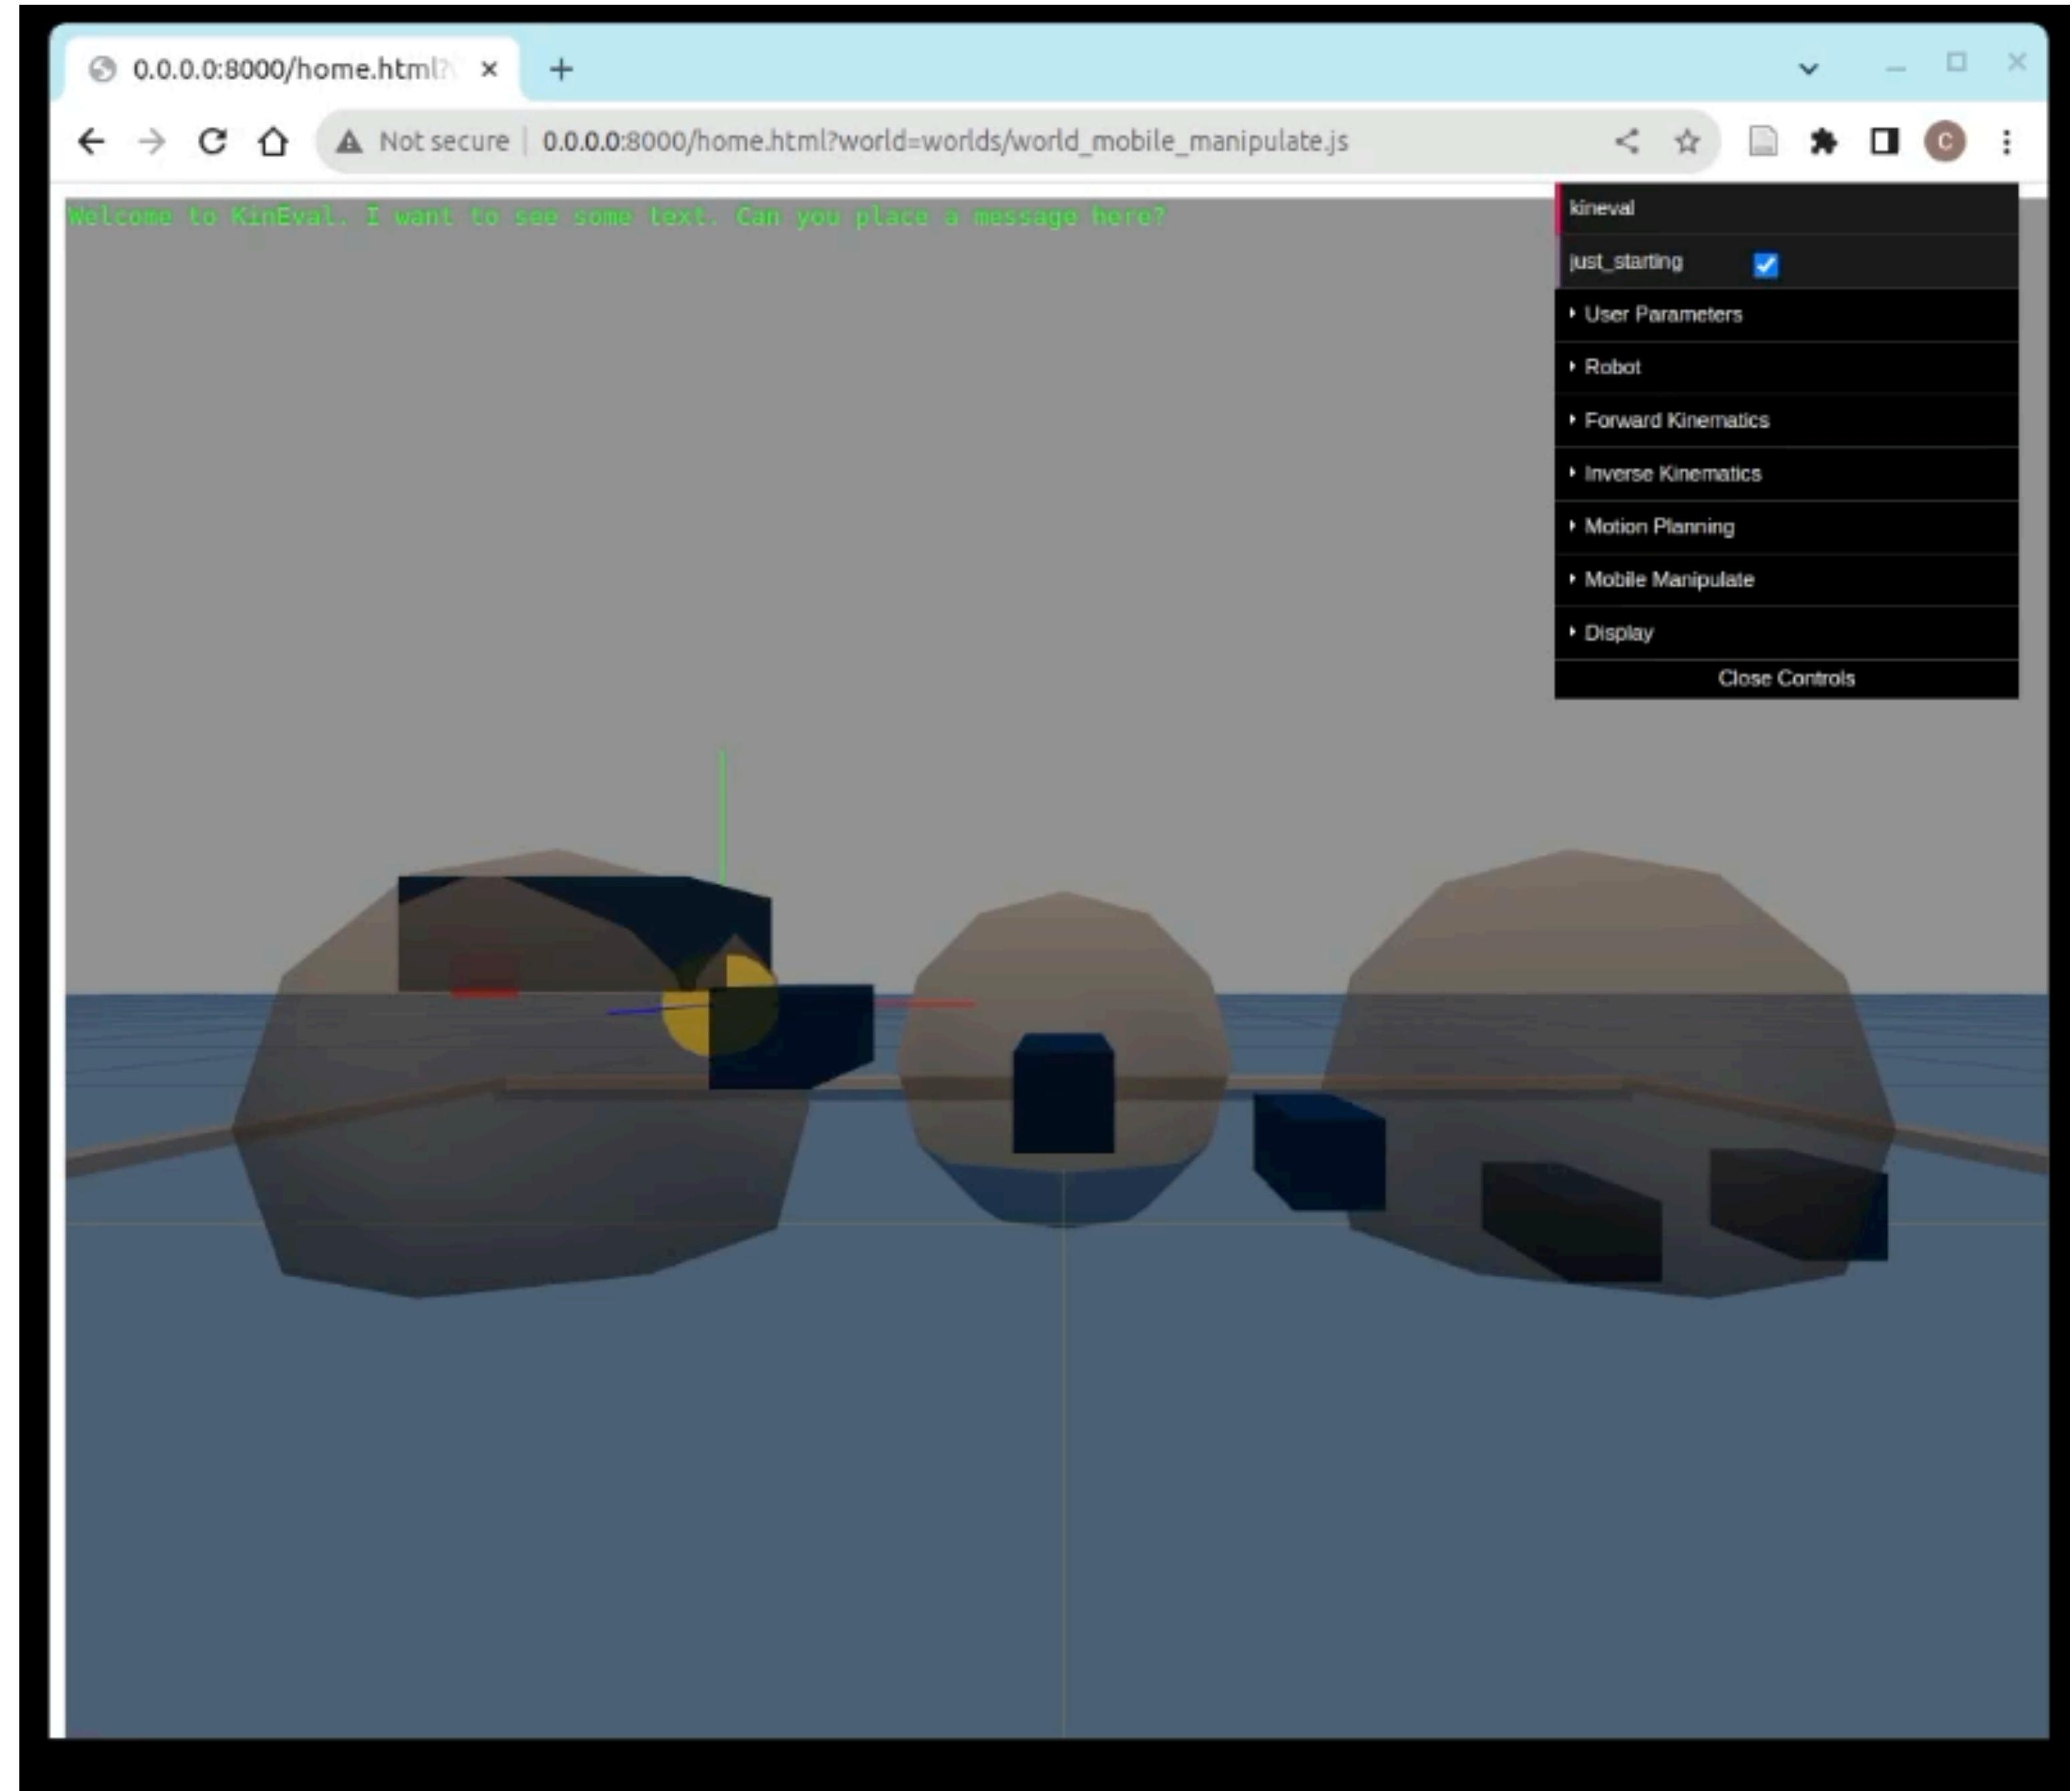
Course logistics

- Quiz 18-19 was posted yesterday and was due before today's lecture
- Project 4 is posted on 10/30 and will be due 11/15 (today)
- Project 5 will be posted today
 - Start early!

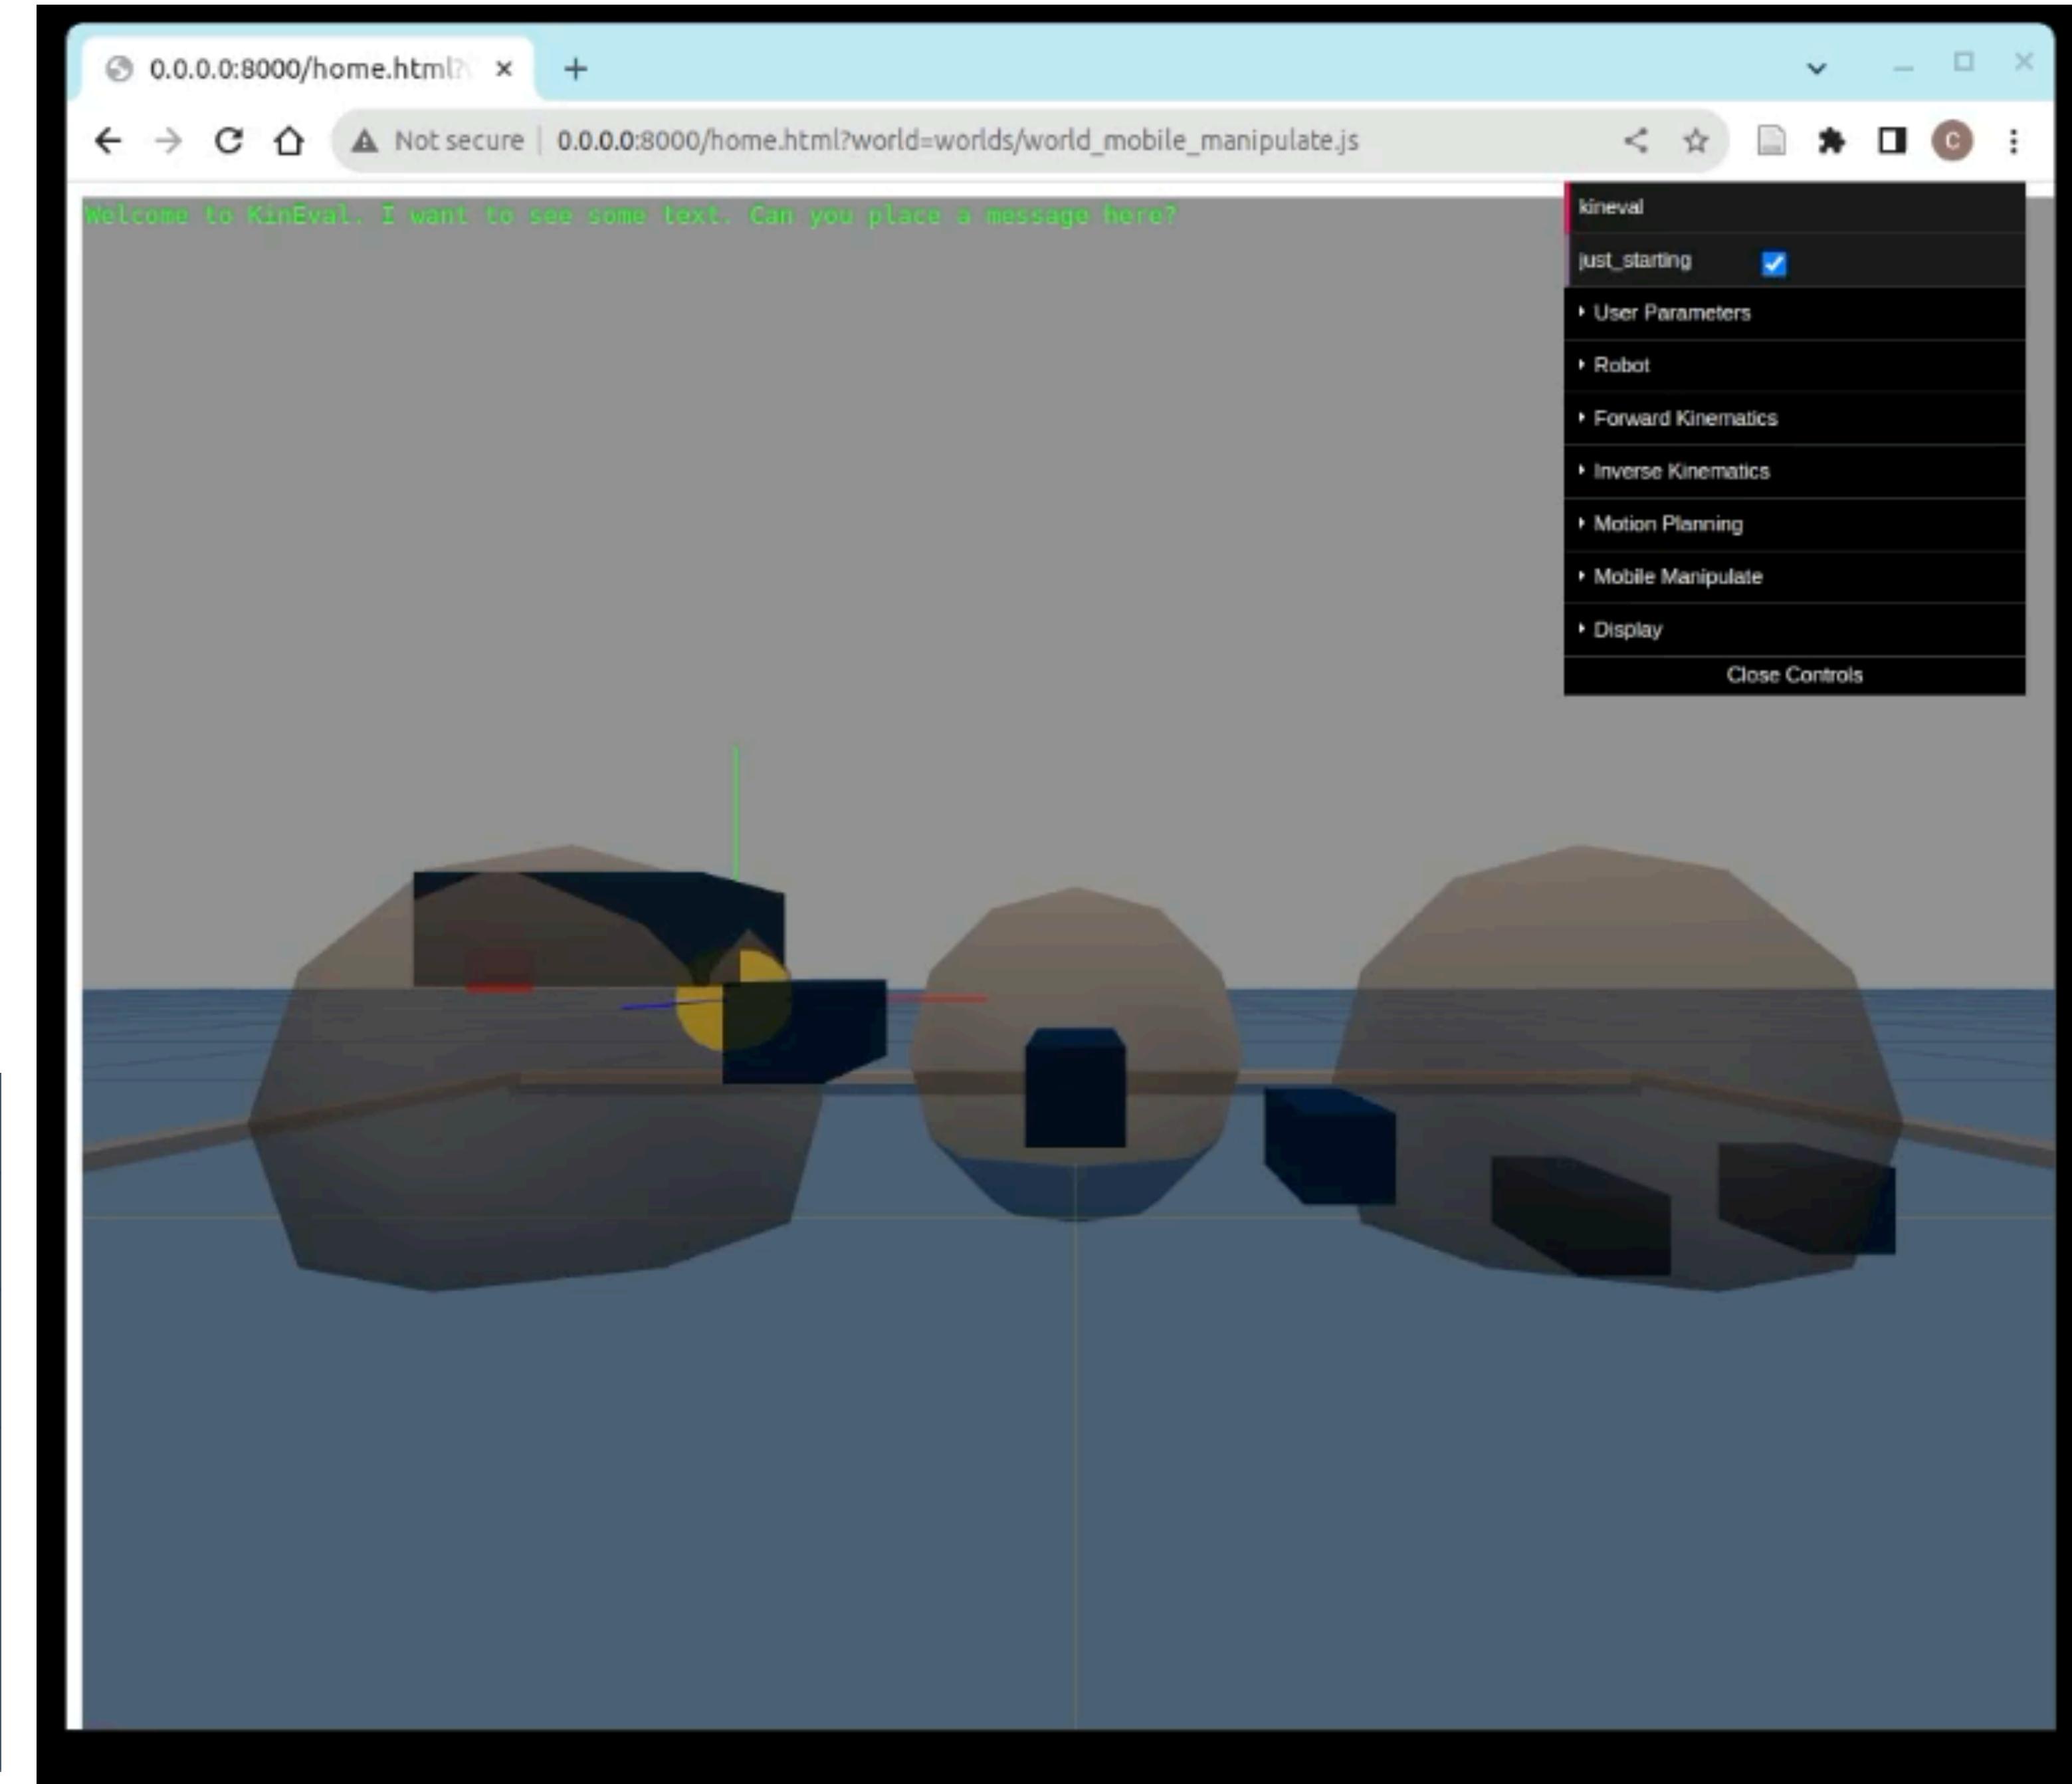
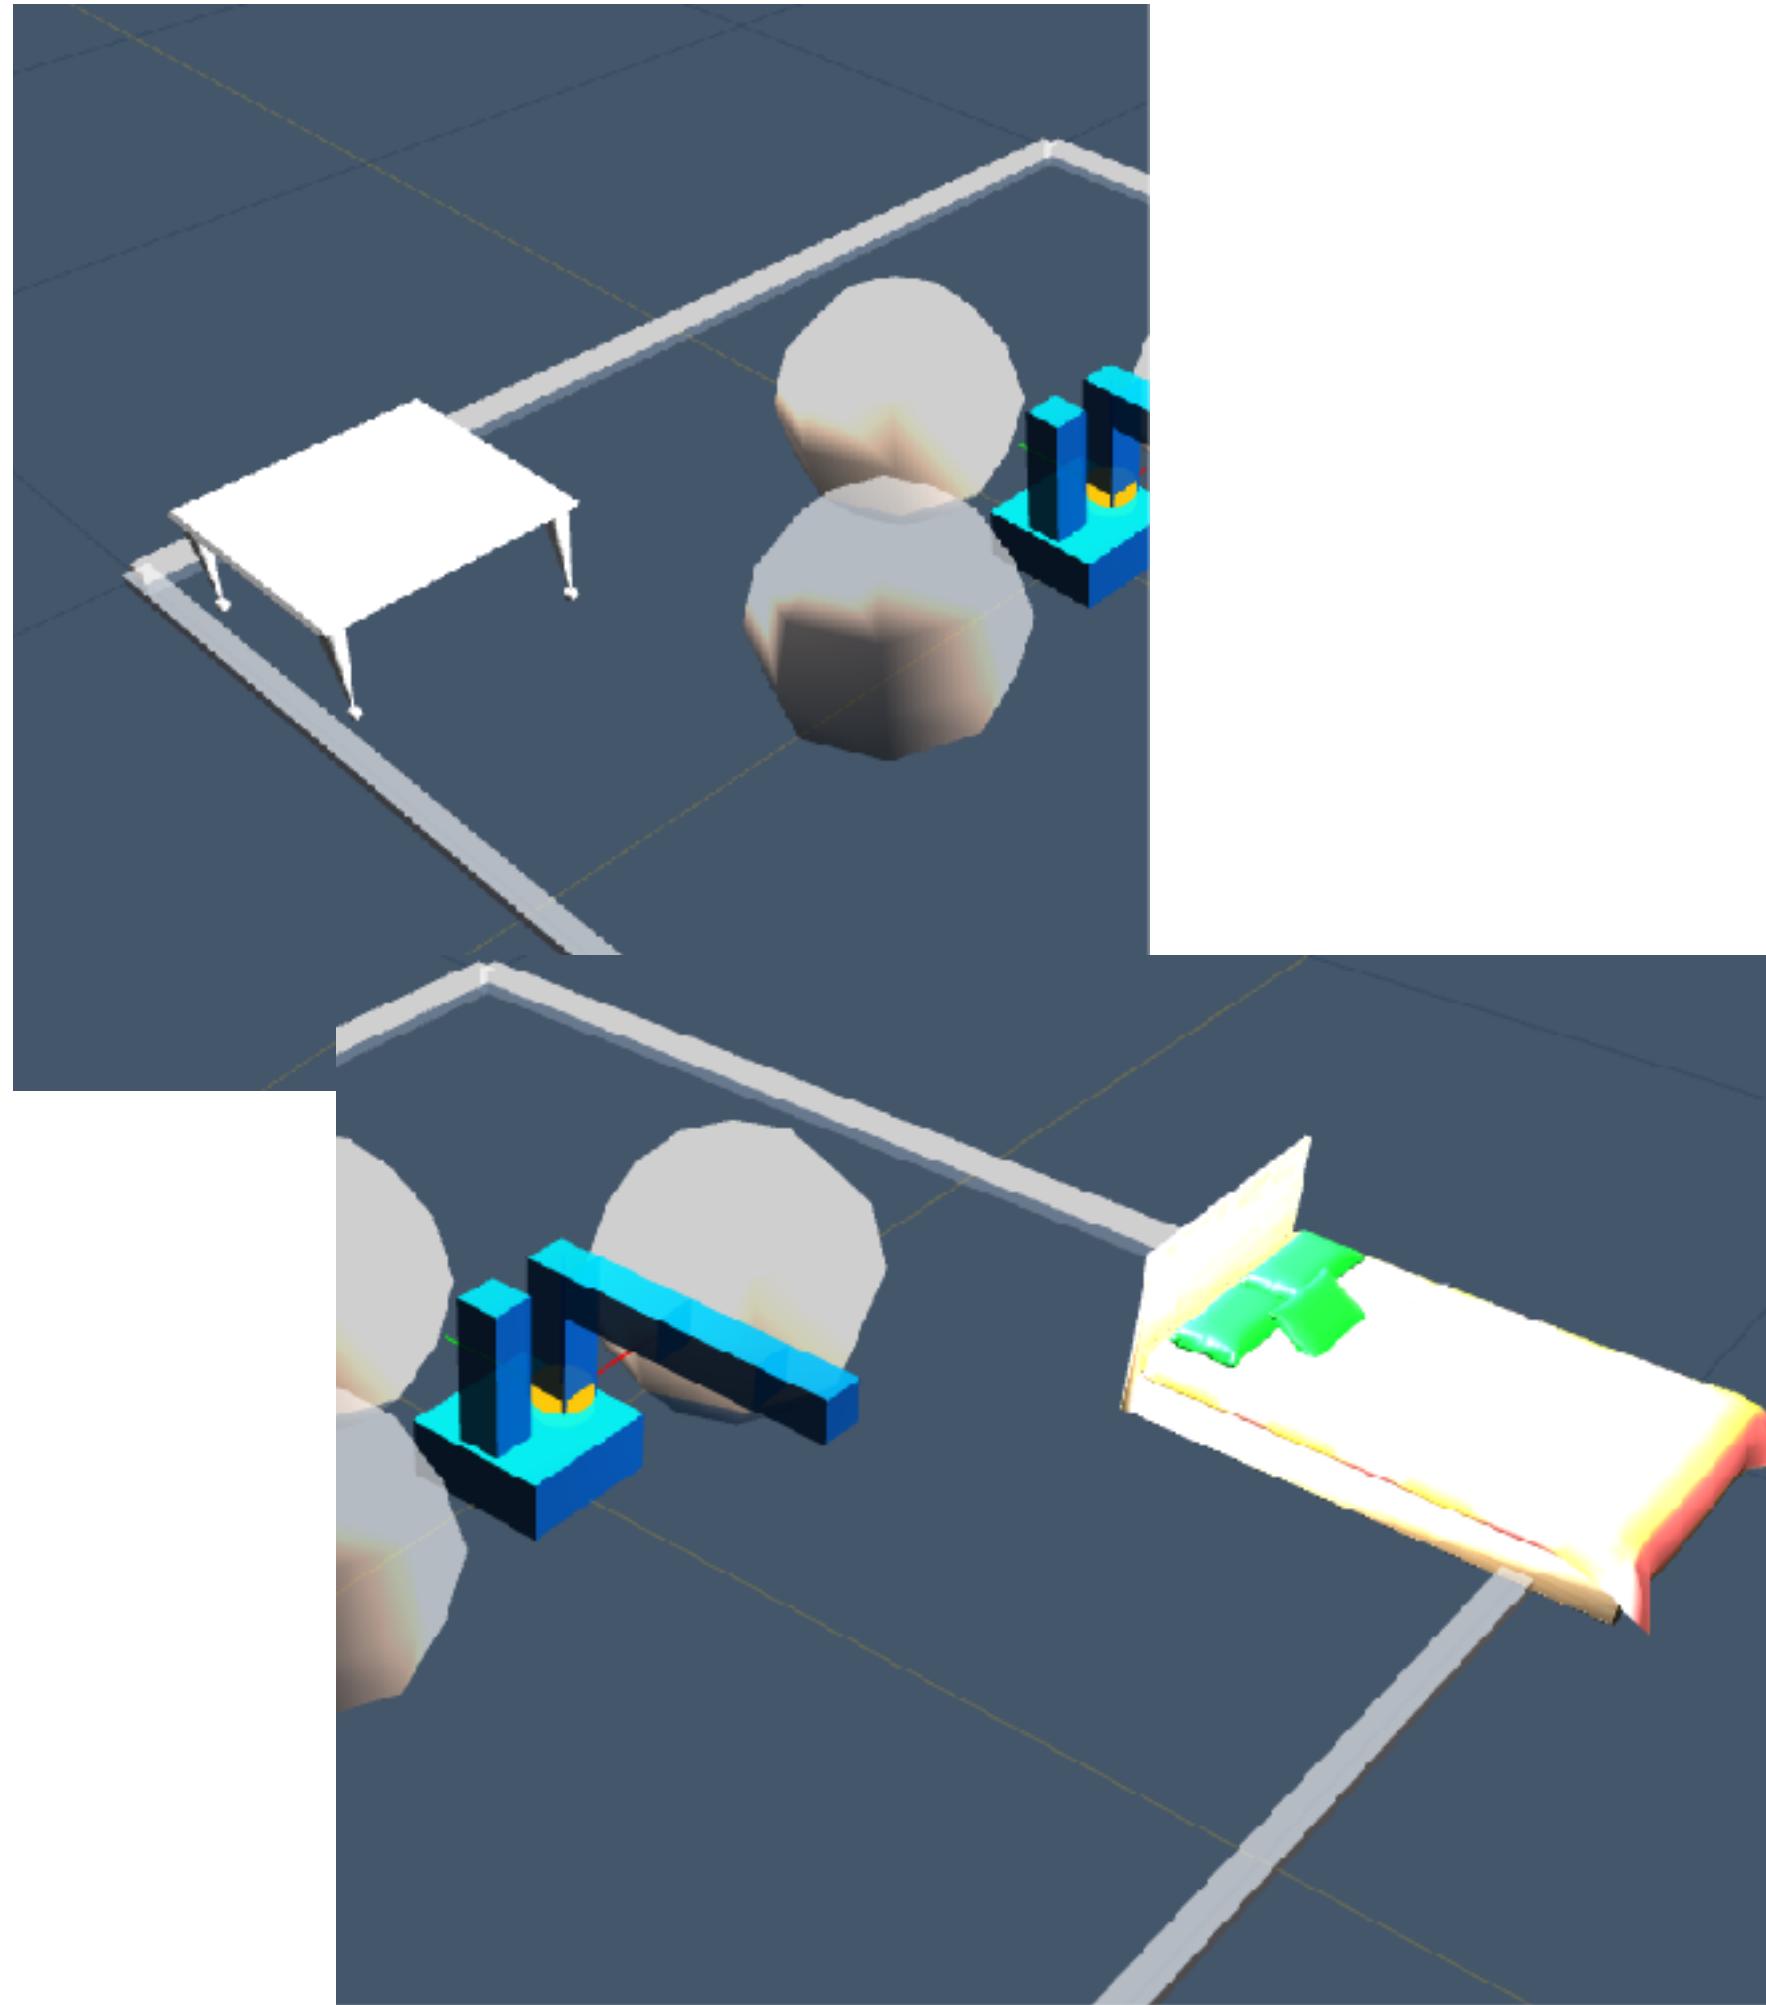


Project-5

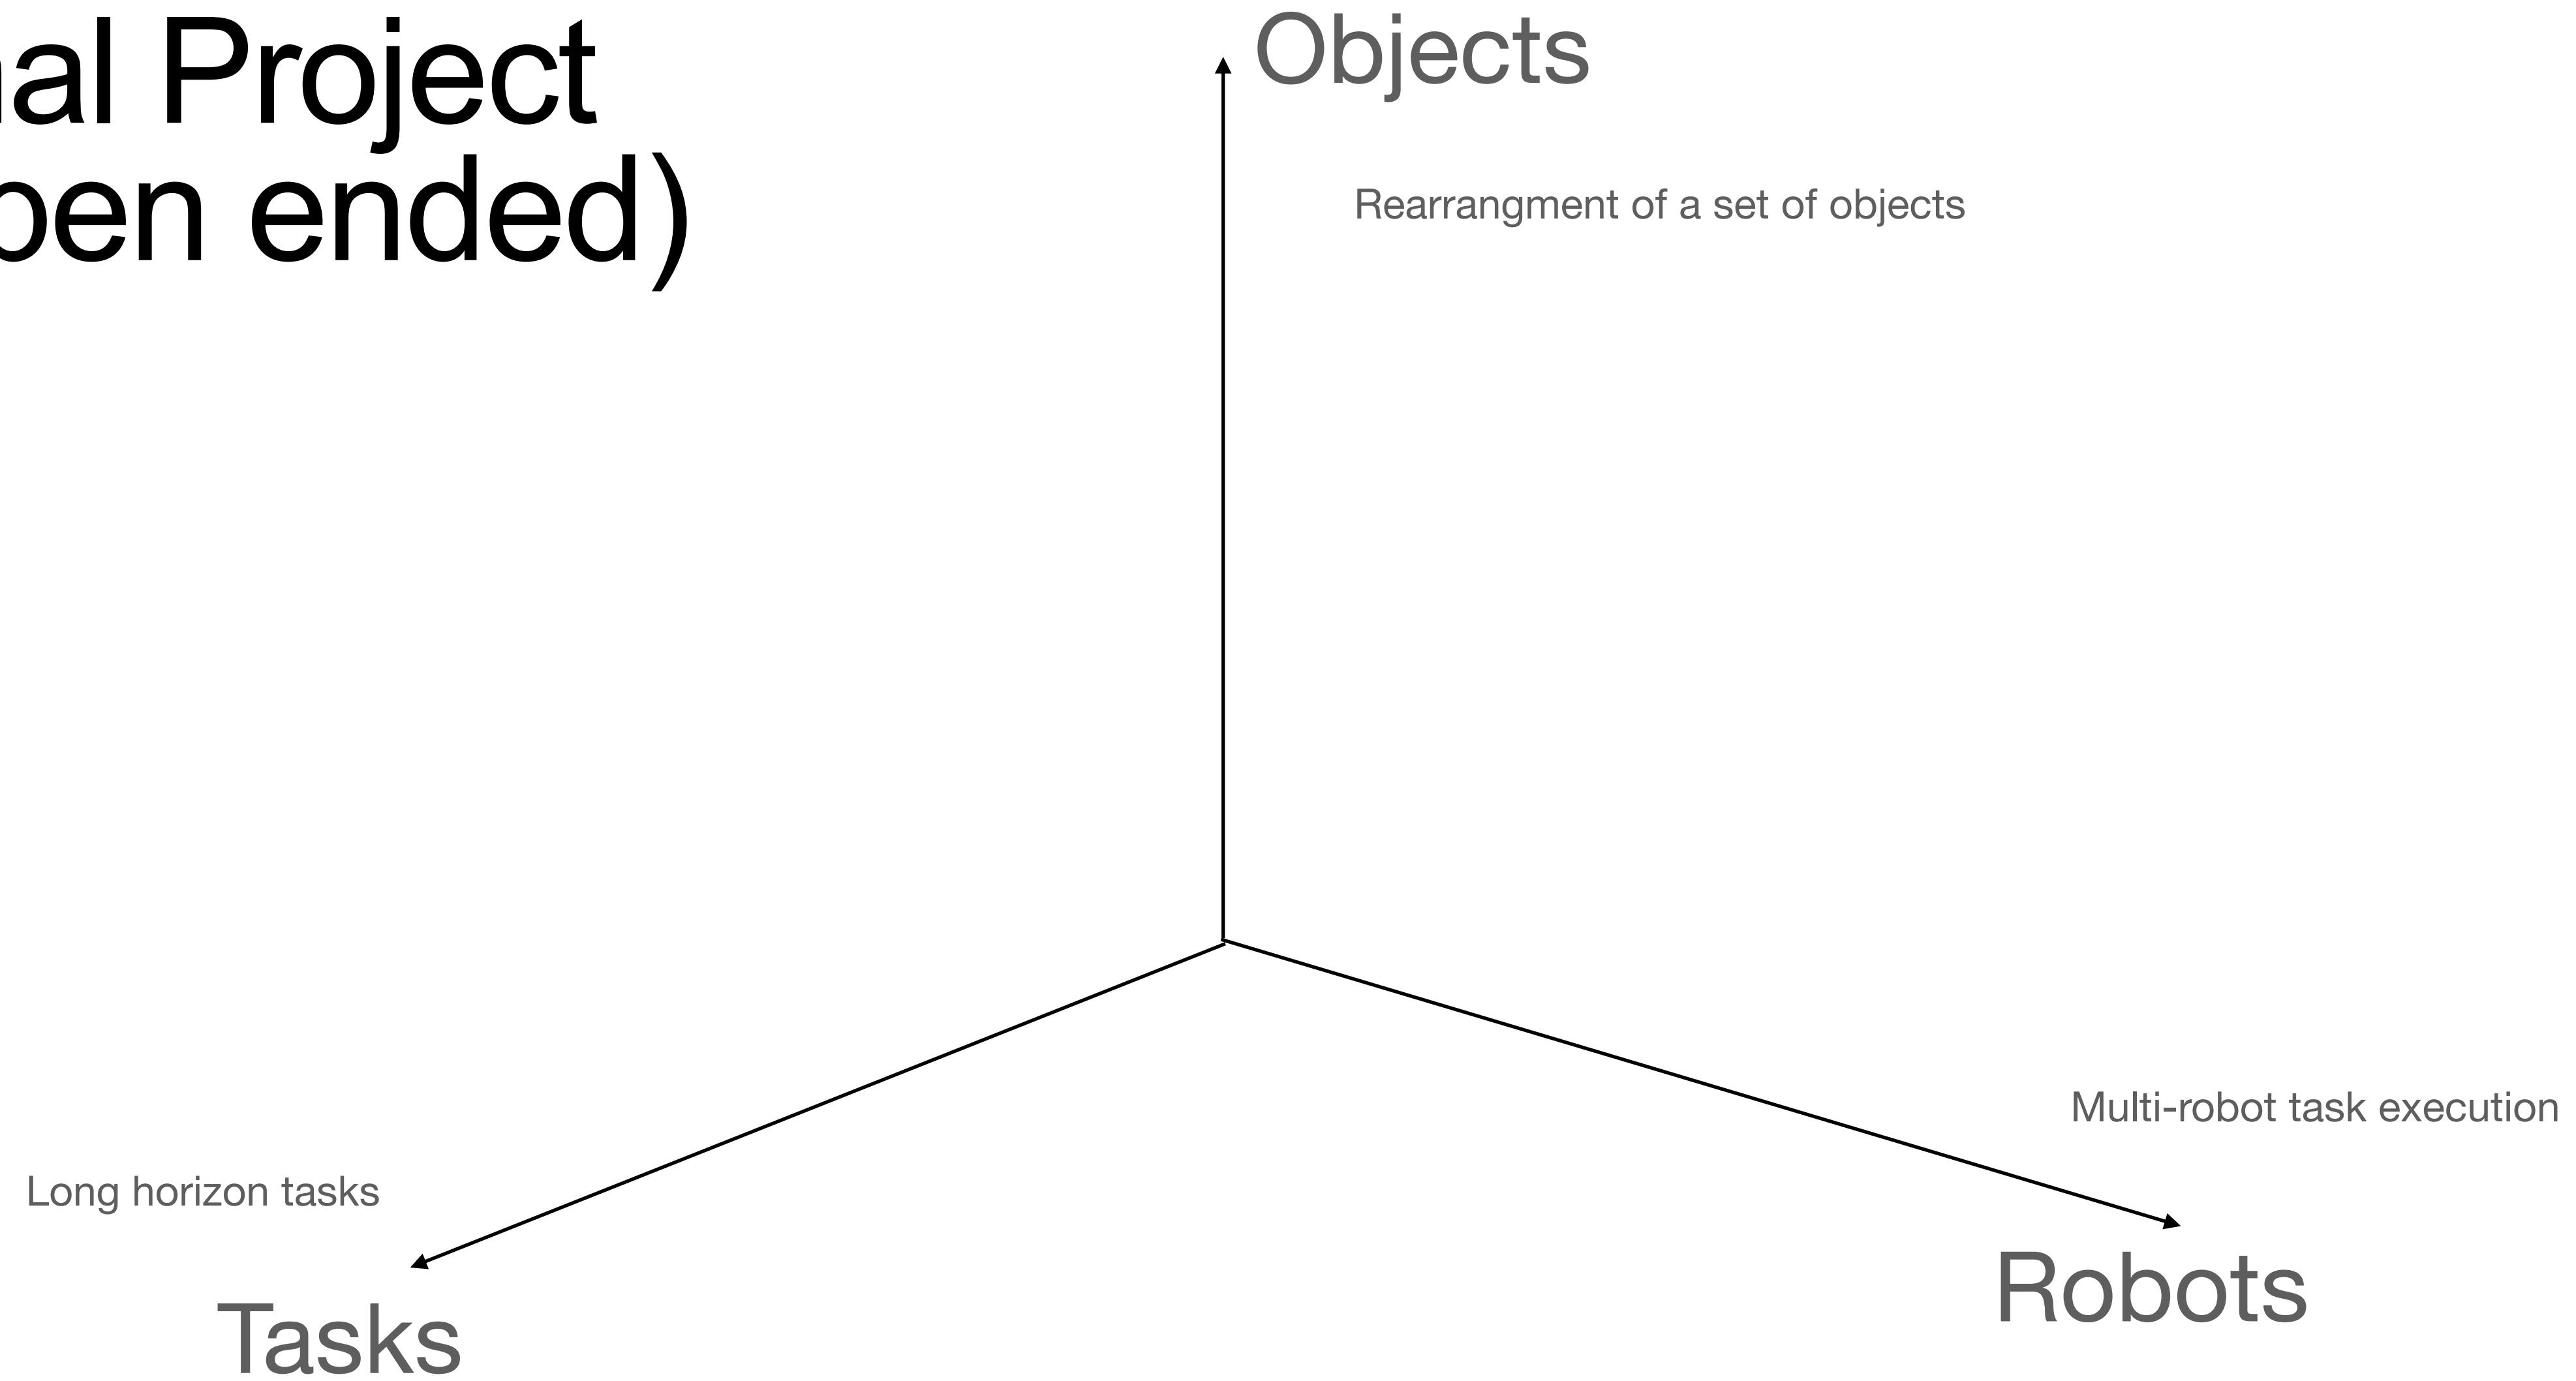
- RRT-Connect
- FSM
- IK



Final Project (Open ended)



Final Project (Open ended)



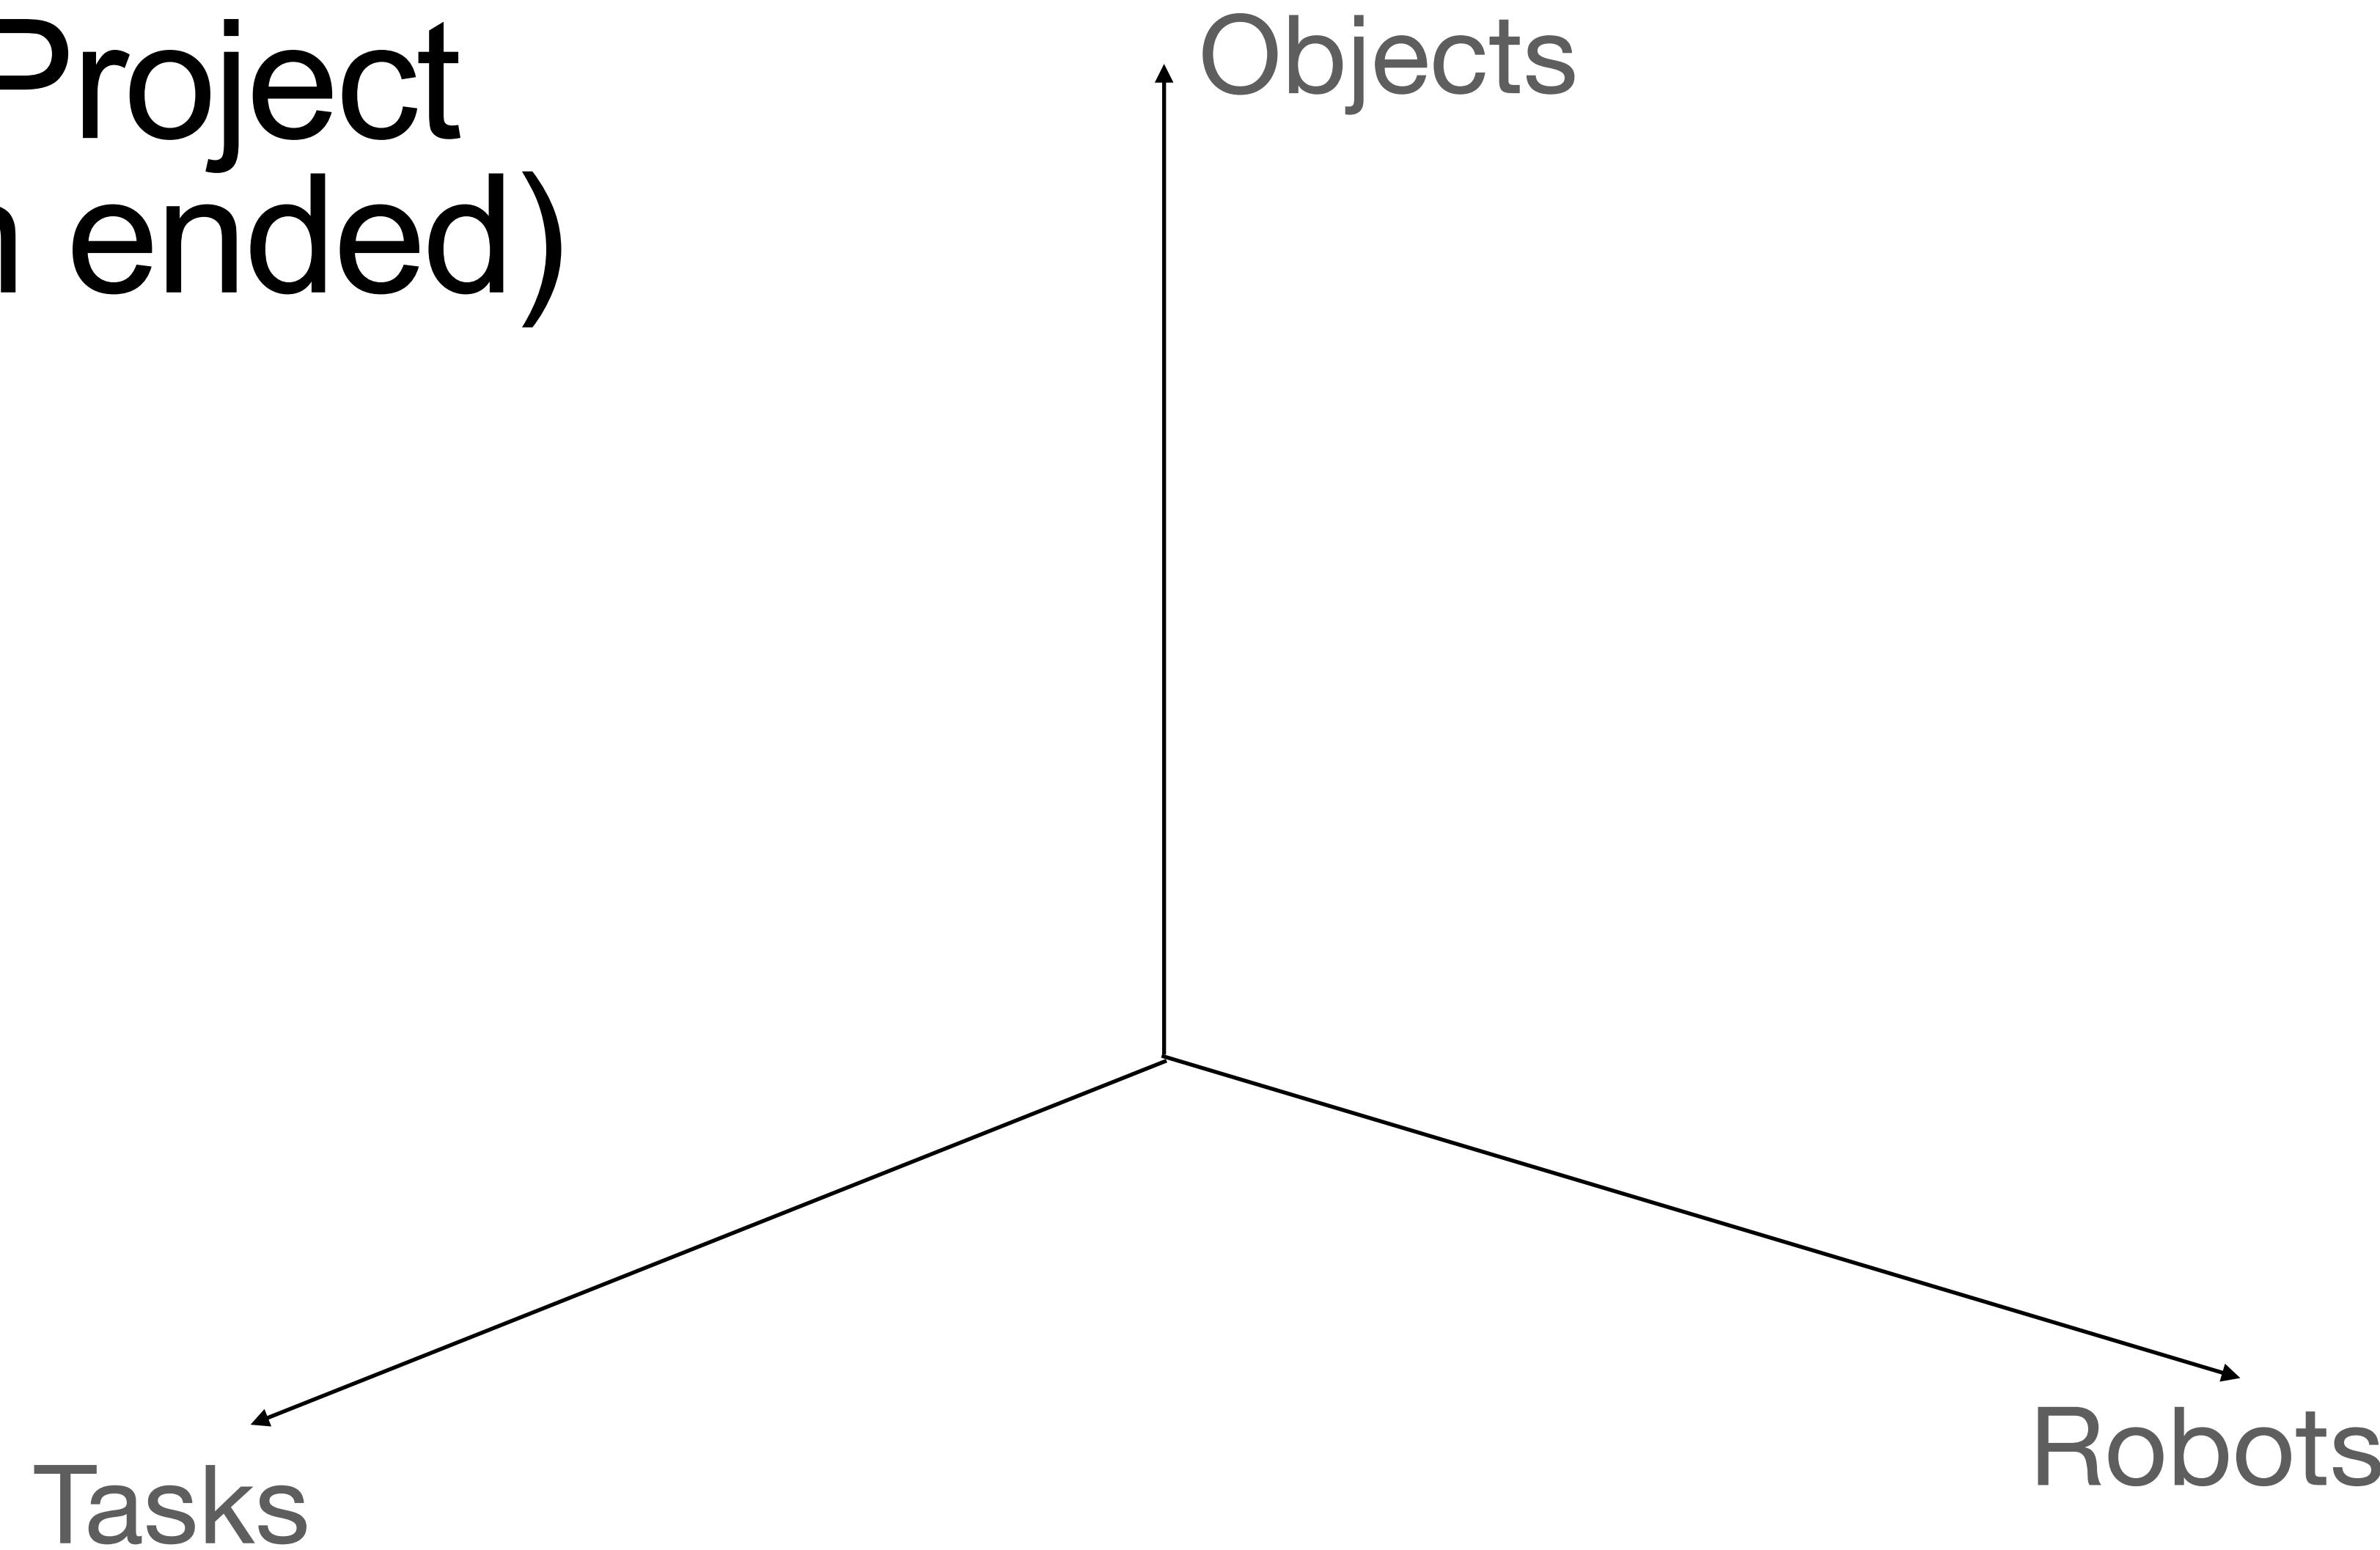
Final Project (Open ended)

For inspiration!



Yang, Zhutian, Caelan Reed Garrett, Tomás Lozano-Pérez, Leslie Kaelbling, and Dieter Fox. "Sequence-based plan feasibility prediction for efficient task and motion planning." *arXiv preprint arXiv:2211.01576* (2022).

Final Project (Open ended)



Continuing previous Lecture

KF and EKF



Discrete Kalman Filter

Estimates the state x of a discrete-time controlled process that is governed by the linear stochastic difference equation

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$$

with a measurement

$$z_t = C_t x_t + \delta_t$$



Components of a Kalman Filter

A_t

Matrix ($n \times n$) that describes how the state evolves from $t-1$ to t without controls or noise.

B_t

Matrix ($n \times l$) that describes how the control u_t changes the state from $t-1$ to t .

C_t

Matrix ($k \times n$) that describes how to map the state x_t to an observation z_t .

ε_t

Random variables representing the process and measurement noise that are assumed to be independent and normally distributed with covariance R_t and Q_t respectively.

δ_t



Kalman Filter Algorithm

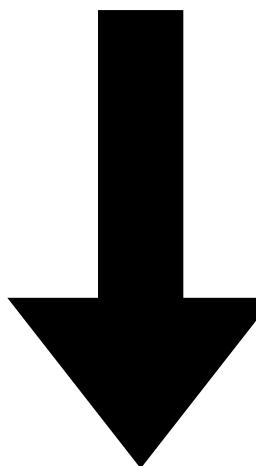
1. Algorithm **Kalman_filter**($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$):
2. Prediction:
3. $\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$
4. $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$
5. Correction:
6. $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$
7. $\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$
8. $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$
9. Return μ_t, Σ_t



Non-linear Dynamic Systems

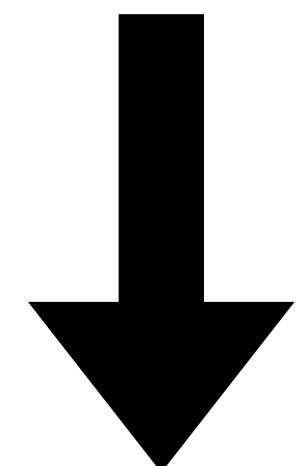
- Most realistic problems involve nonlinear functions

$$x_t = A_t x_{t-1} + B_t u_t + \epsilon_t$$



$$x_t = g(u_t, x_{t-1}) + \epsilon_t$$

$$z_t = C_t x_t + \delta_t$$



$$z_t = h(x_t) + \delta_t$$

Linearization

$$x_t = g(u_t, x_{t-1}) + \varepsilon_t$$

$$\begin{aligned} g(u_t, x_{t-1}) &\approx g(u_t, \mu_{t-1}) + \underbrace{g'(u_t, \mu_{t-1})}_{=: G_t} (x_{t-1} - \mu_{t-1}) \\ &= g(u_t, \mu_{t-1}) + G_t (x_{t-1} - \mu_{t-1}) \end{aligned}$$

$$z_t = h(x_t) + \delta_t$$

$$h(x_t) \approx h(\bar{\mu}_t) + \underbrace{\frac{\partial h(\bar{\mu}_t)}{\partial x_t}}_{=: H_t} (x_t - \bar{\mu}_t)$$



EKF Algorithm

1. **Extended_Kalman_filter**($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$):

2. Prediction:

$$\begin{array}{ll} 3. \quad \bar{\mu}_t = g(u_t, \mu_{t-1}) & \longleftarrow \quad \bar{\mu}_t = A_t \mu_{t-1} + B_t u_t \\ 4. \quad \bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t & \longleftarrow \quad \bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t \end{array}$$

5. Correction:

$$\begin{array}{ll} 6. \quad K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1} & \longleftarrow \quad K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1} \\ 7. \quad \mu_t = \bar{\mu}_t + K_t (z_t - h(\bar{\mu}_t)) & \longleftarrow \quad \mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t) \\ 8. \quad \Sigma_t = (I - K_t H_t) \bar{\Sigma}_t & \longleftarrow \quad \Sigma_t = (I - K_t C_t) \bar{\Sigma}_t \end{array}$$

9. Return μ_t, Σ_t

$$H_t = \frac{\partial h(\bar{\mu}_t)}{\partial x_t} \qquad G_t = \frac{\partial g(u_t, \mu_{t-1})}{\partial x_{t-1}}$$



Localization

“Using sensory information to locate the robot in its environment is the most fundamental problem to providing a mobile robot with autonomous capabilities.” [Cox ’91]

- **Given**

- Map of the environment.
- Sequence of sensor measurements.

- **Wanted**

- Estimate of the robot’s position.

- **Problem classes**

- Position tracking
- Global localization
- Kidnapped robot problem (recovery)



EKF Summary

- Highly efficient: Polynomial in measurement dimensionality k and state dimensionality n :
 $O(k^{2.376} + n^2)$
- Not optimal!
- Can diverge if nonlinearities are large!
- Works surprisingly well even when all assumptions are violated!



Particle Filter

A Bayesian Filter Implementation

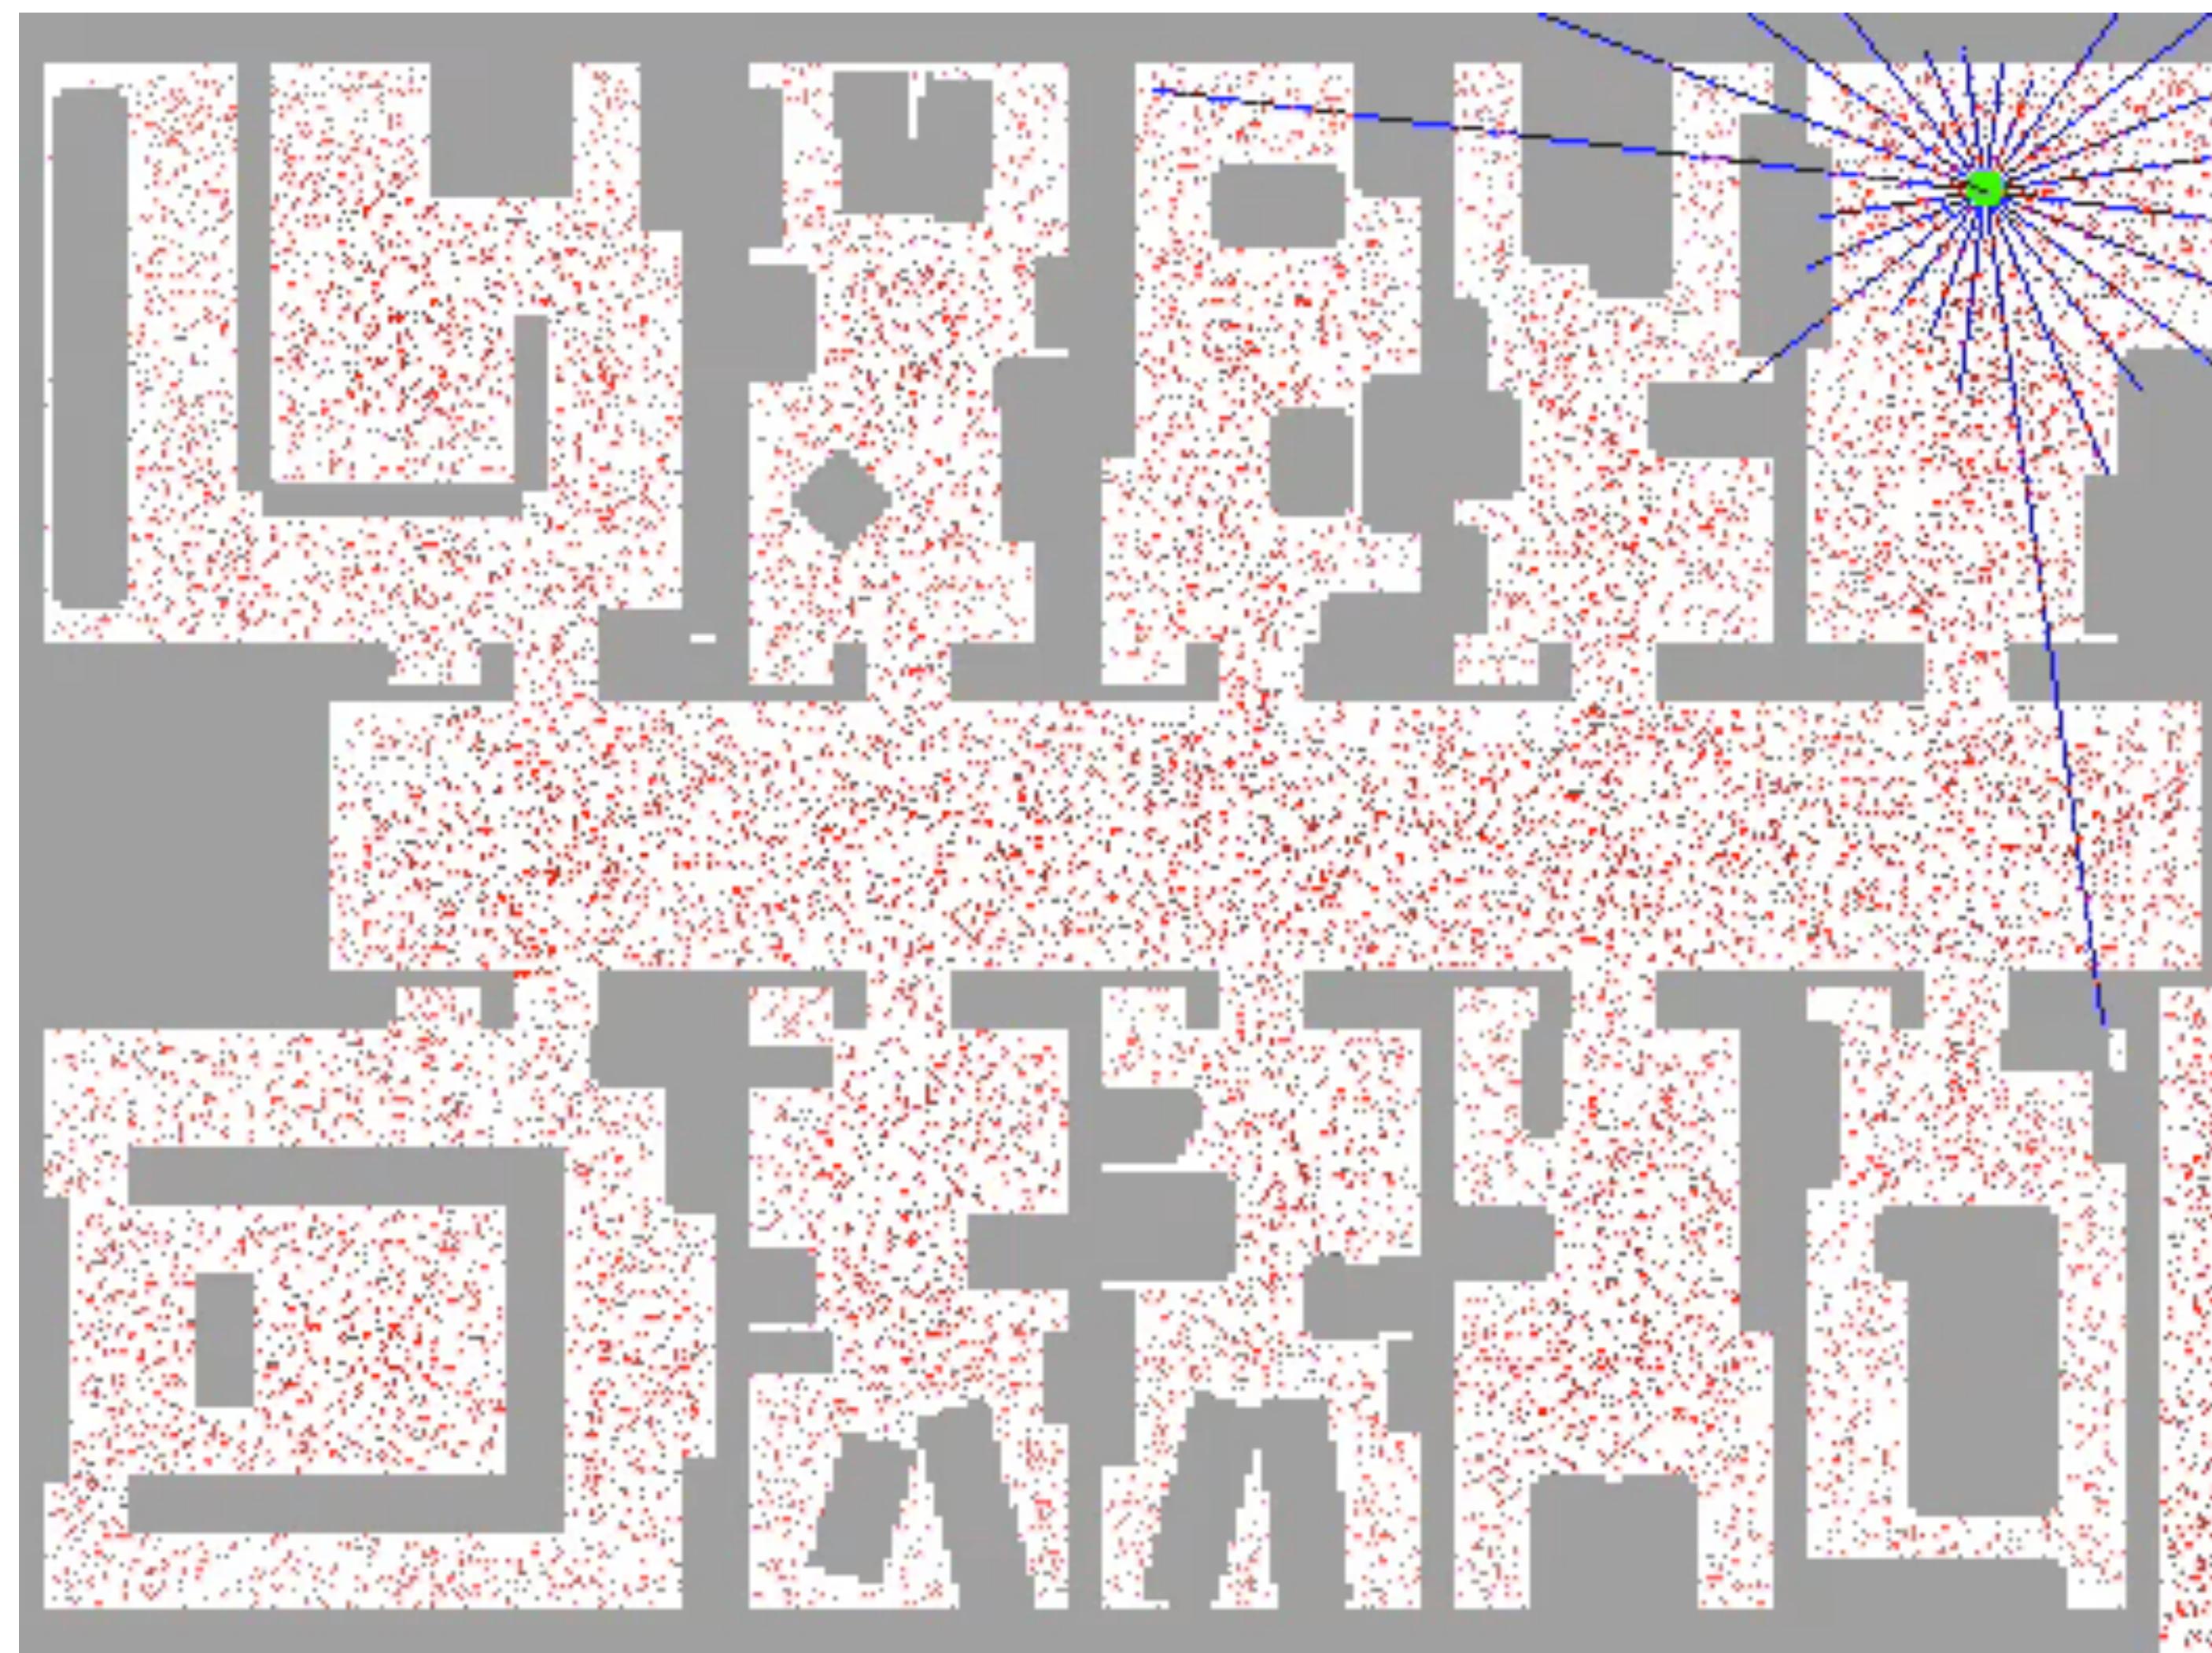


Motivation

- So far, we discussed the
 - Kalman filter: Gaussian, linearization problems, multi-modal beliefs
- Particle filters are a way to **efficiently** represent **non-Gaussian distributions**
- Basic principle
 - Set of state hypotheses ("particles")
 - Survival-of-the-fittest

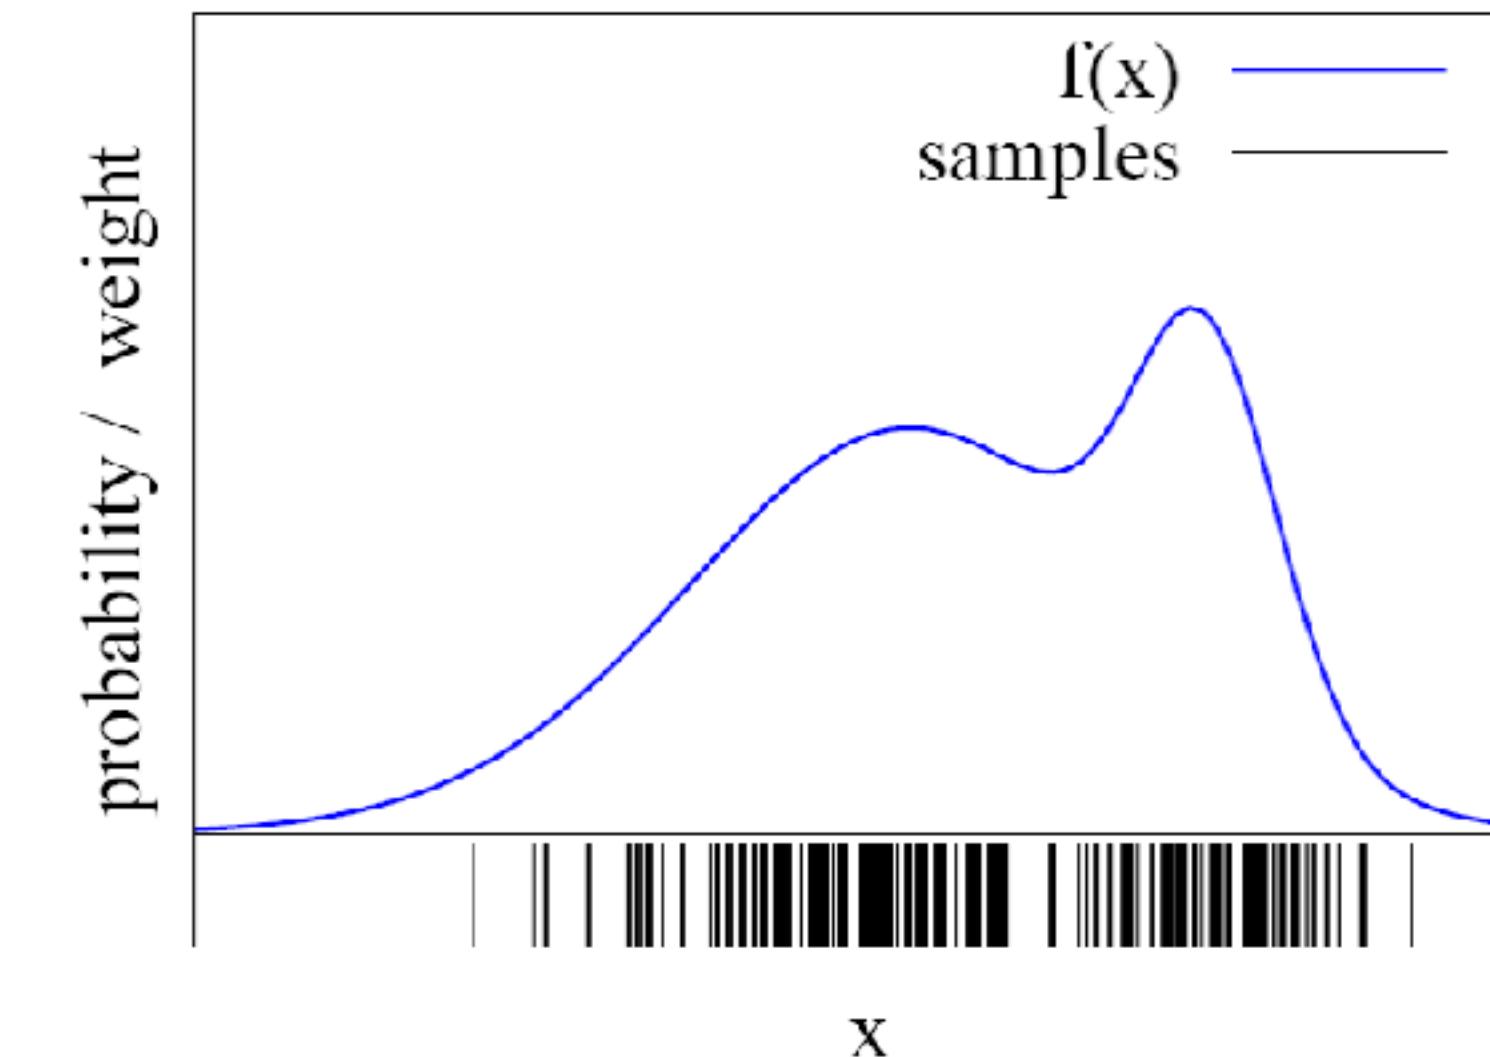
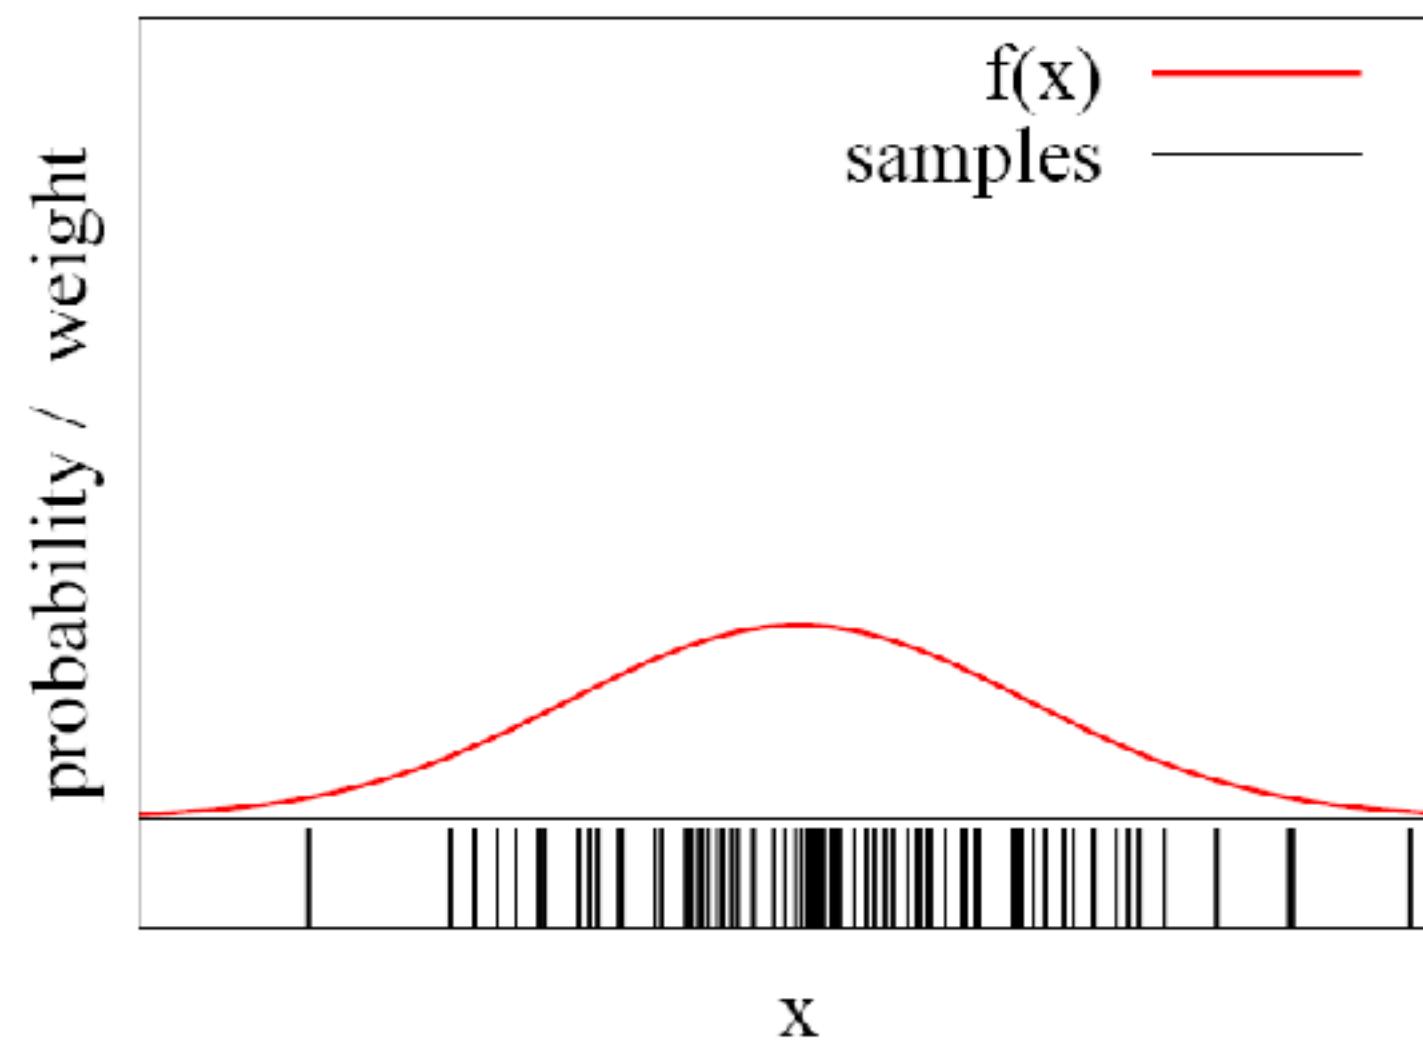


Sample-based Localization (sonar)



Density Approximation

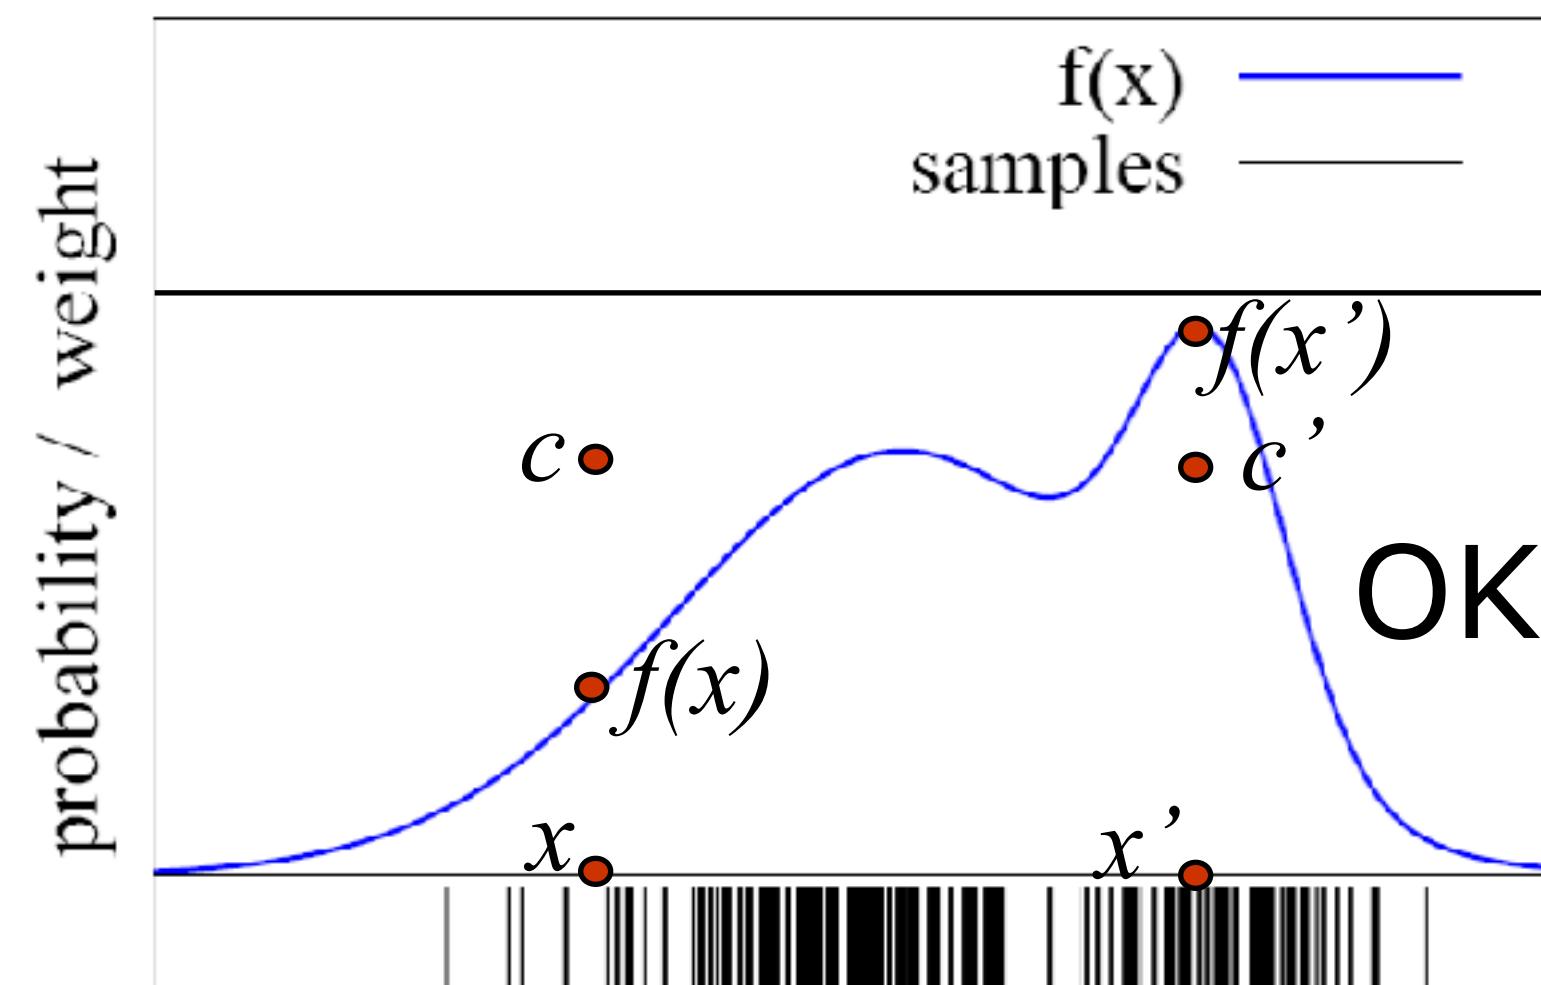
- Particle sets can be used to approximate densities



- The more particles fall into an interval, the higher the probability of that interval
- How to draw samples from a function/distribution?

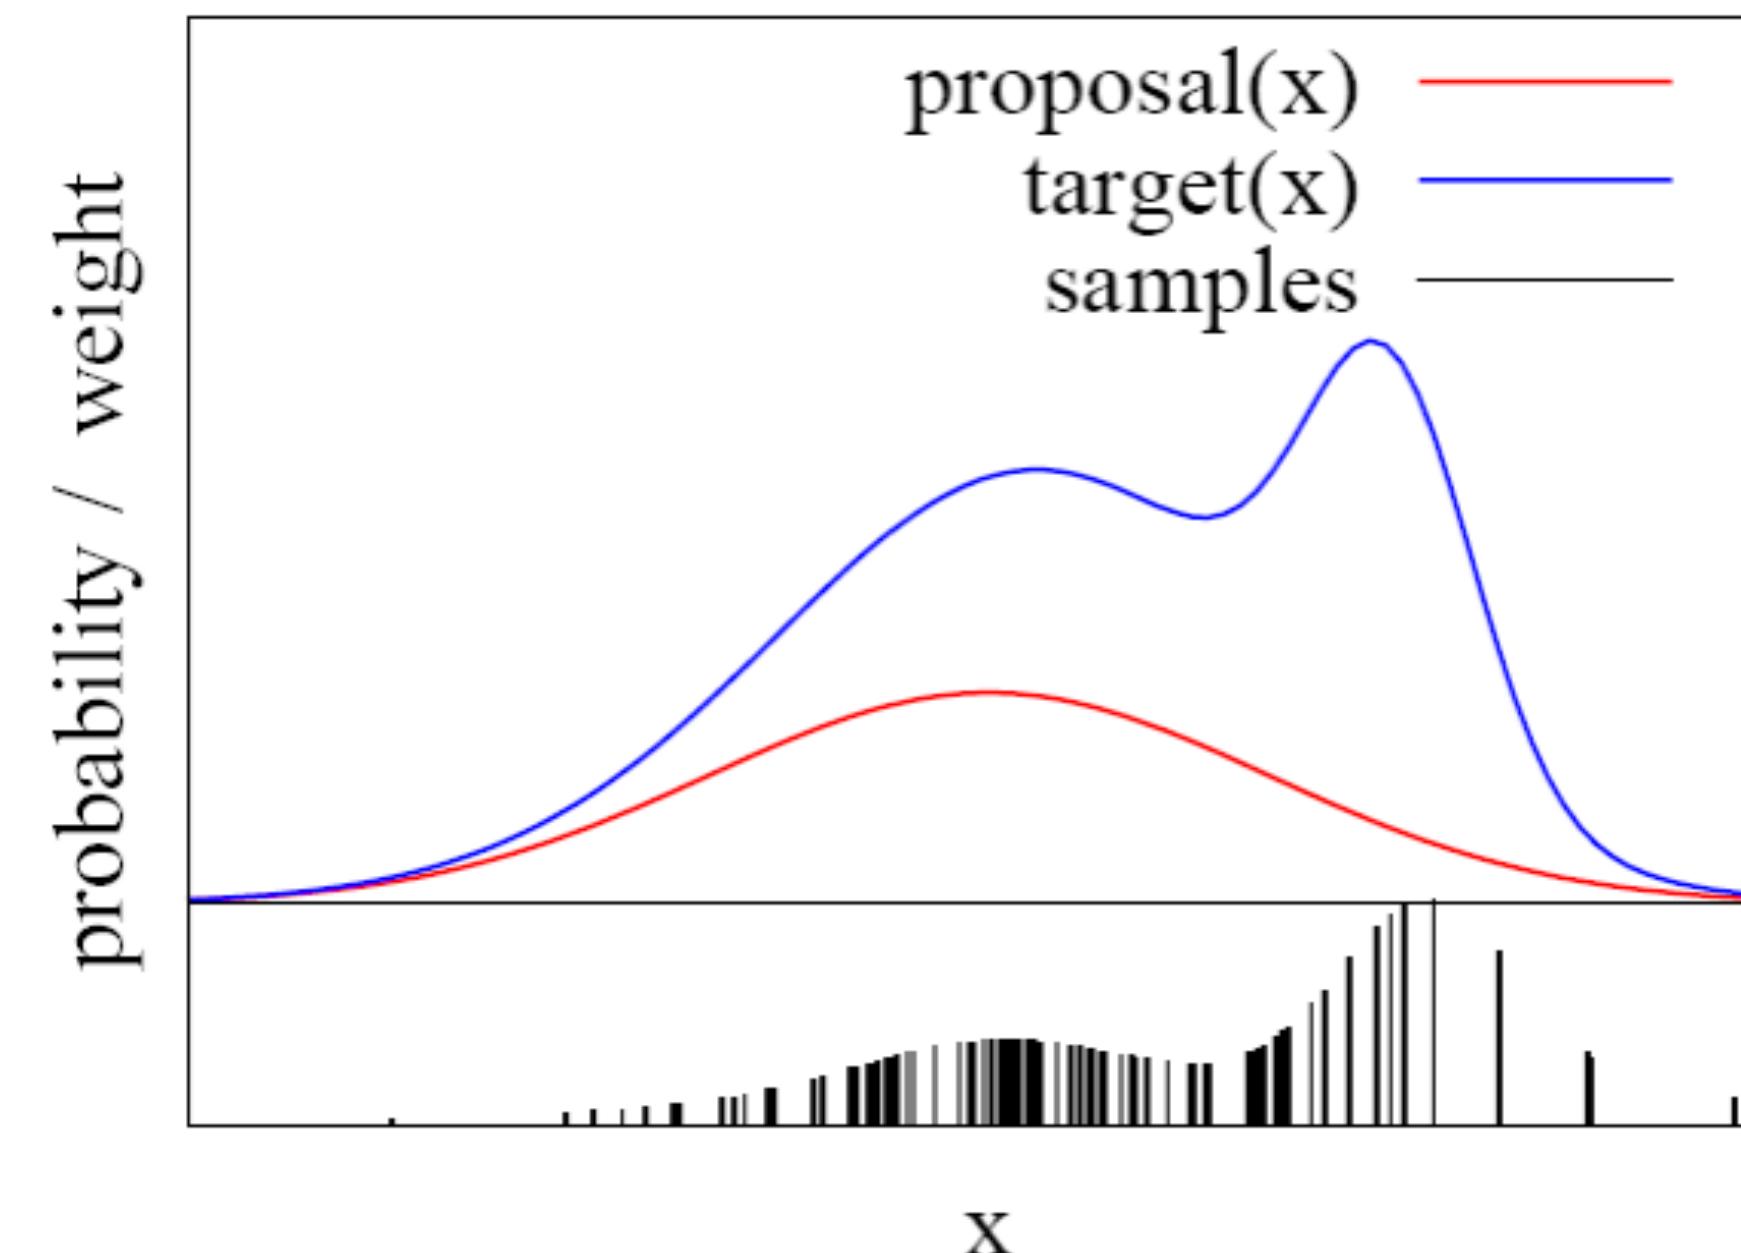
Rejection Sampling

- Let us assume that $f(x) \leq 1$ for all x
- Sample x from a uniform distribution
- Sample c from $[0,1]$
- if $f(x) > c$ keep the sample
otherwise reject the sample



Importance Sampling Principle

- We can even use a different distribution g to generate samples from f
- By introducing an importance weight w , we can account for the “differences between g and f ”
- $w = f/g$
- f is often called target
- g is often called proposal



Particle Filter for State estimation

- Non-parametric approach
- Recursive Bayes Filter
- Models the distribution by samples
- **Prediction:** draw from the proposal g
- **Correction:** weighting by the ratio of the target f and the proposal g

The more samples we use, the better is the estimate



Particle Filter Algorithm

1. Sample the particles using the proposal distribution.

$$x_t^{[j]} \sim \text{proposal}(x_t | \dots)$$

2. Compute the importance weights

$$w_t^{[j]} = \frac{\text{target}(x_t^{[j]})}{\text{proposal}(x_t^{[j]})}$$

3. Resampling: Draw samples i with probability $w_t^{[i]}$ and repeat J times



Particle Filter Algorithm

Particle_filter($\mathcal{X}_{t-1}, u_t, z_t$):

```
1:       $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
2:      for  $j = 1$  to  $J$  do
3:          sample  $x_t^{[j]} \sim \pi(x_t)$ 
4:           $w_t^{[j]} = \frac{p(x_t^{[j]})}{\pi(x_t^{[j]})}$ 
5:           $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[j]}, w_t^{[j]} \rangle$ 
6:      endfor
7:      for  $j = 1$  to  $J$  do
8:          draw  $i \in 1, \dots, J$  with probability  $\propto w_t^{[i]}$ 
9:          add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
10:     endfor
11:     return  $\mathcal{X}_t$ 
```



Particle Filter Algorithm

$$Bel(x_t) = \eta p(z_t | x_t) \int p(x_t | x_{t-1}, u_{t-1}) Bel(x_{t-1}) dx_{t-1}$$

- draw x_{t-1}^i from $Bel(x_{t-1})$
- draw x_t^i from $p(x_t | x_{t-1}^i, u_{t-1})$
- Importance factor for x_t^i :

$$\begin{aligned} w_t^i &= \frac{\text{target distribution}}{\text{proposal distribution}} \\ &= \frac{\eta p(z_t | x_t) p(x_t | x_{t-1}, u_{t-1}) Bel(x_{t-1})}{p(x_t | x_{t-1}, u_{t-1}) Bel(x_{t-1})} \\ &\propto p(z_t | x_t) \end{aligned}$$

Particle Filter

Particle Filter for Localization

Particle_filter($\mathcal{X}_{t-1}, u_t, z_t$):

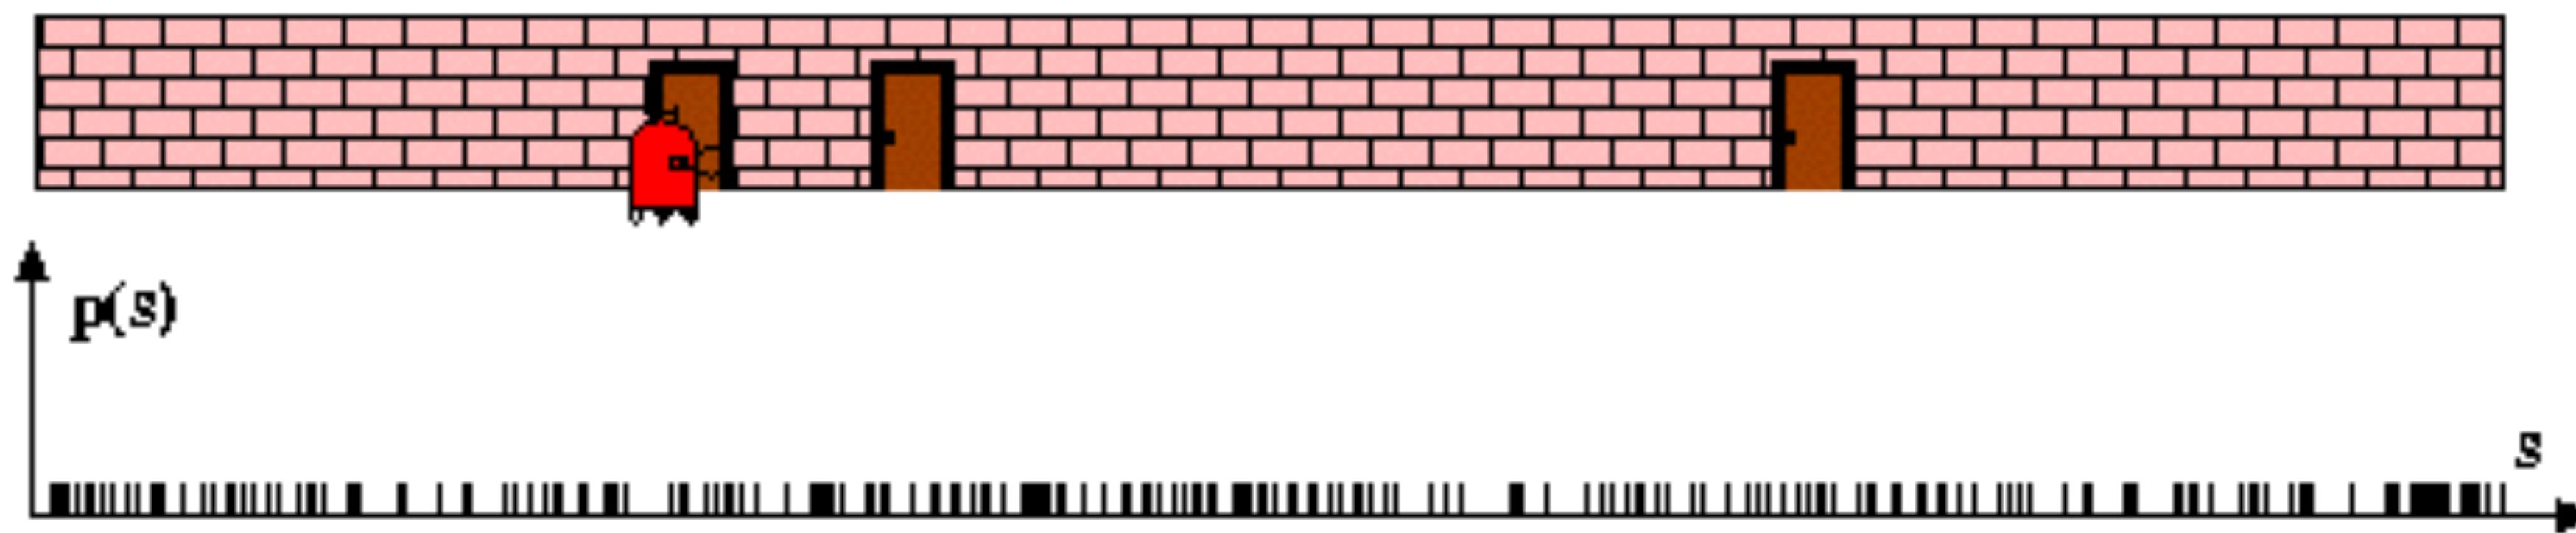
```
1:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
2:   for  $j = 1$  to  $J$  do
3:     sample  $x_t^{[j]} \sim \pi(x_t)$ 
4:      $w_t^{[j]} = \frac{p(x_t^{[j]})}{\pi(x_t^{[j]})}$ 
5:      $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[j]}, w_t^{[j]} \rangle$ 
6:   endfor
7:   for  $j = 1$  to  $J$  do
8:     draw  $i \in 1, \dots, J$  with probability  $\propto w_t^{[i]}$ 
9:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
10:  endfor
11:  return  $\mathcal{X}_t$ 
```

Particle_filter($\mathcal{X}_{t-1}, u_t, z_t$):

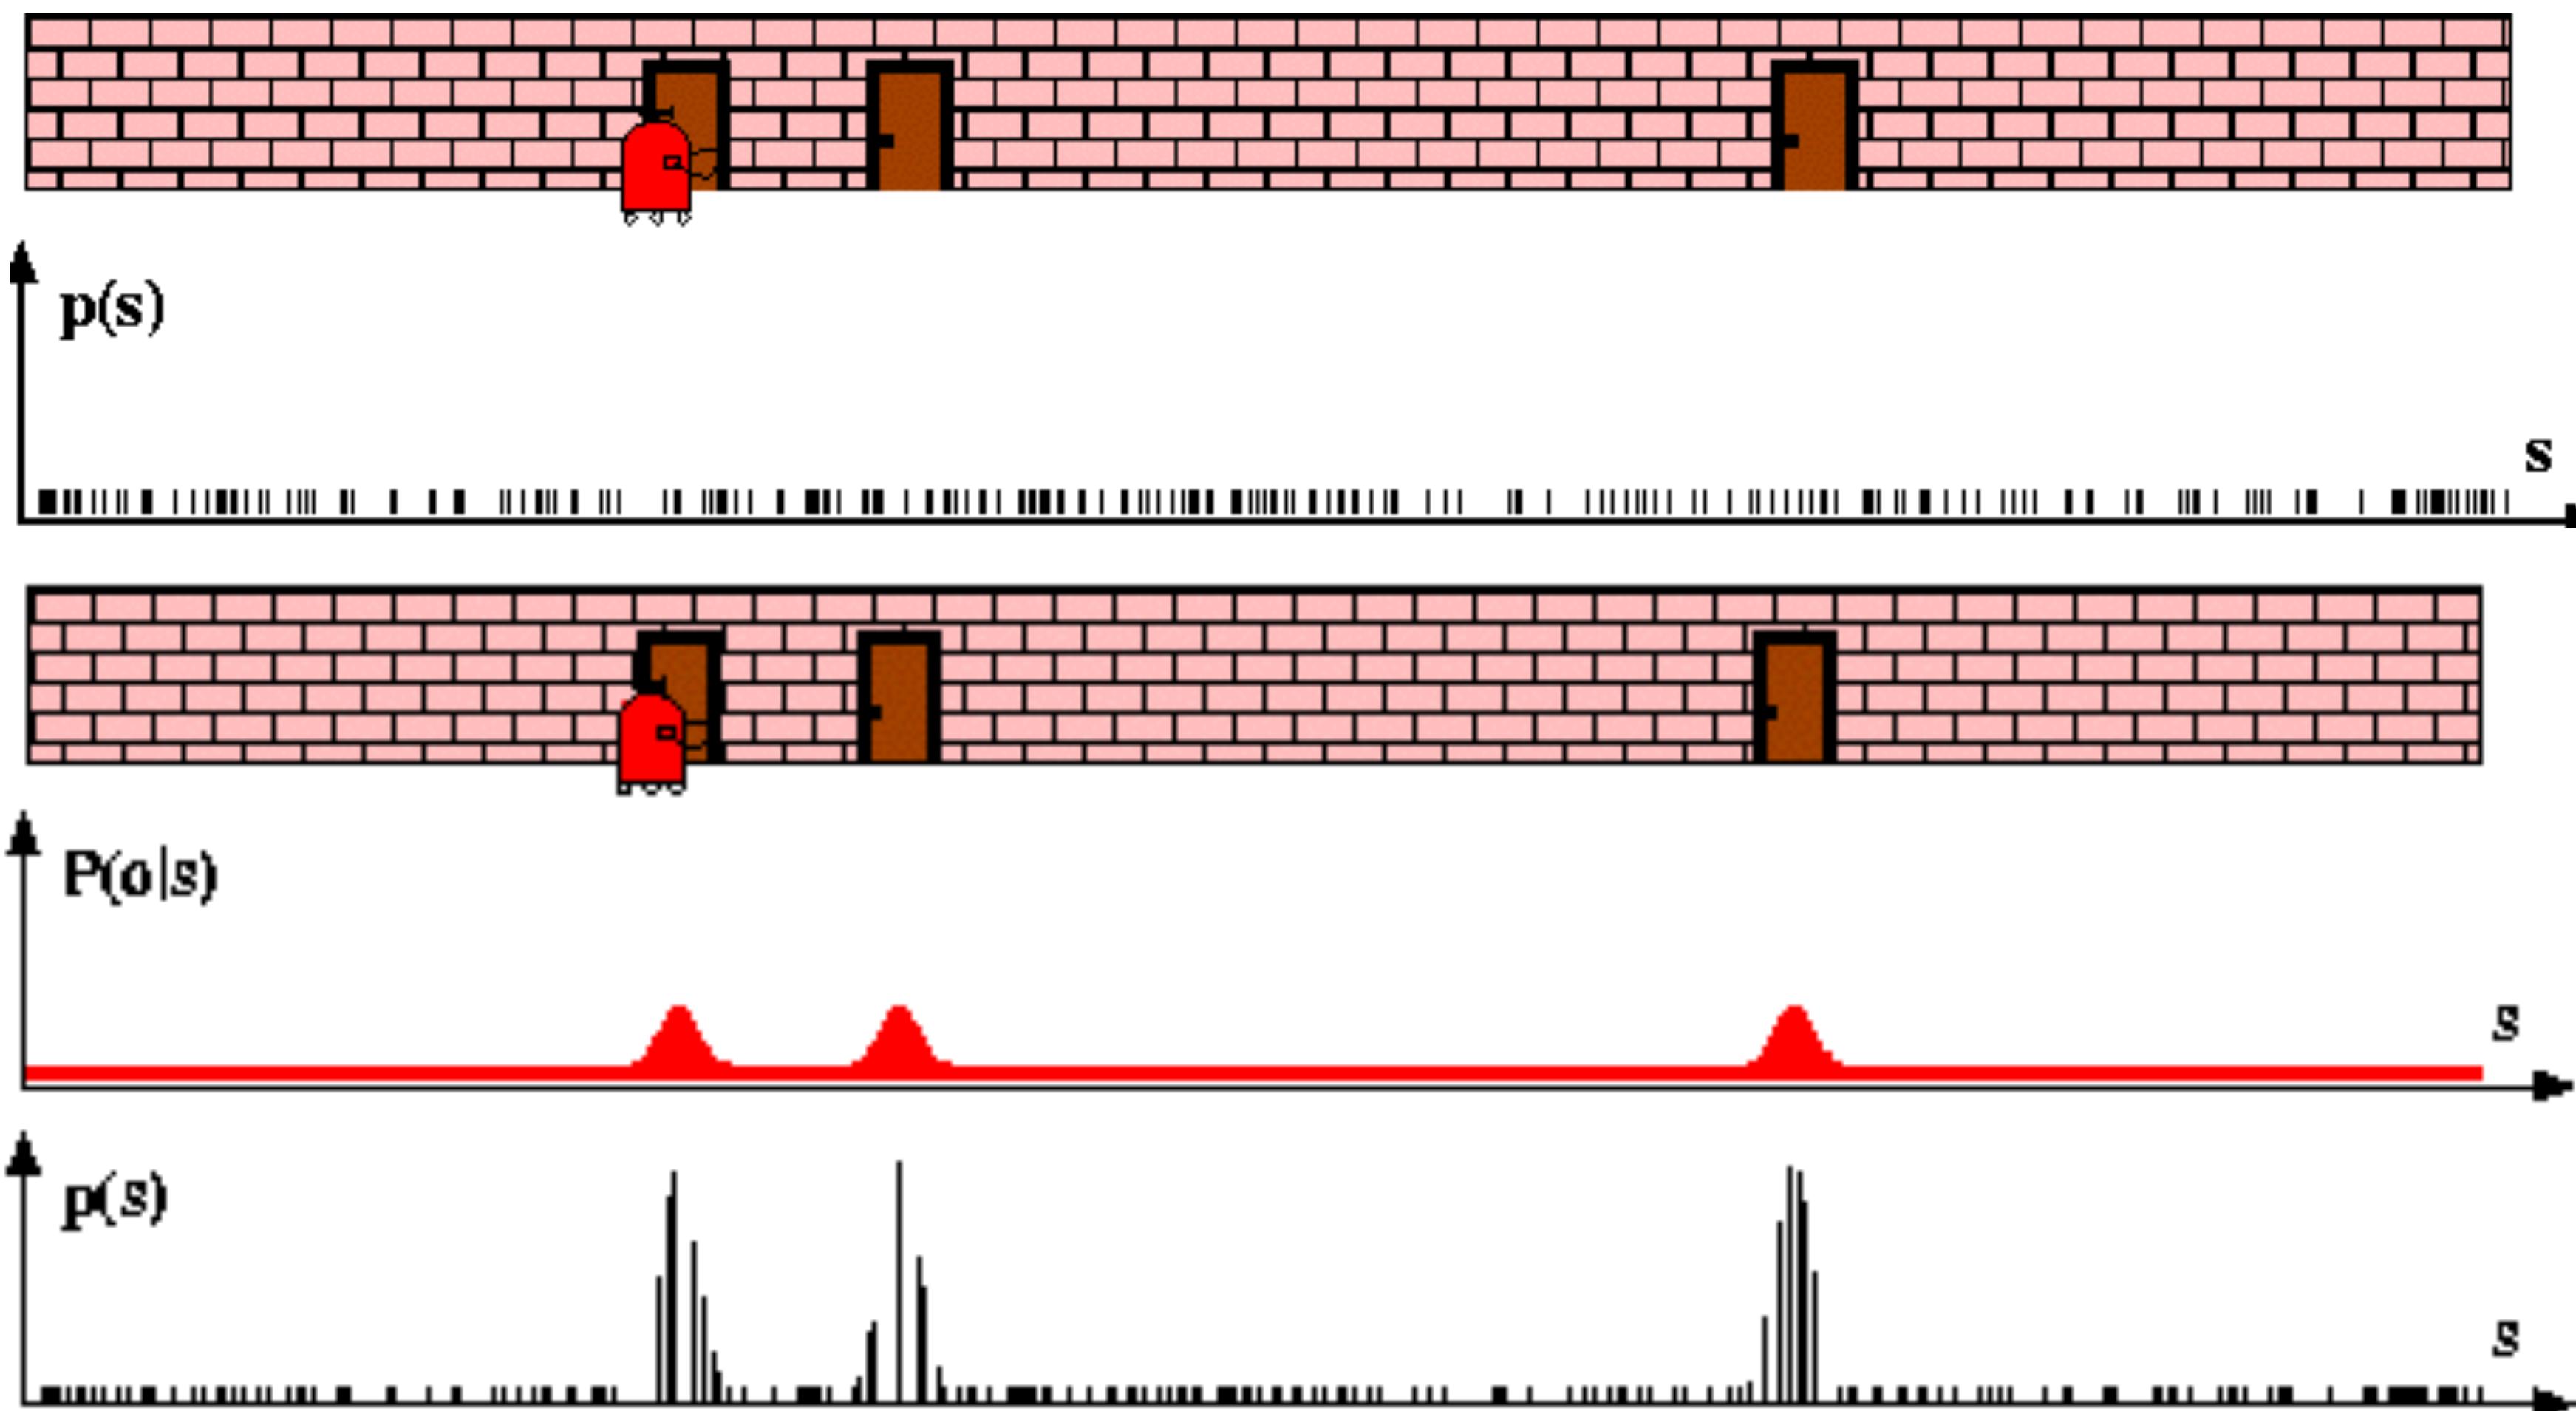
```
1:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
2:   for  $j = 1$  to  $J$  do
3:     sample  $x_t^{[j]} \sim p(x_t | u_t, x_{t-1}^{[j]})$ 
4:      $w_t^{[j]} = p(z_t | x_t^{[j]})$ 
5:      $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[j]}, w_t^{[j]} \rangle$ 
6:   endfor
7:   for  $j = 1$  to  $J$  do
8:     draw  $i \in 1, \dots, J$  with probability  $\propto w_t^{[i]}$ 
9:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
10:  endfor
11:  return  $\mathcal{X}_t$ 
```



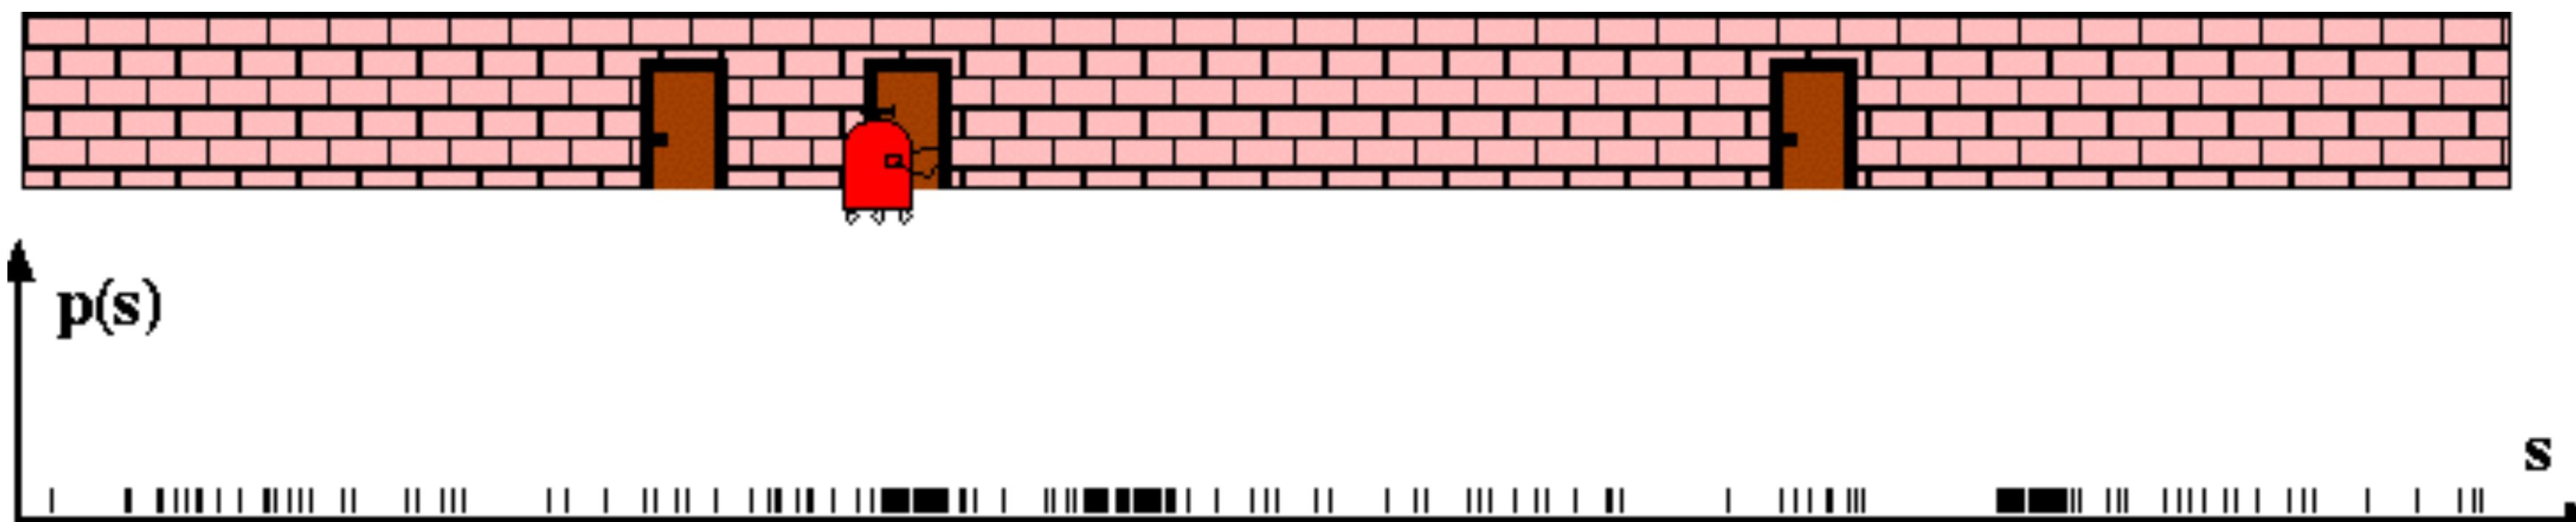
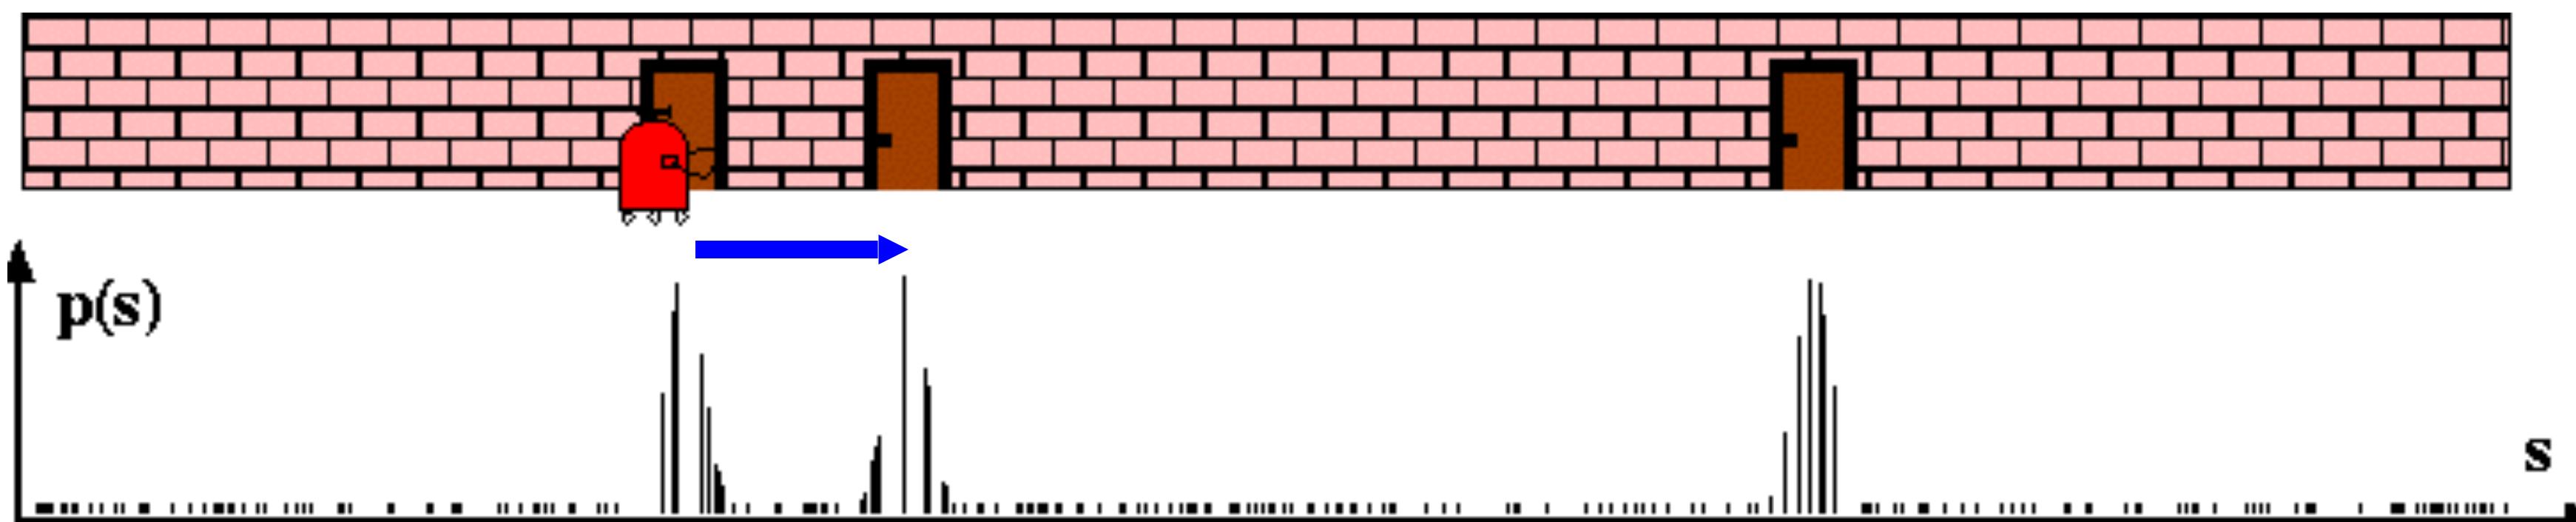
Particle Filters



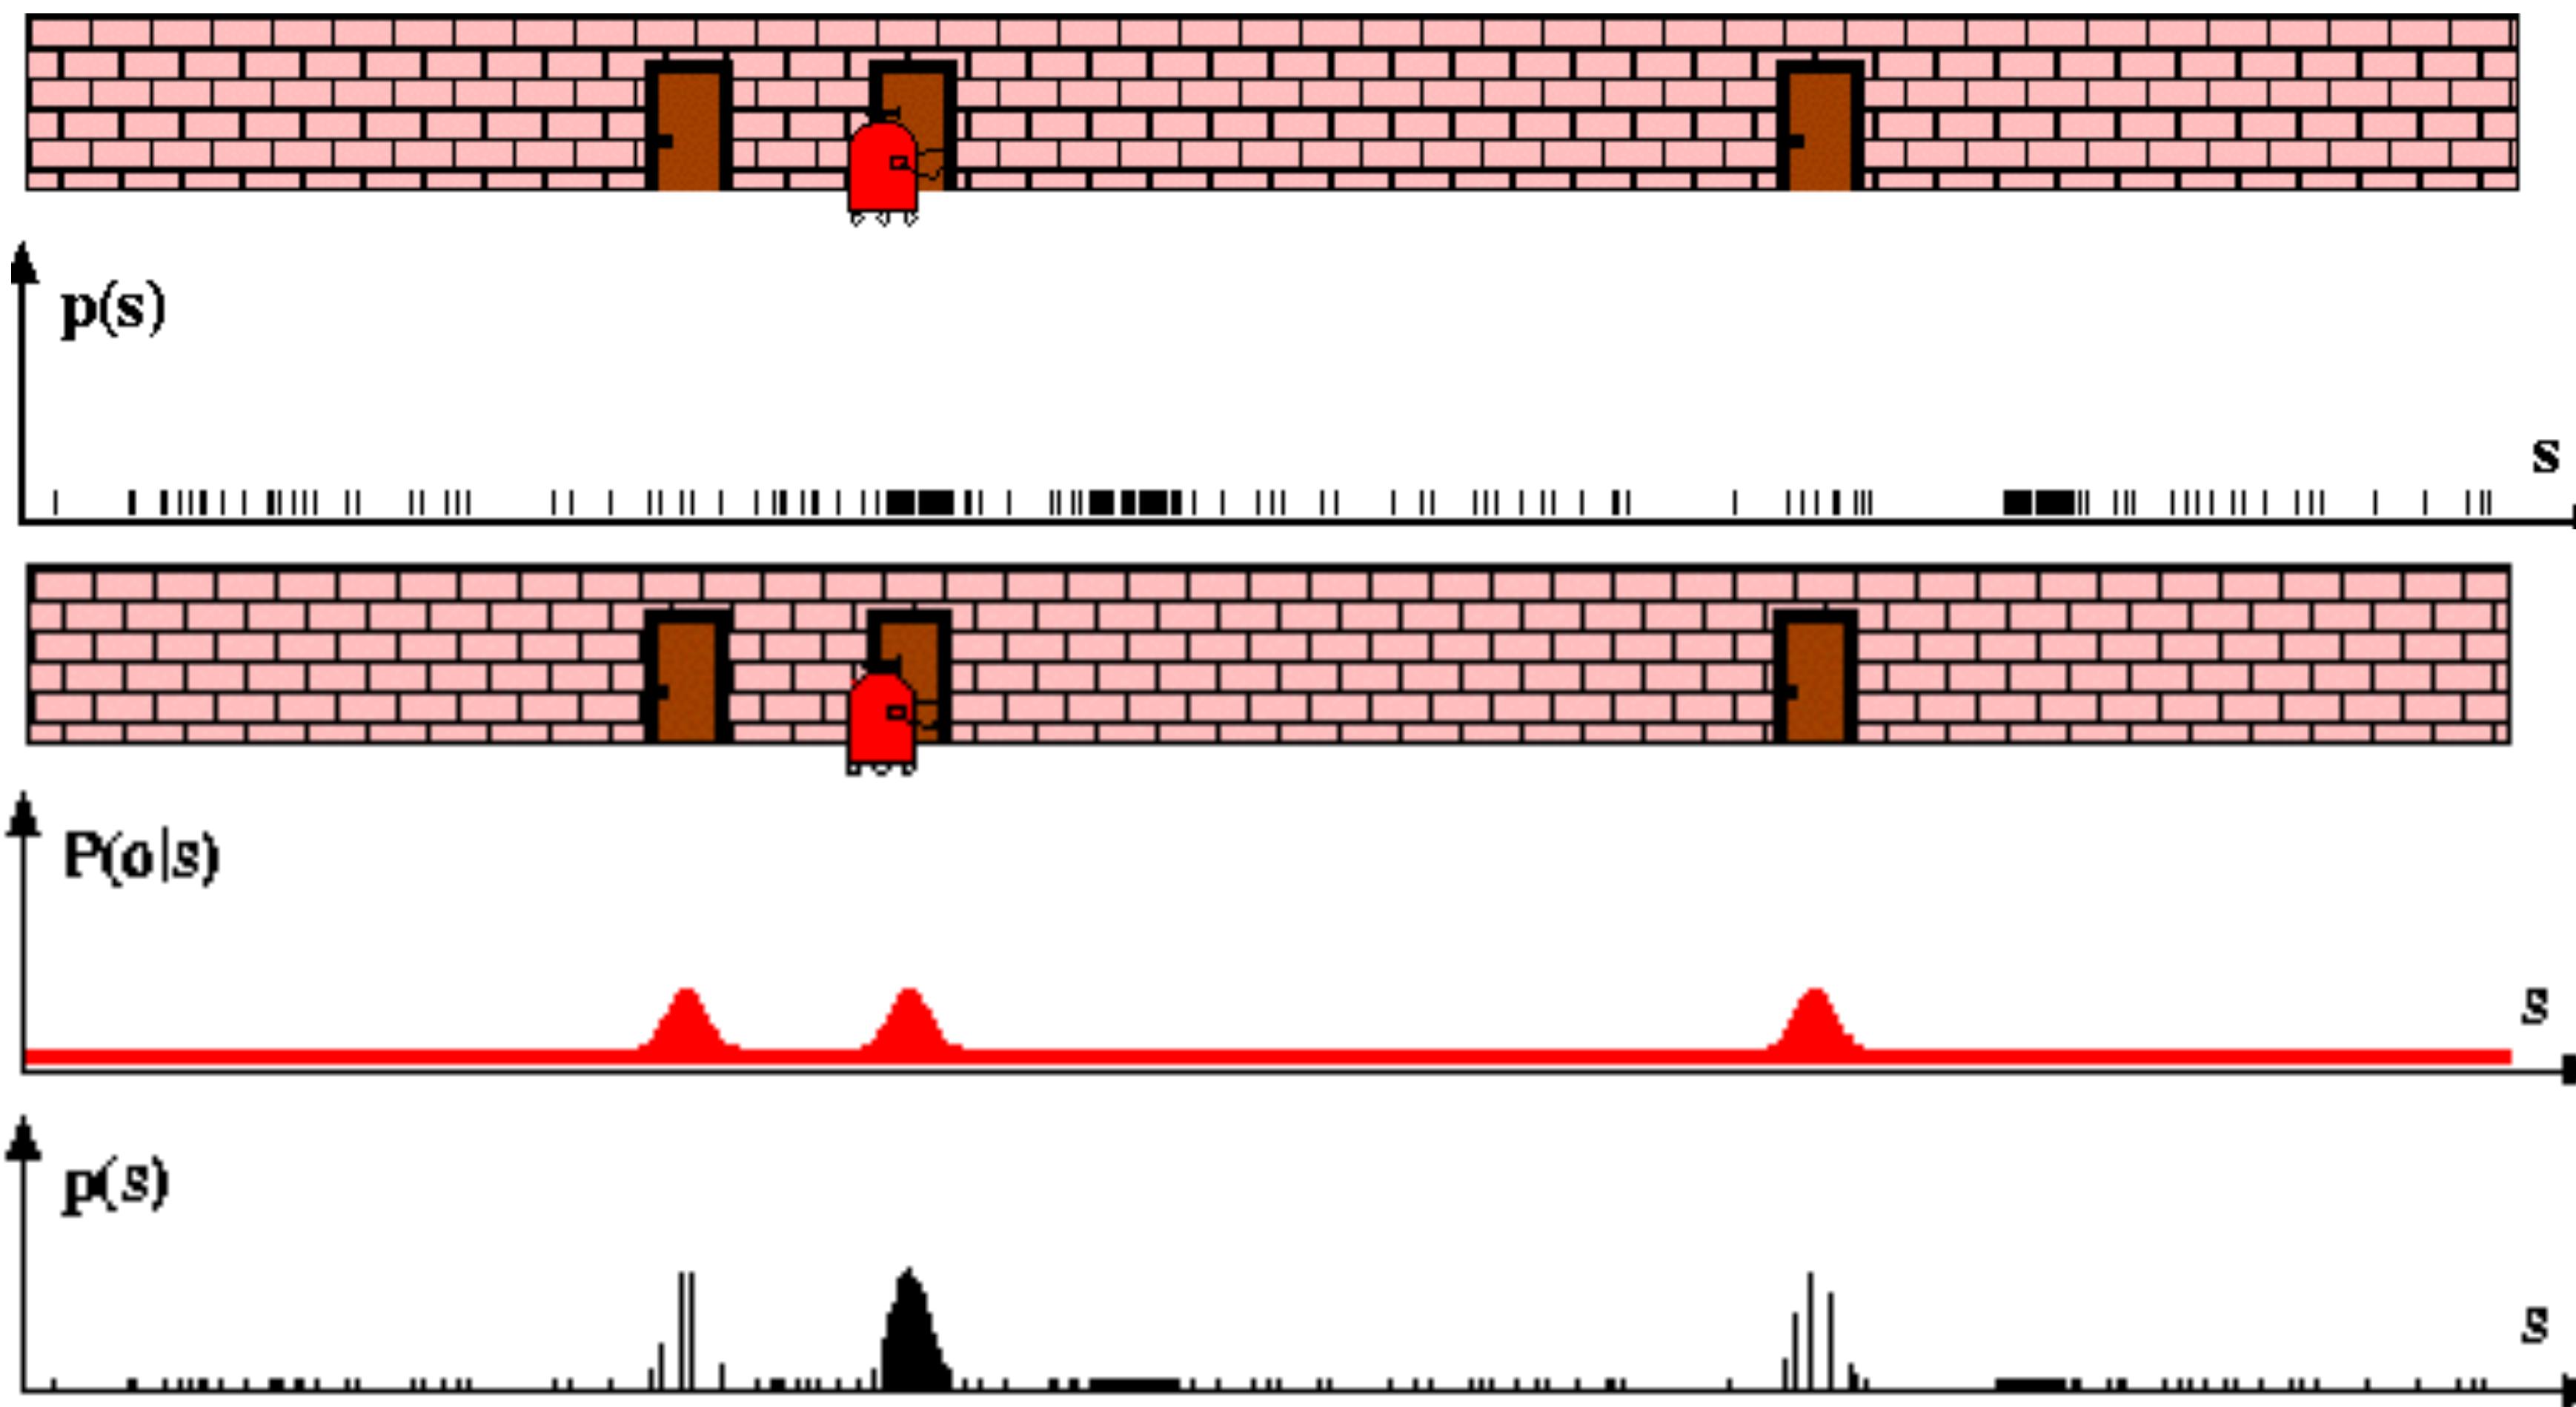
Sensor Information: Importance Sampling



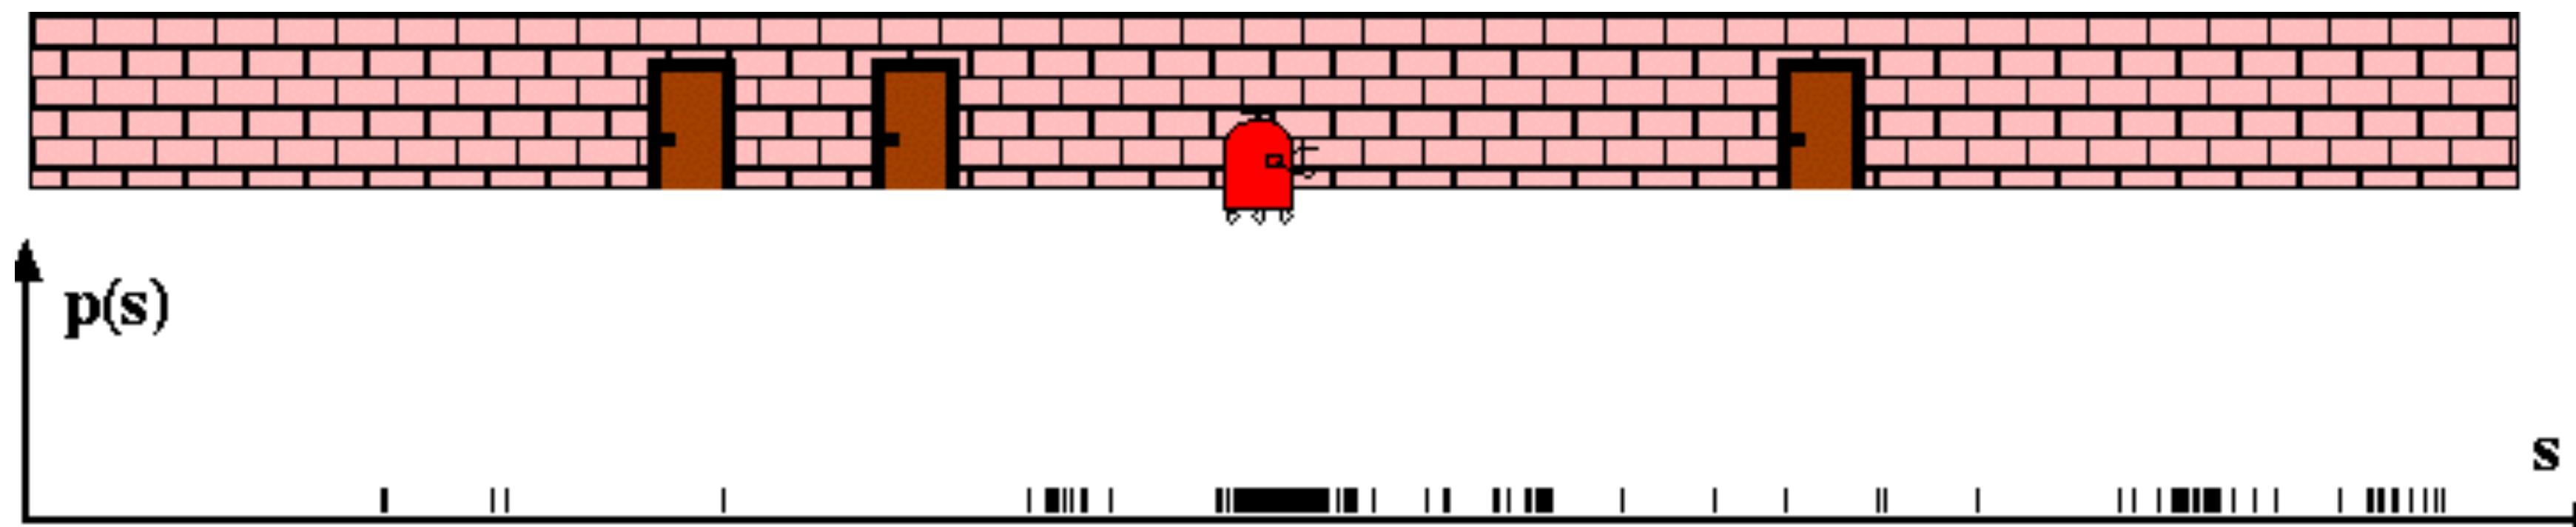
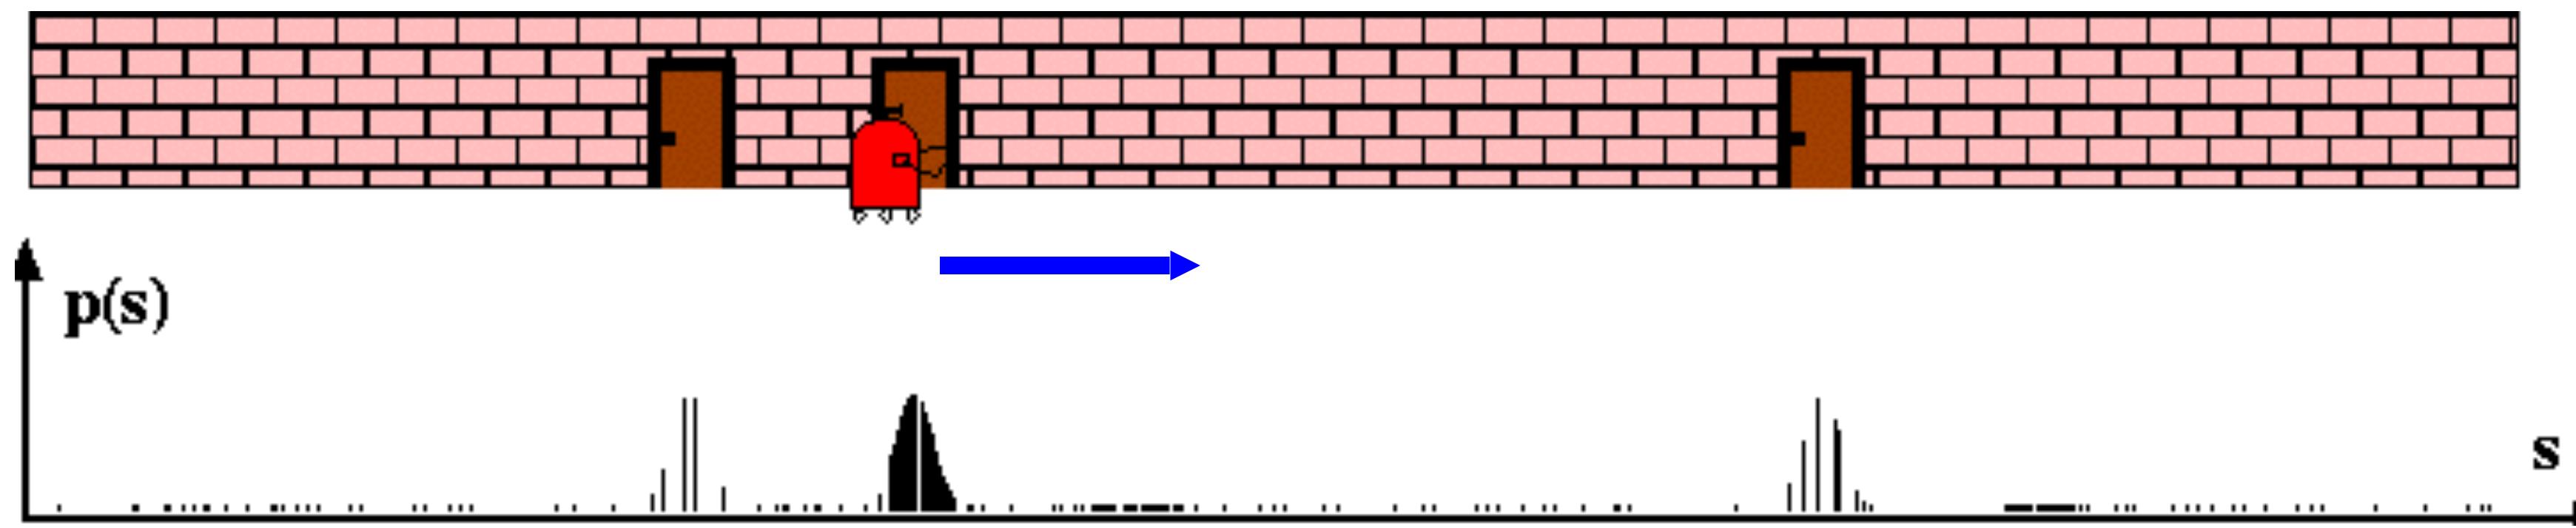
Robot Motion

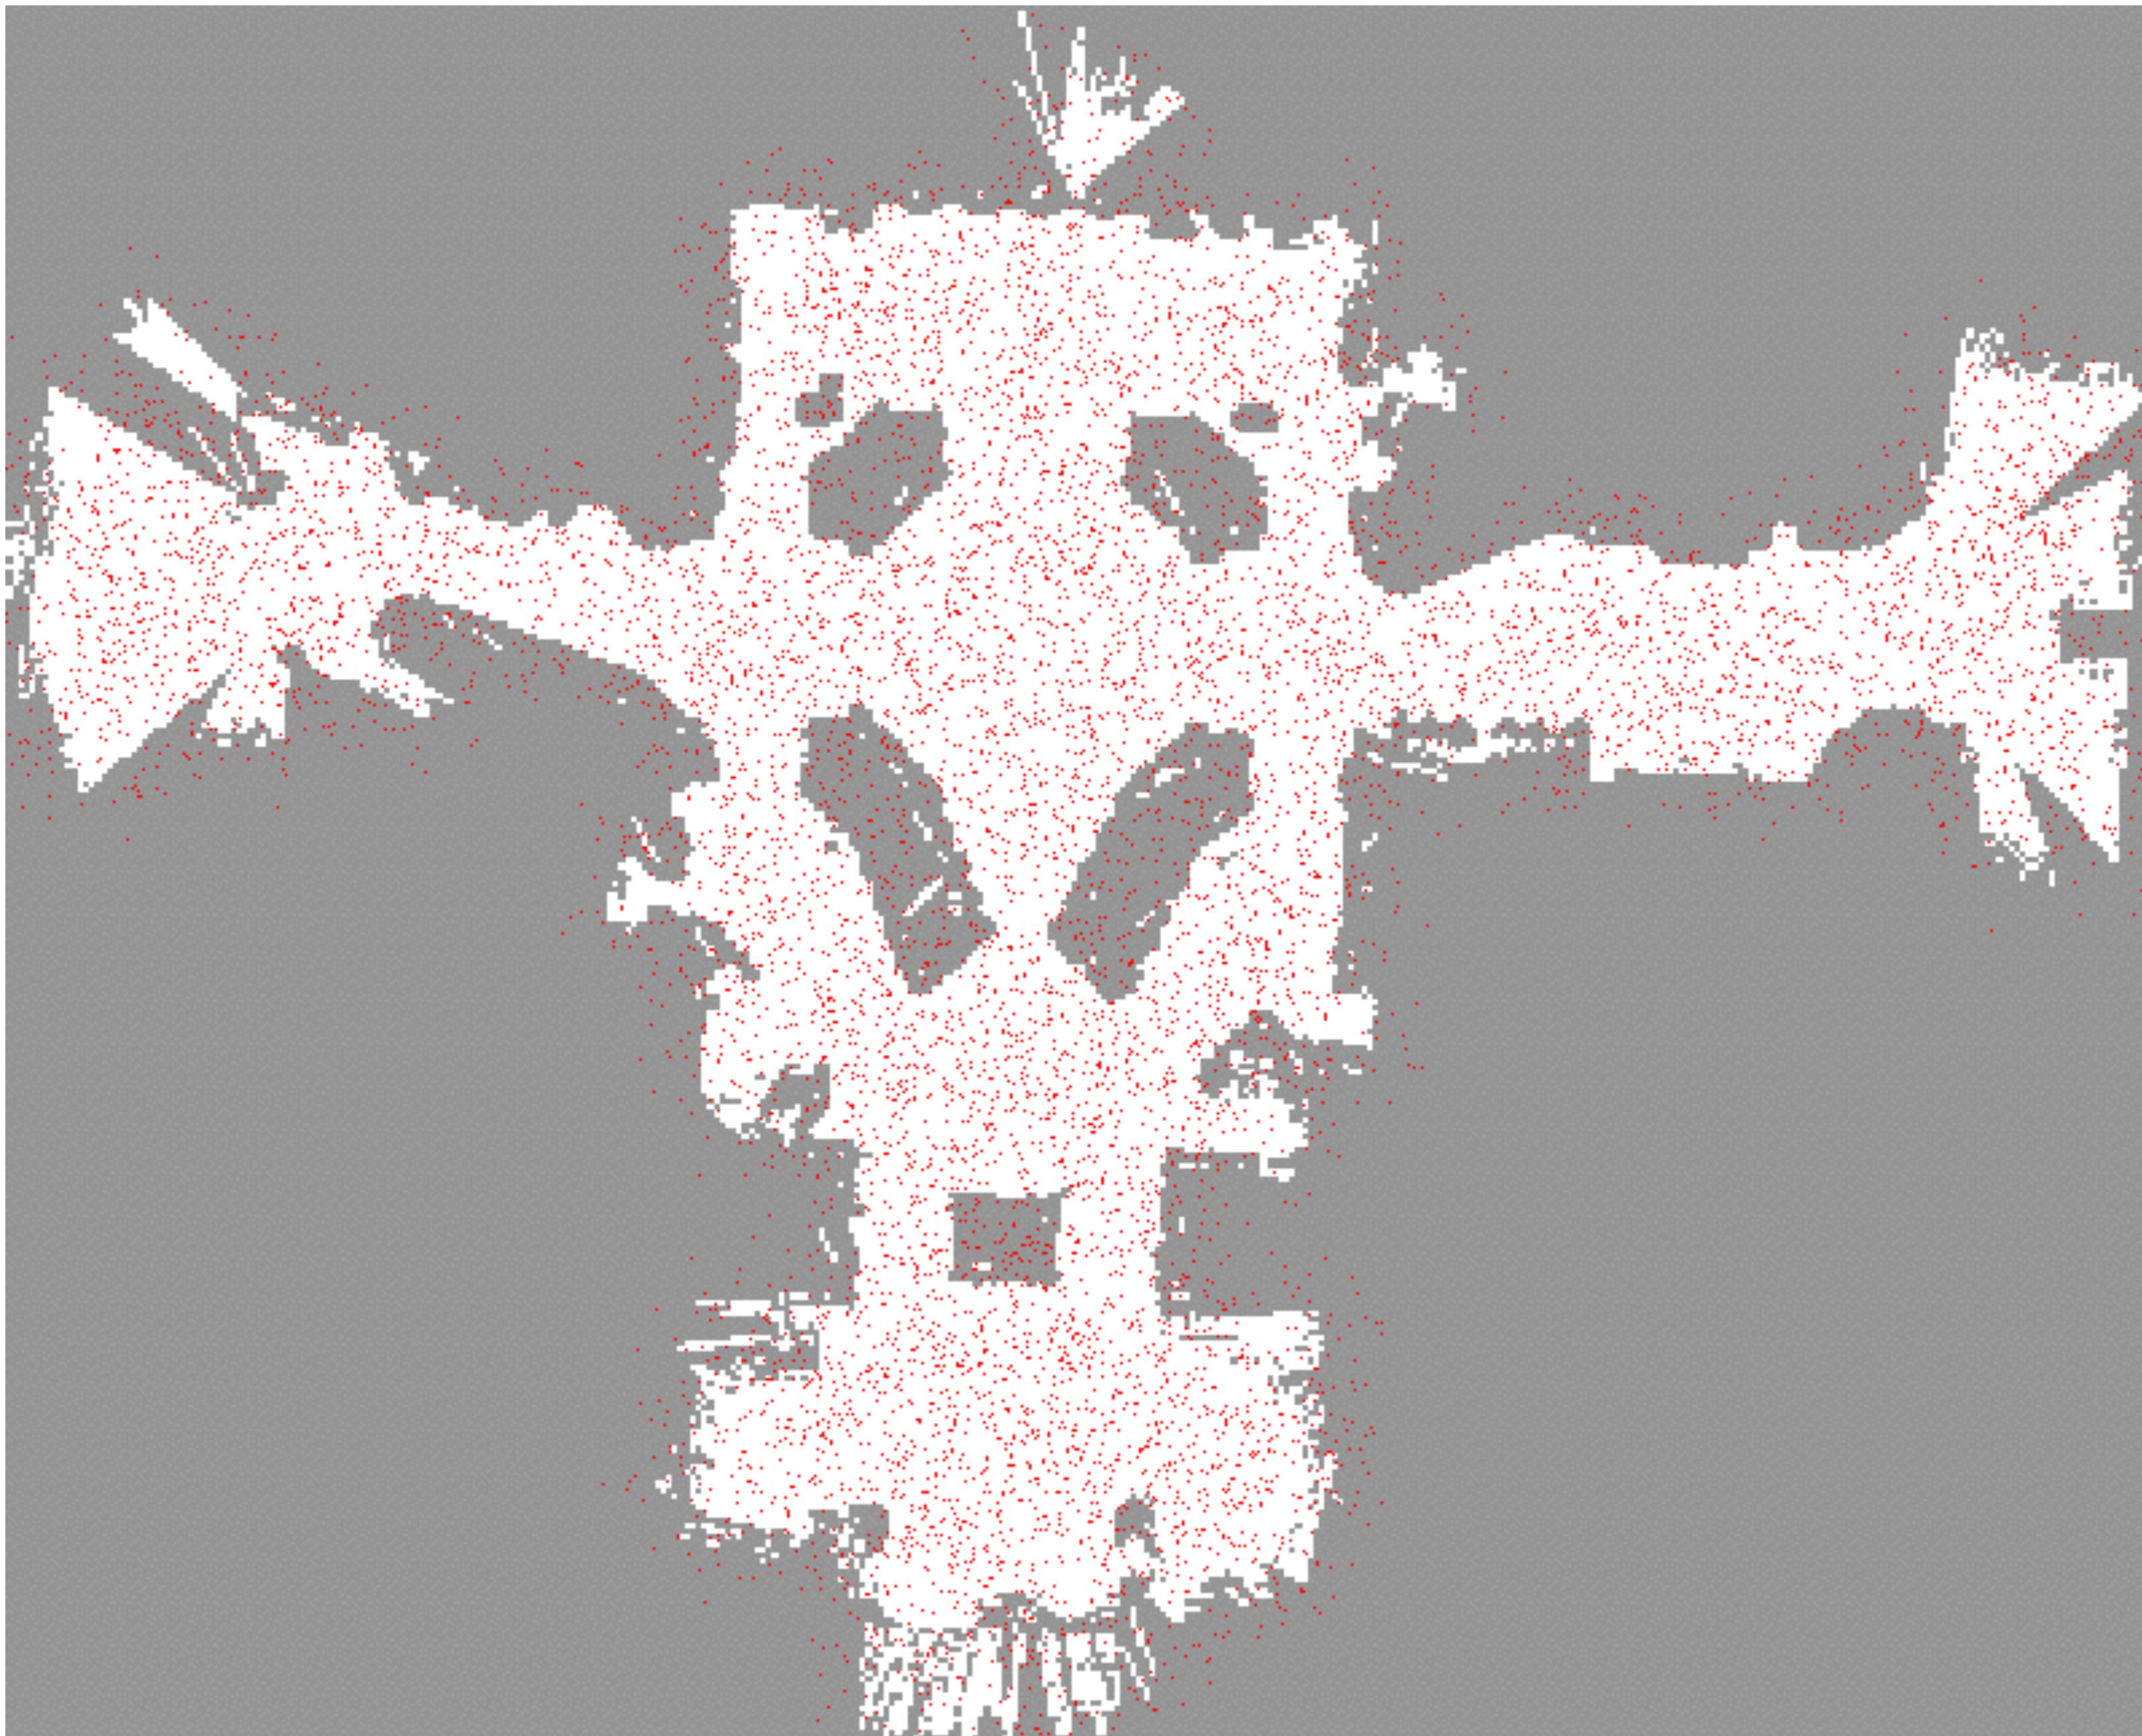


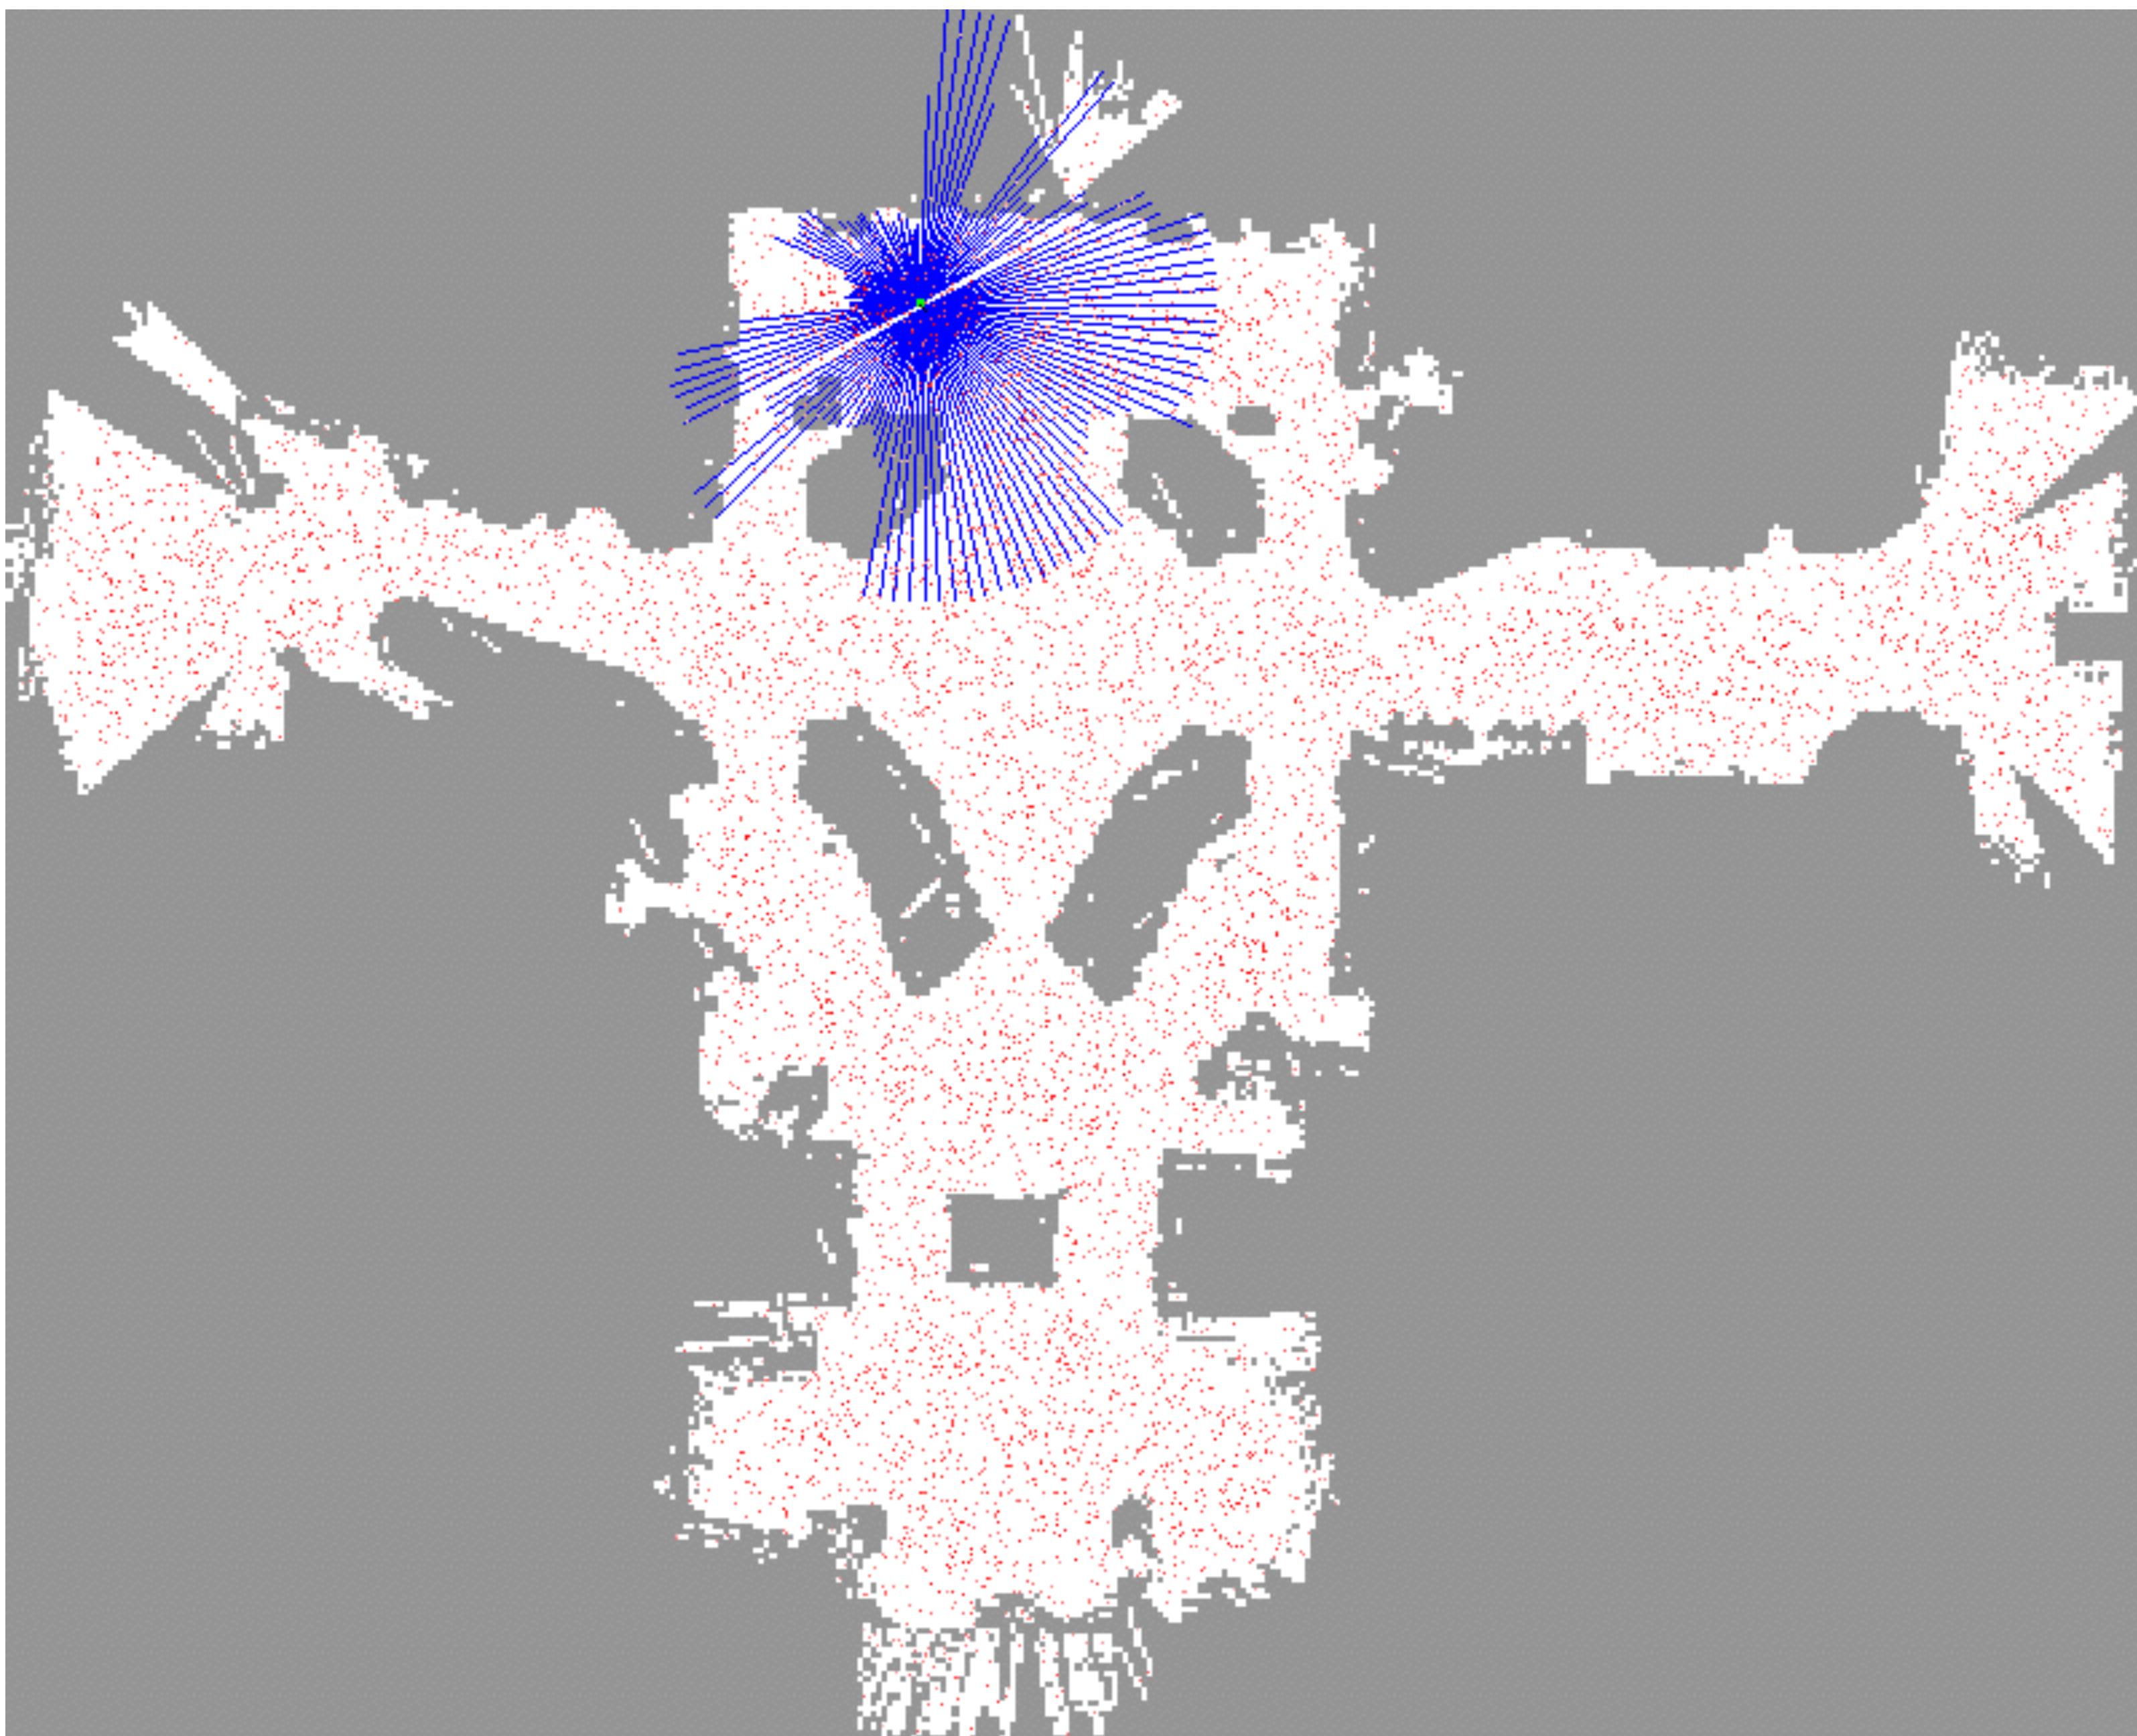
Sensor Information: Importance Sampling

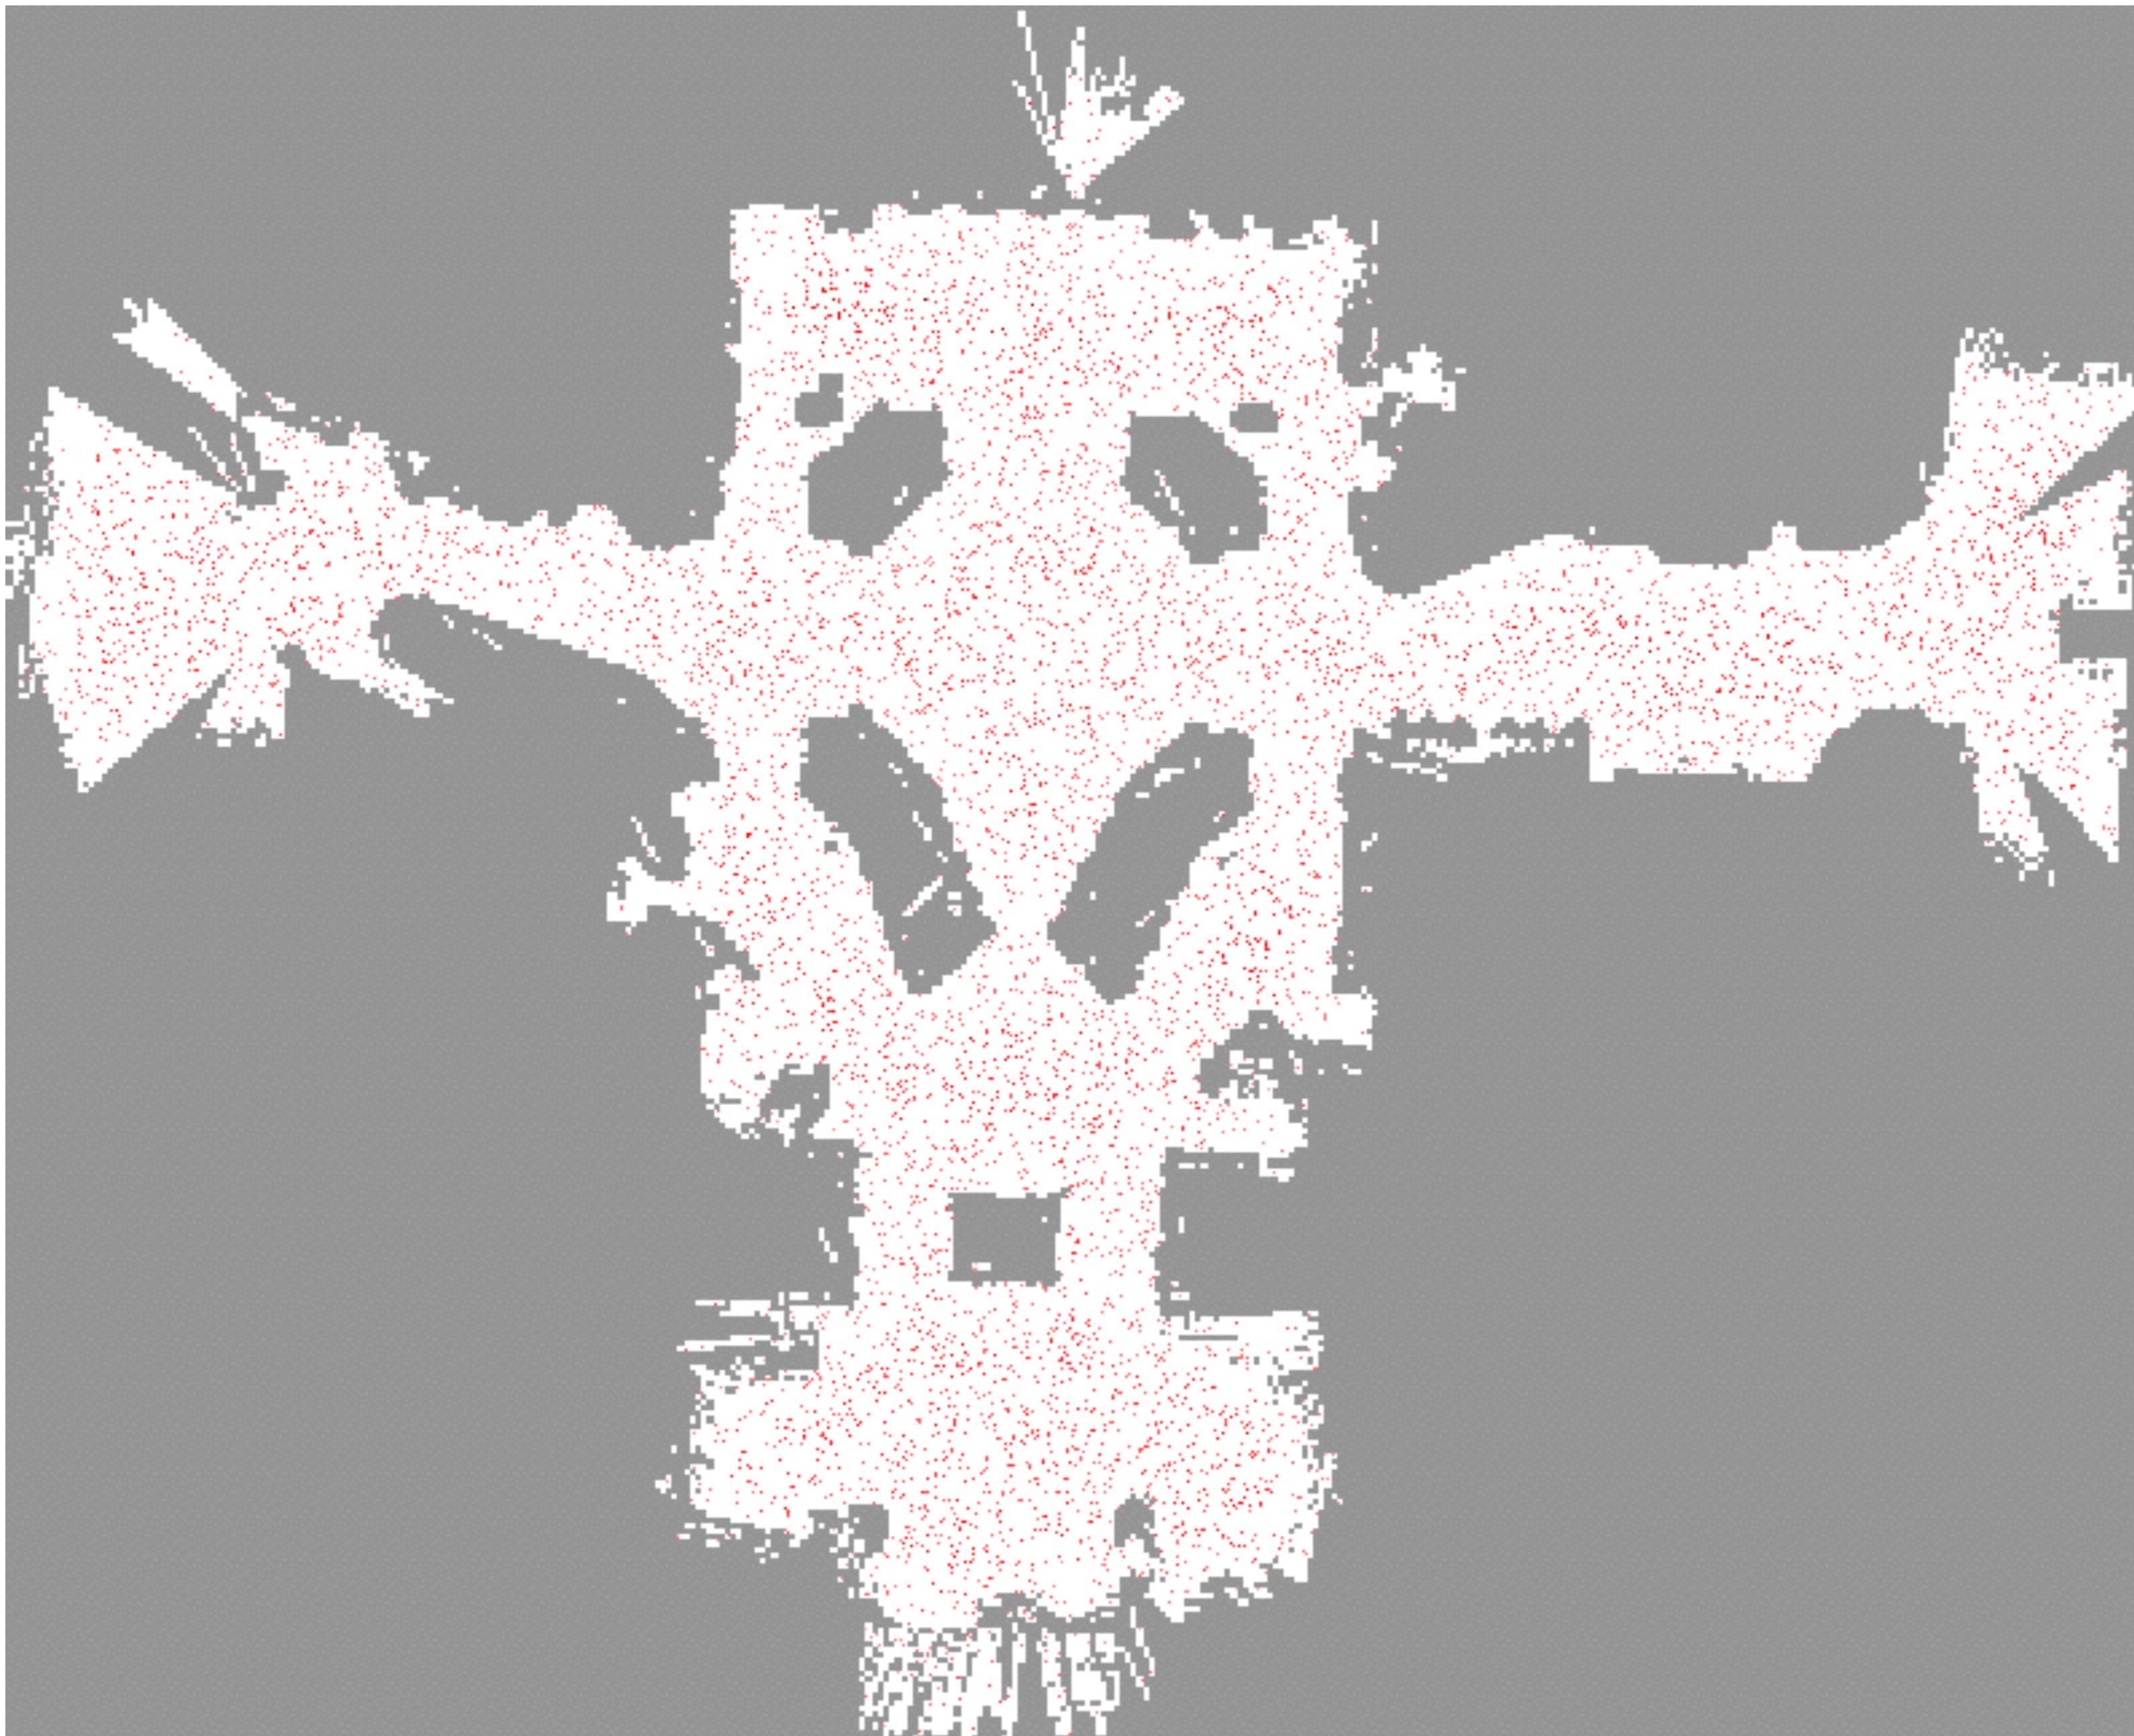


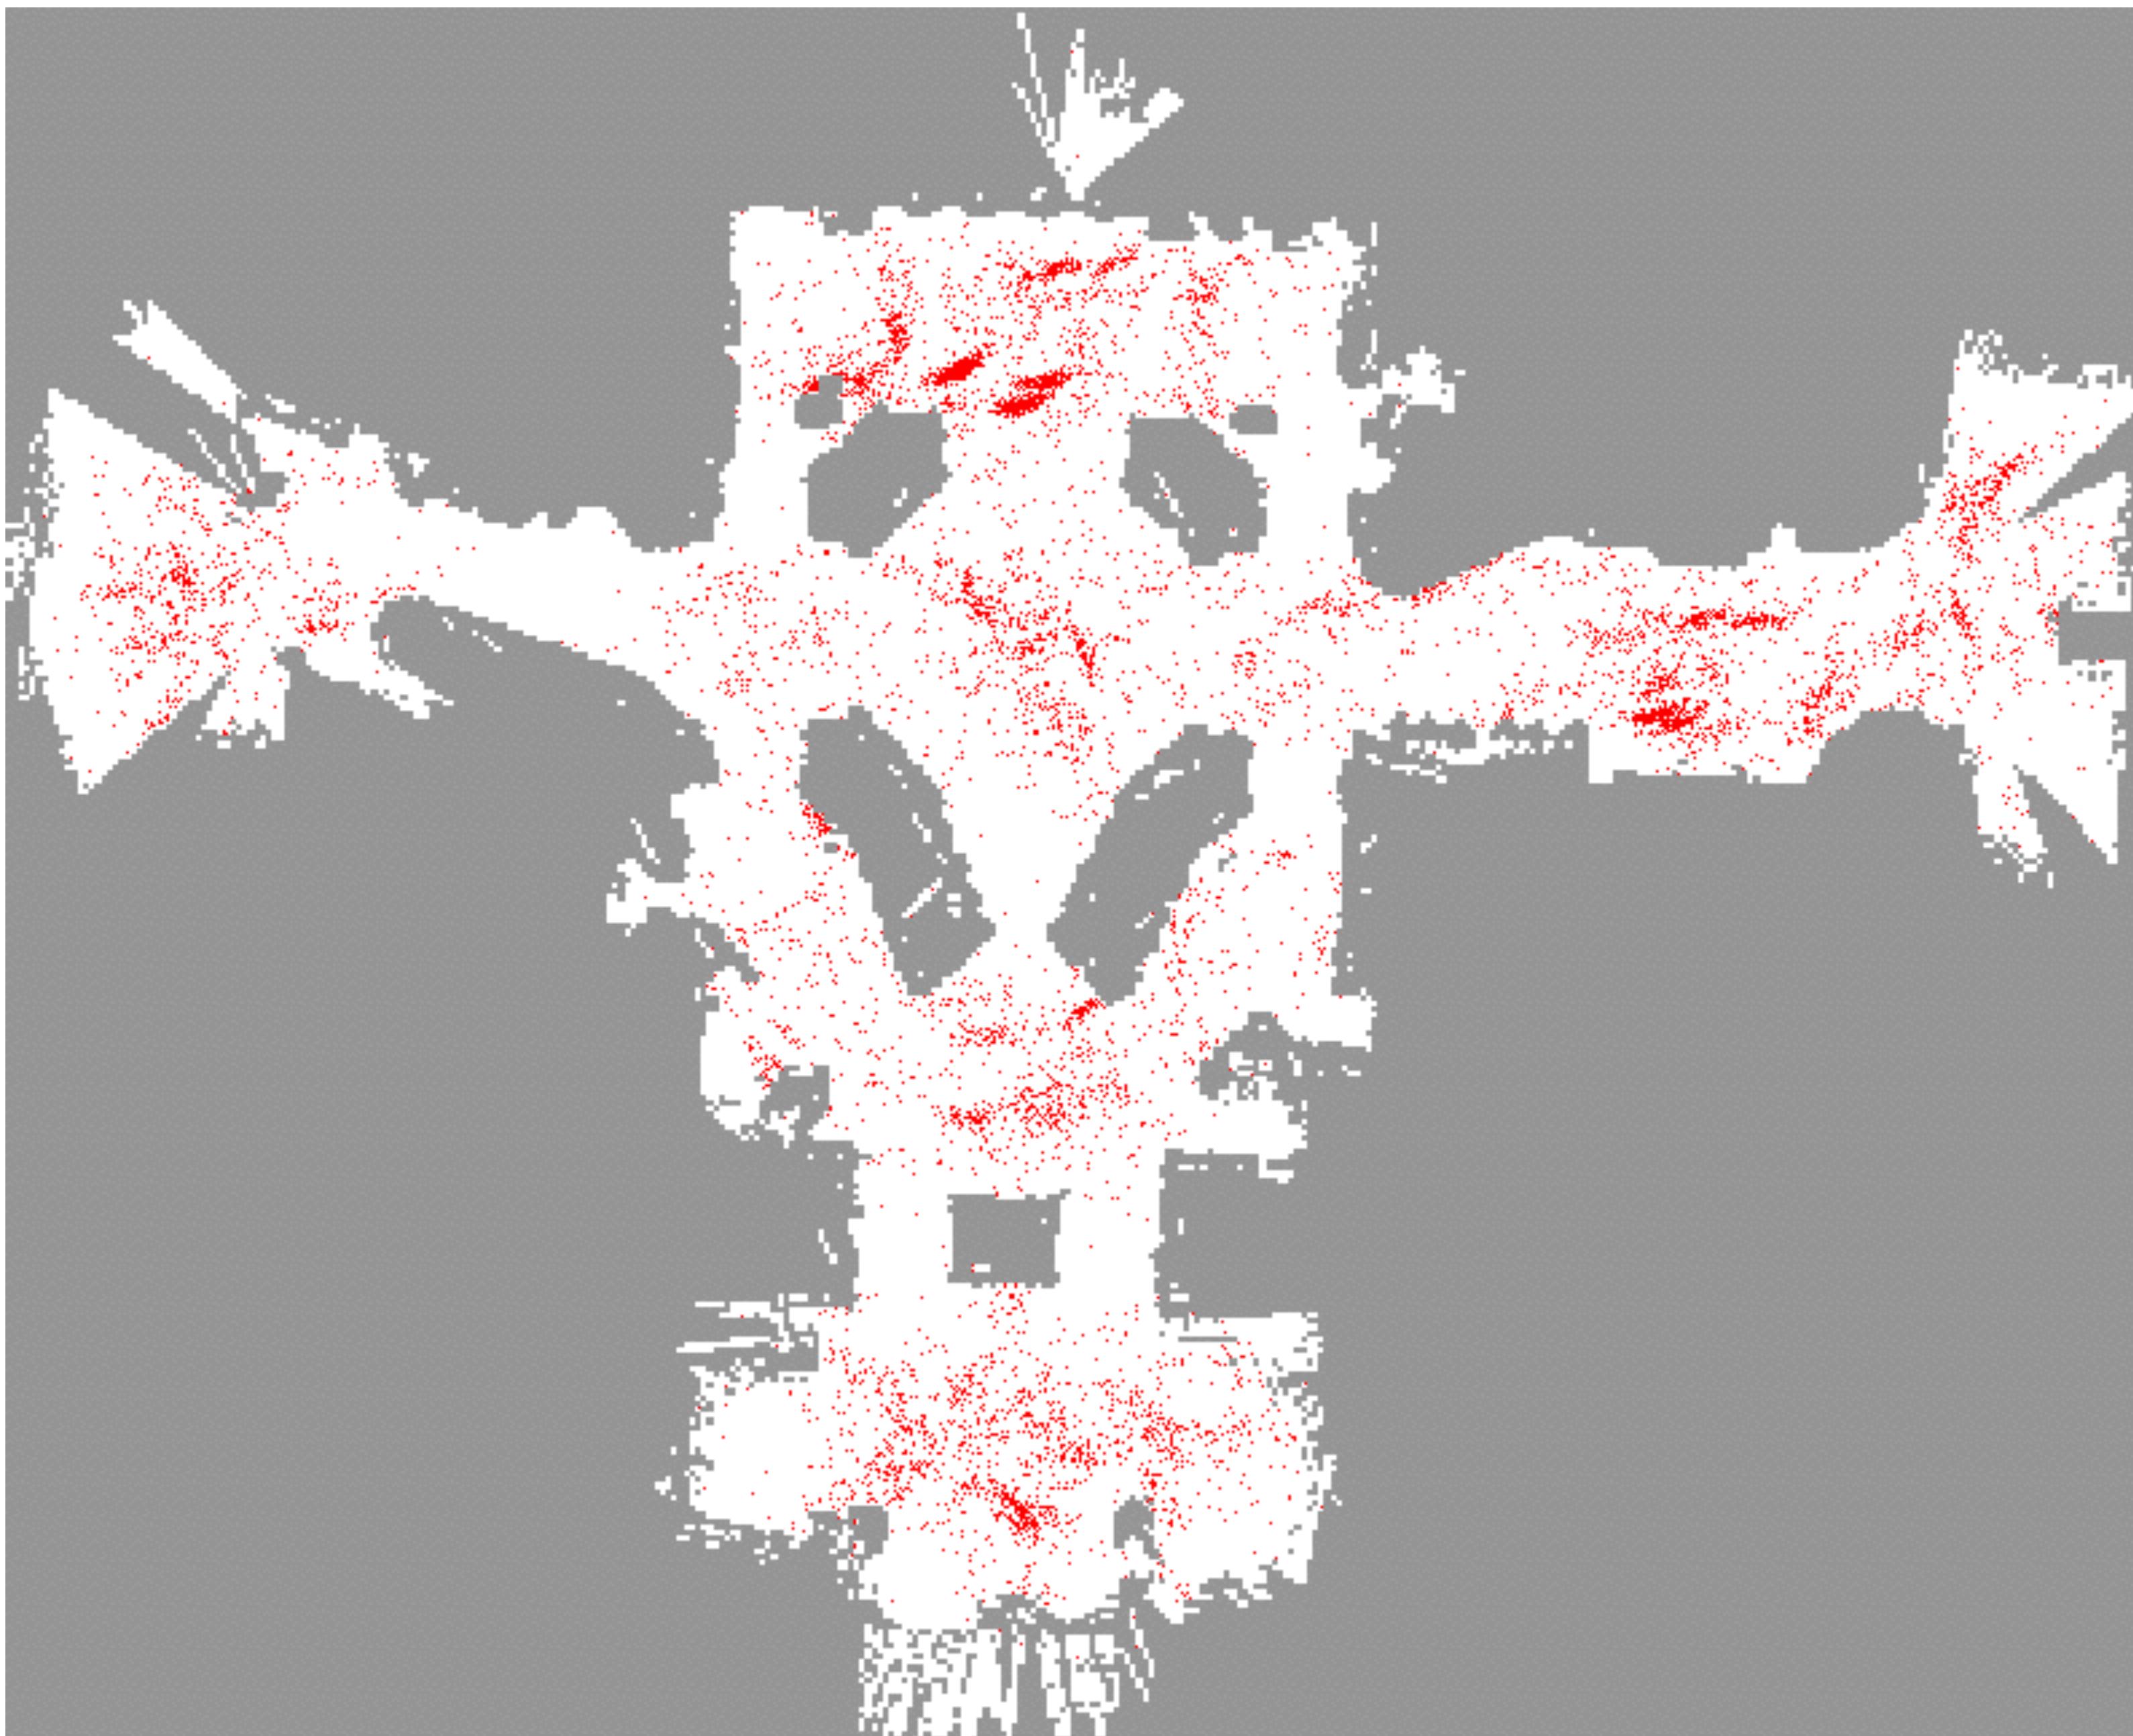
Robot Motion

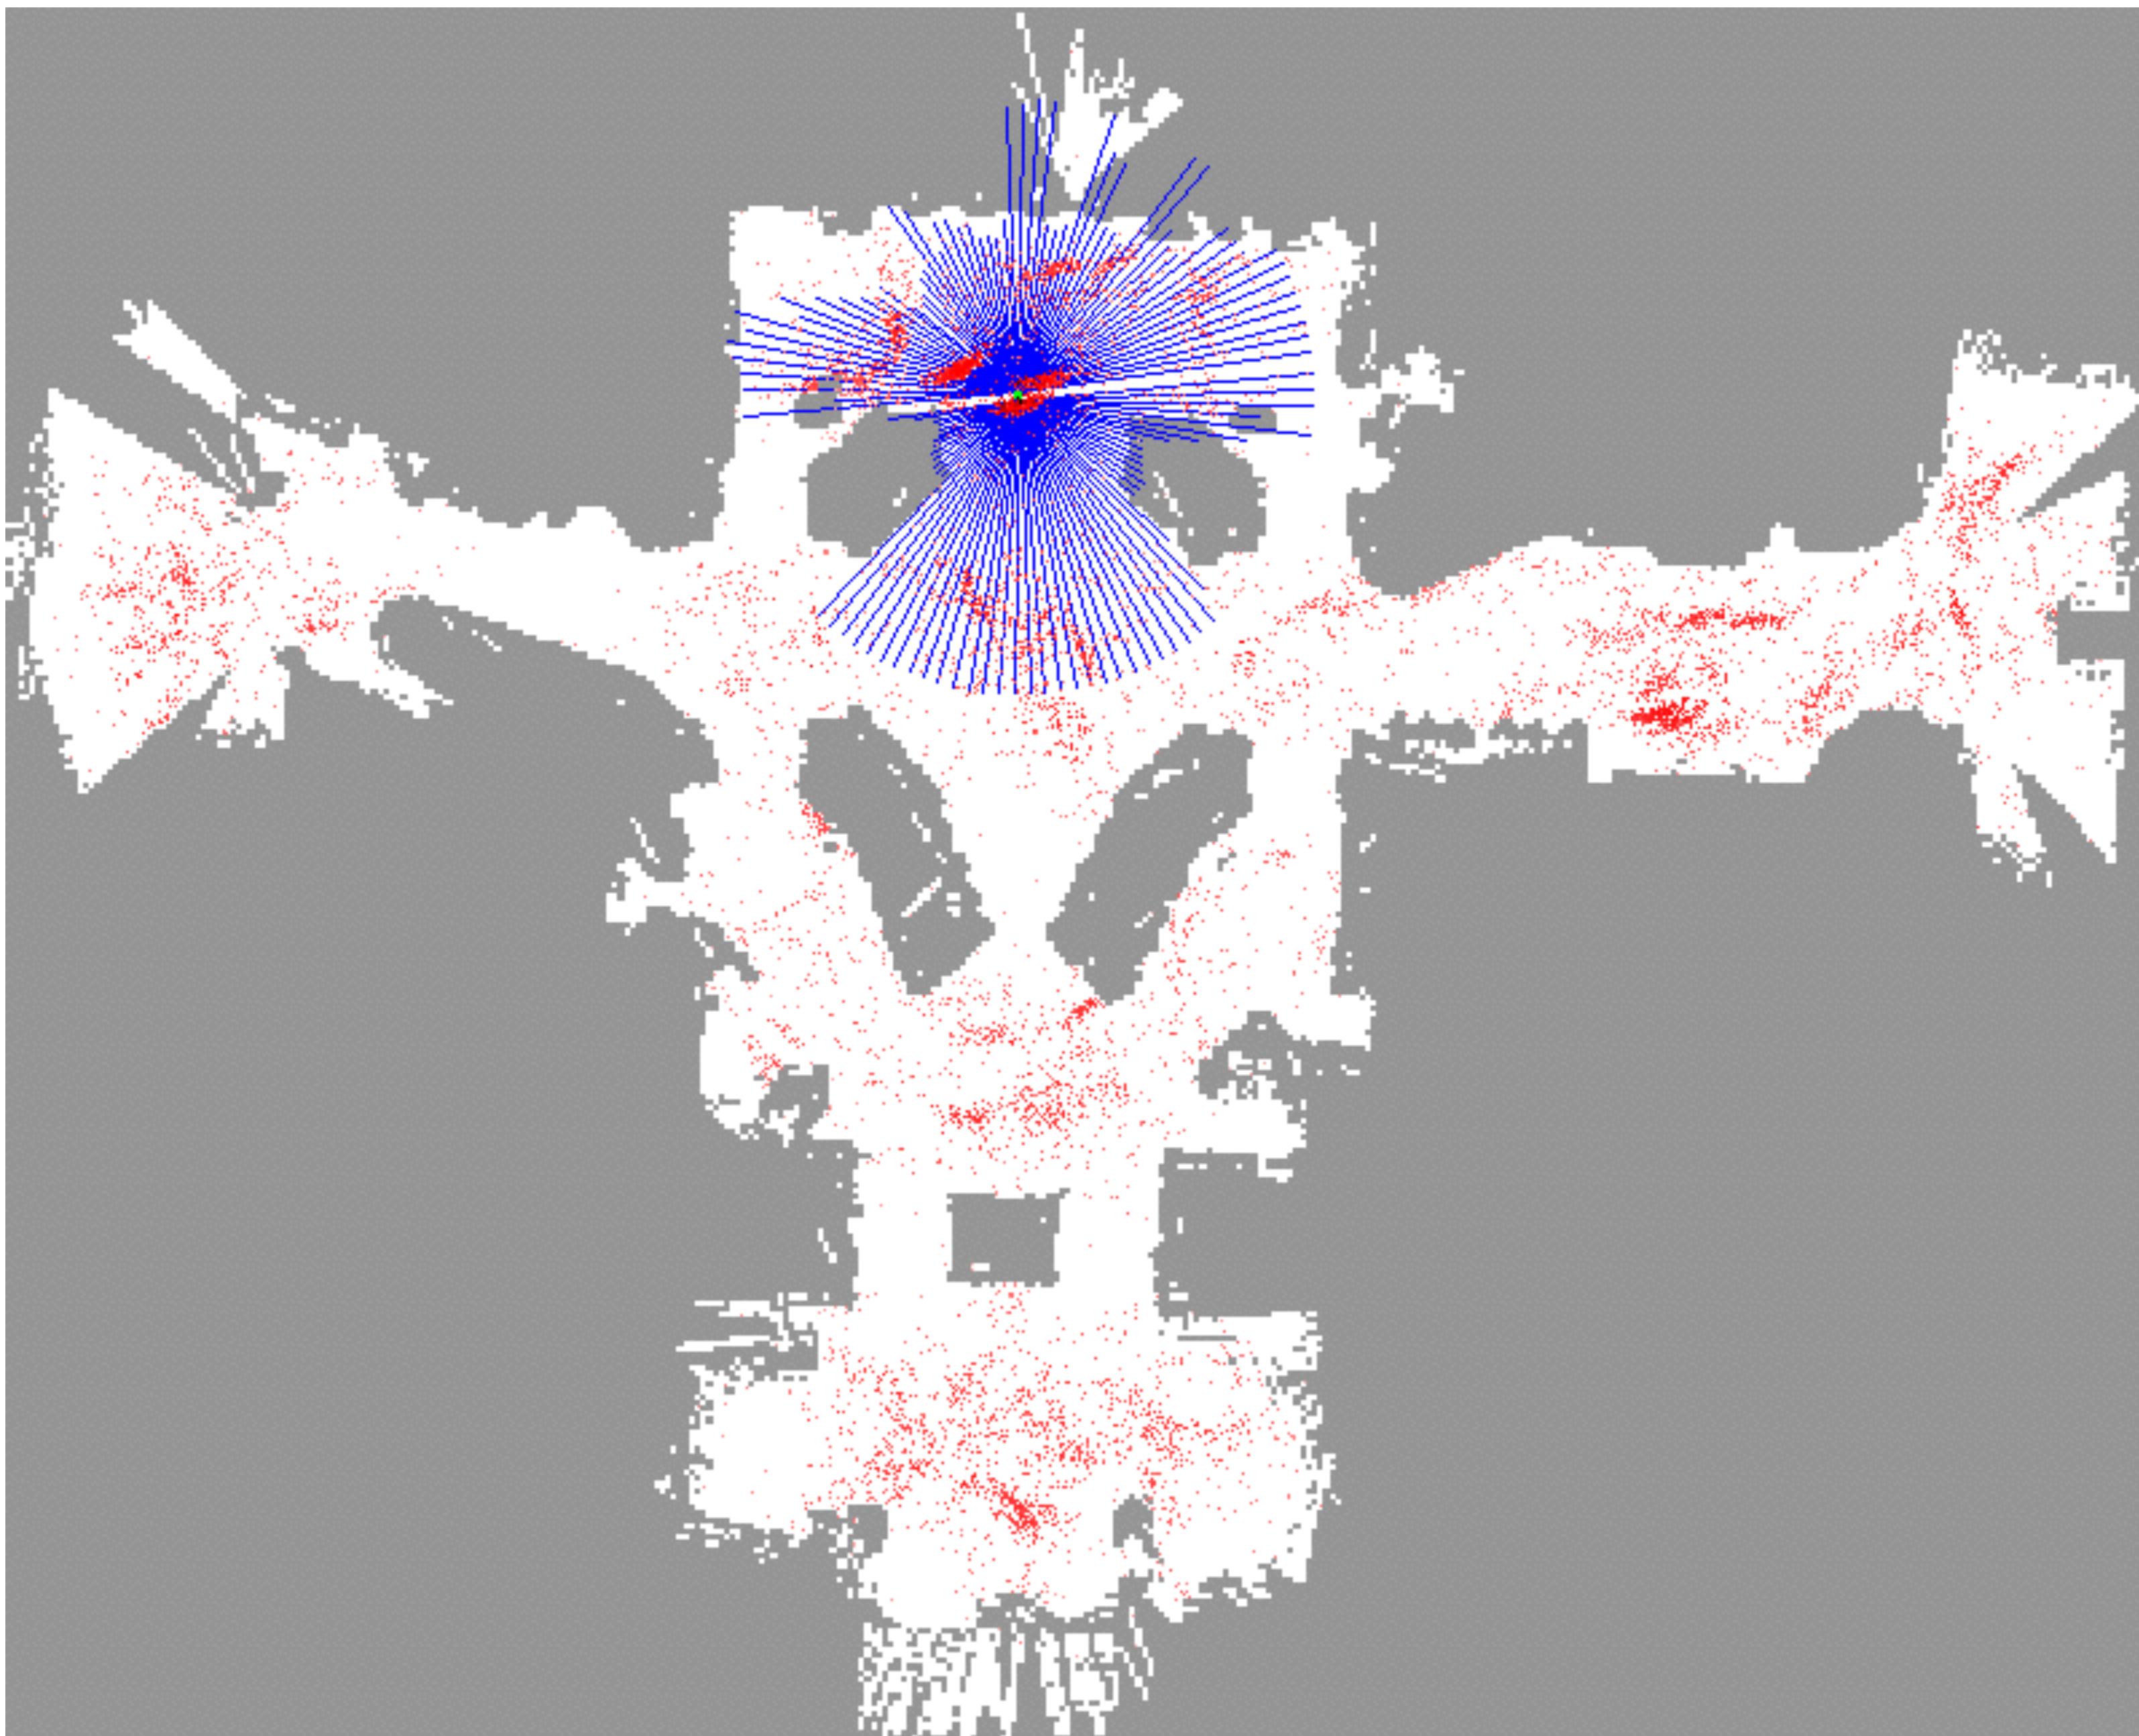


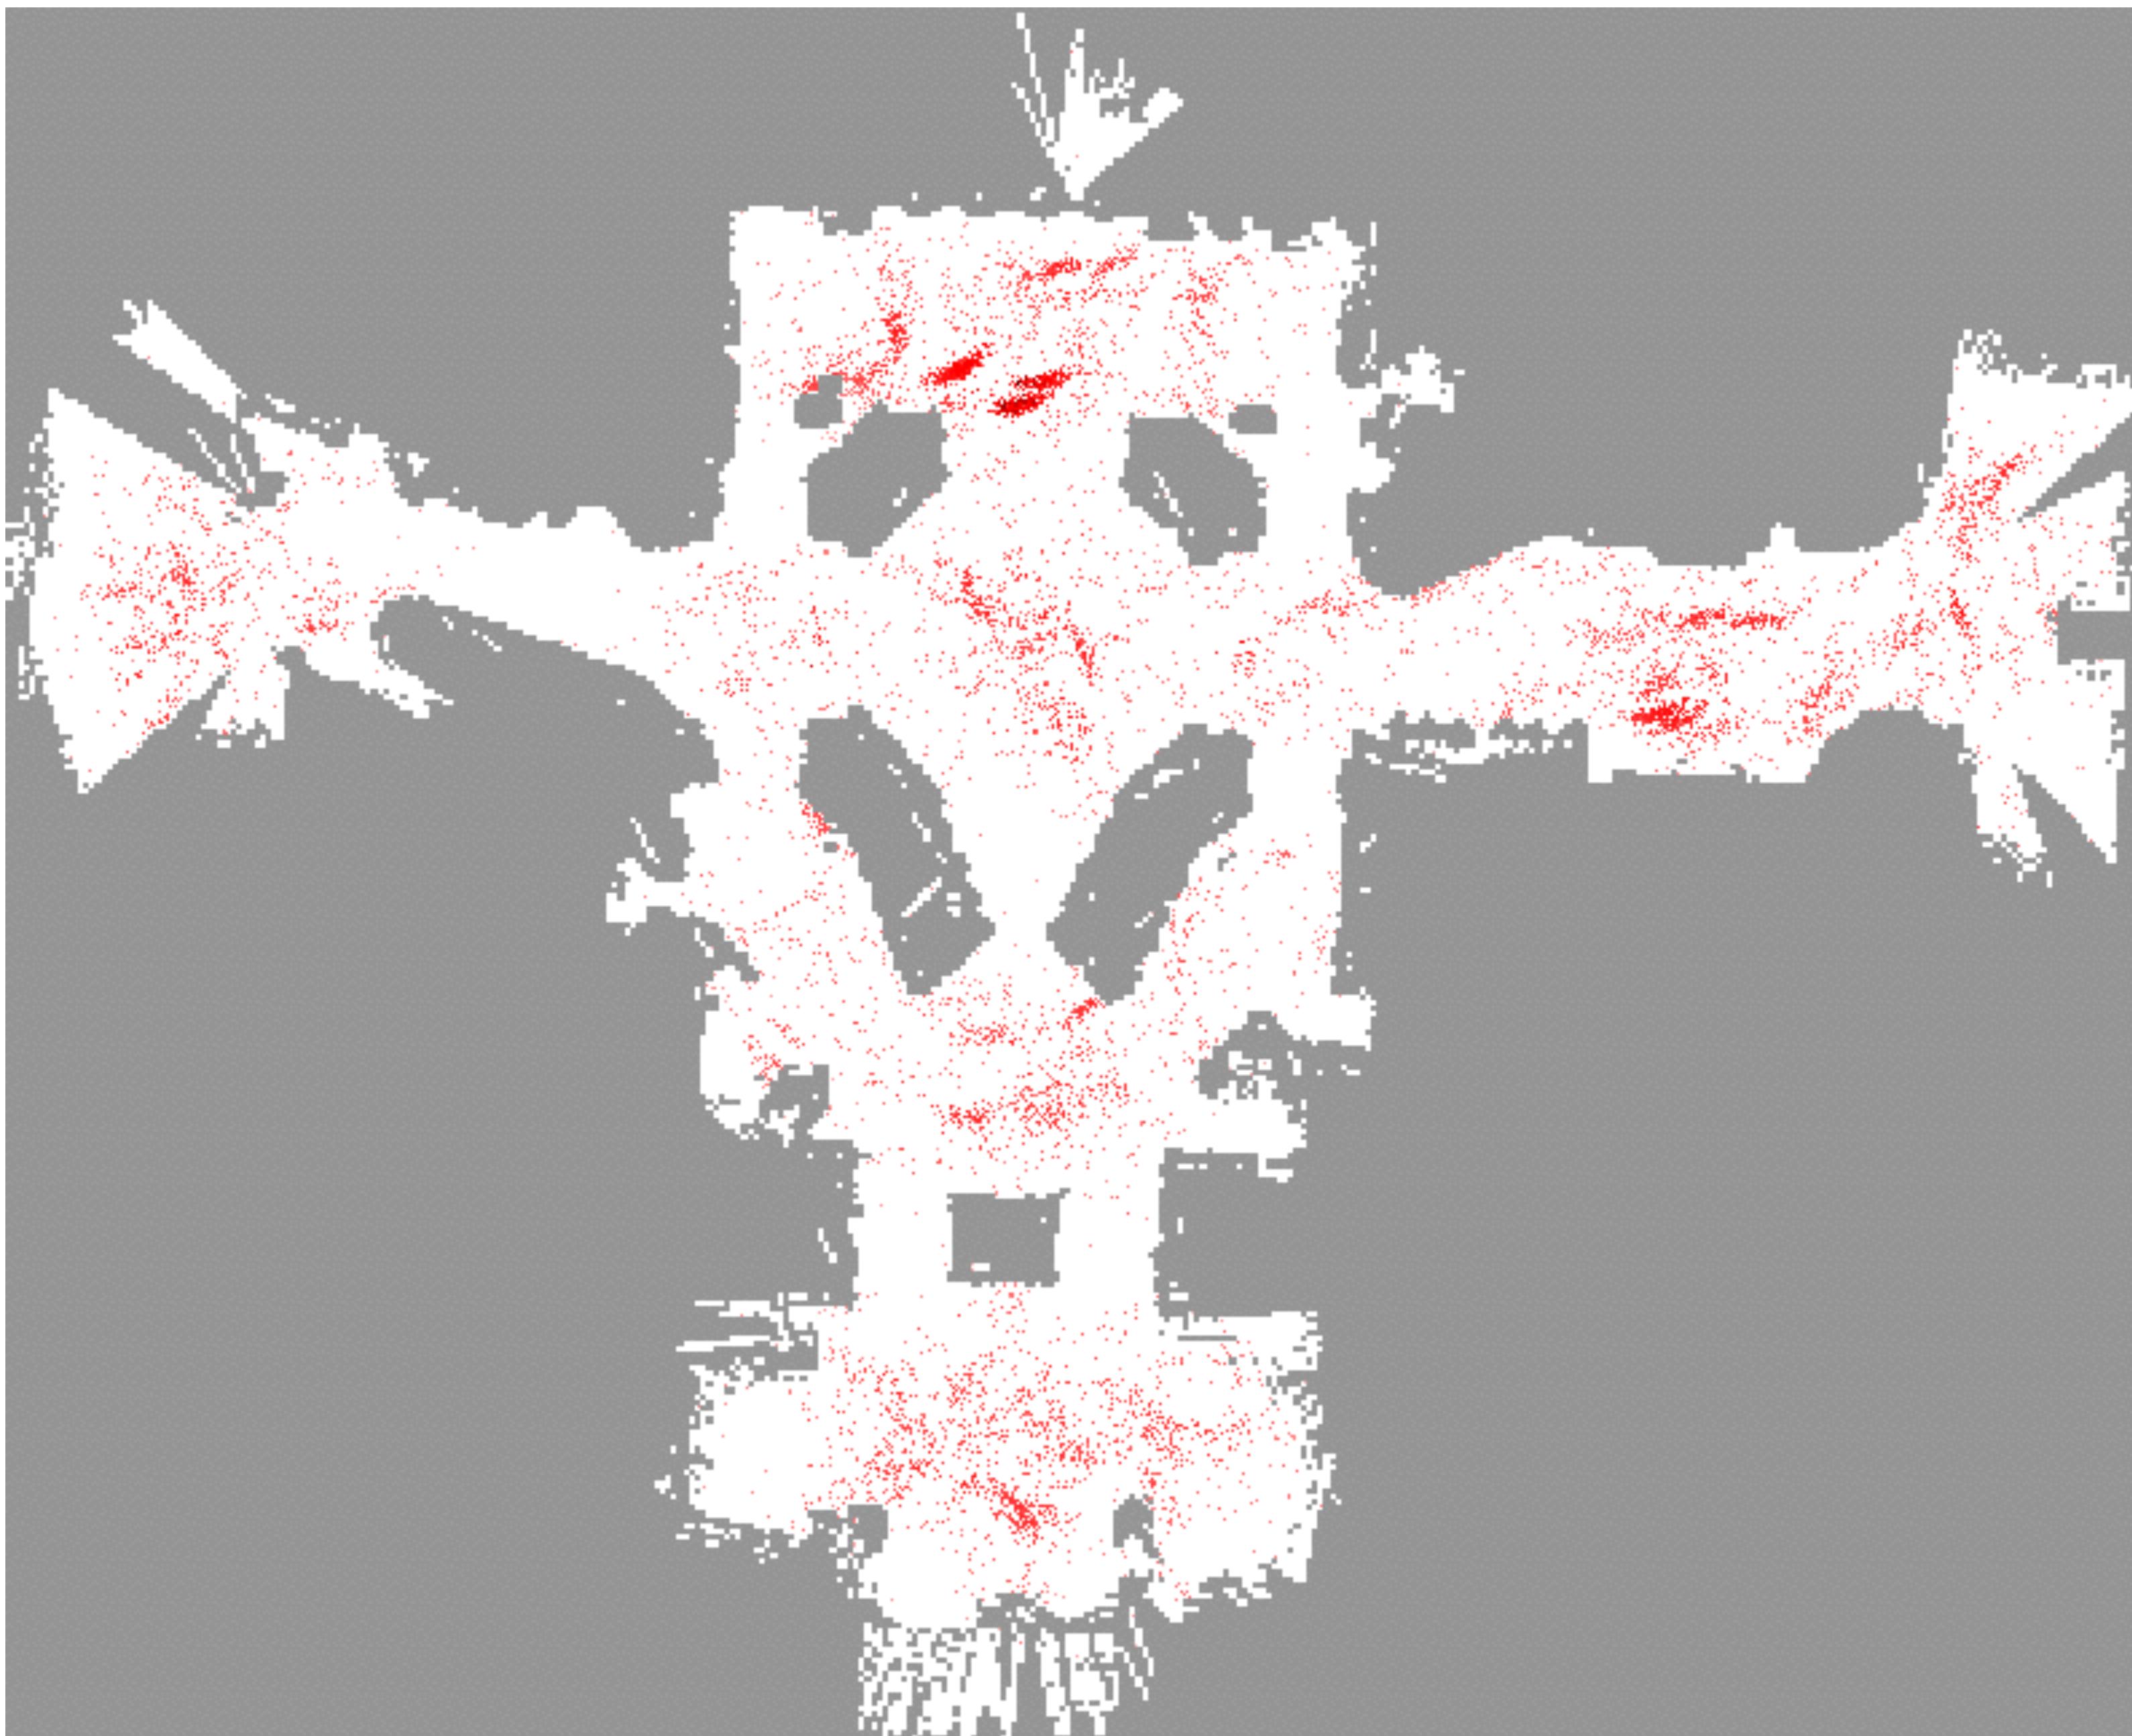


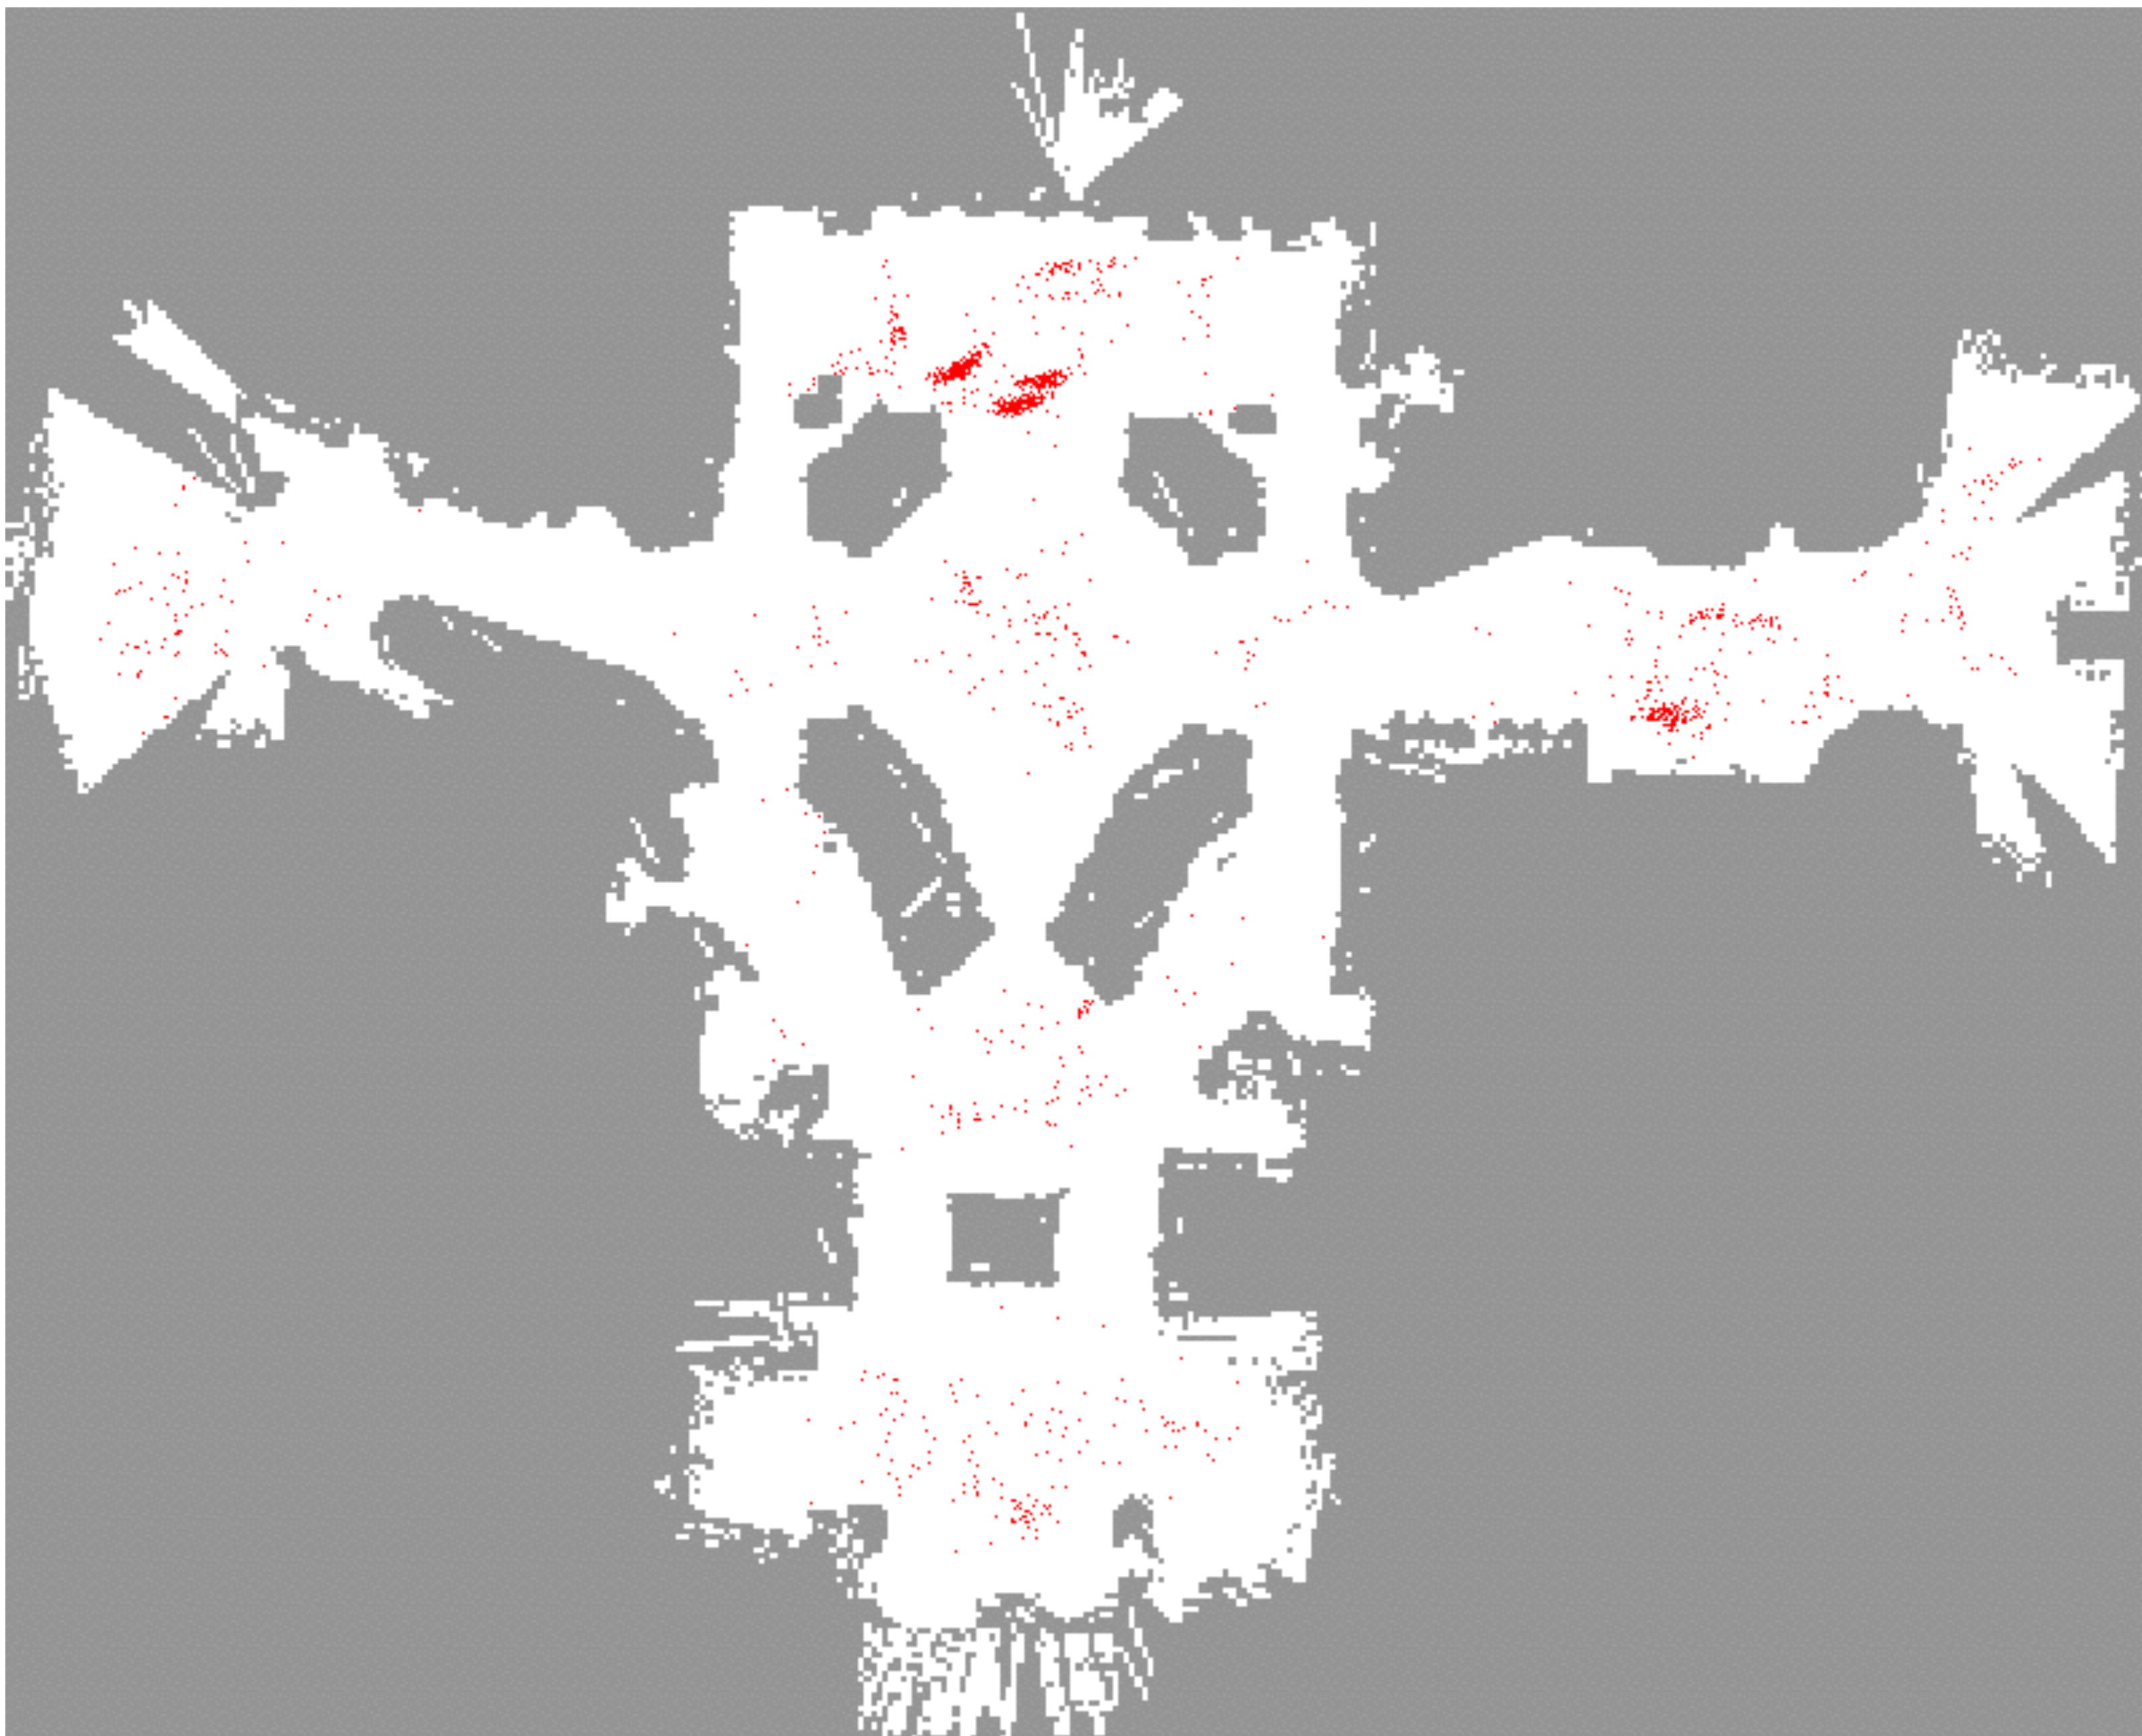




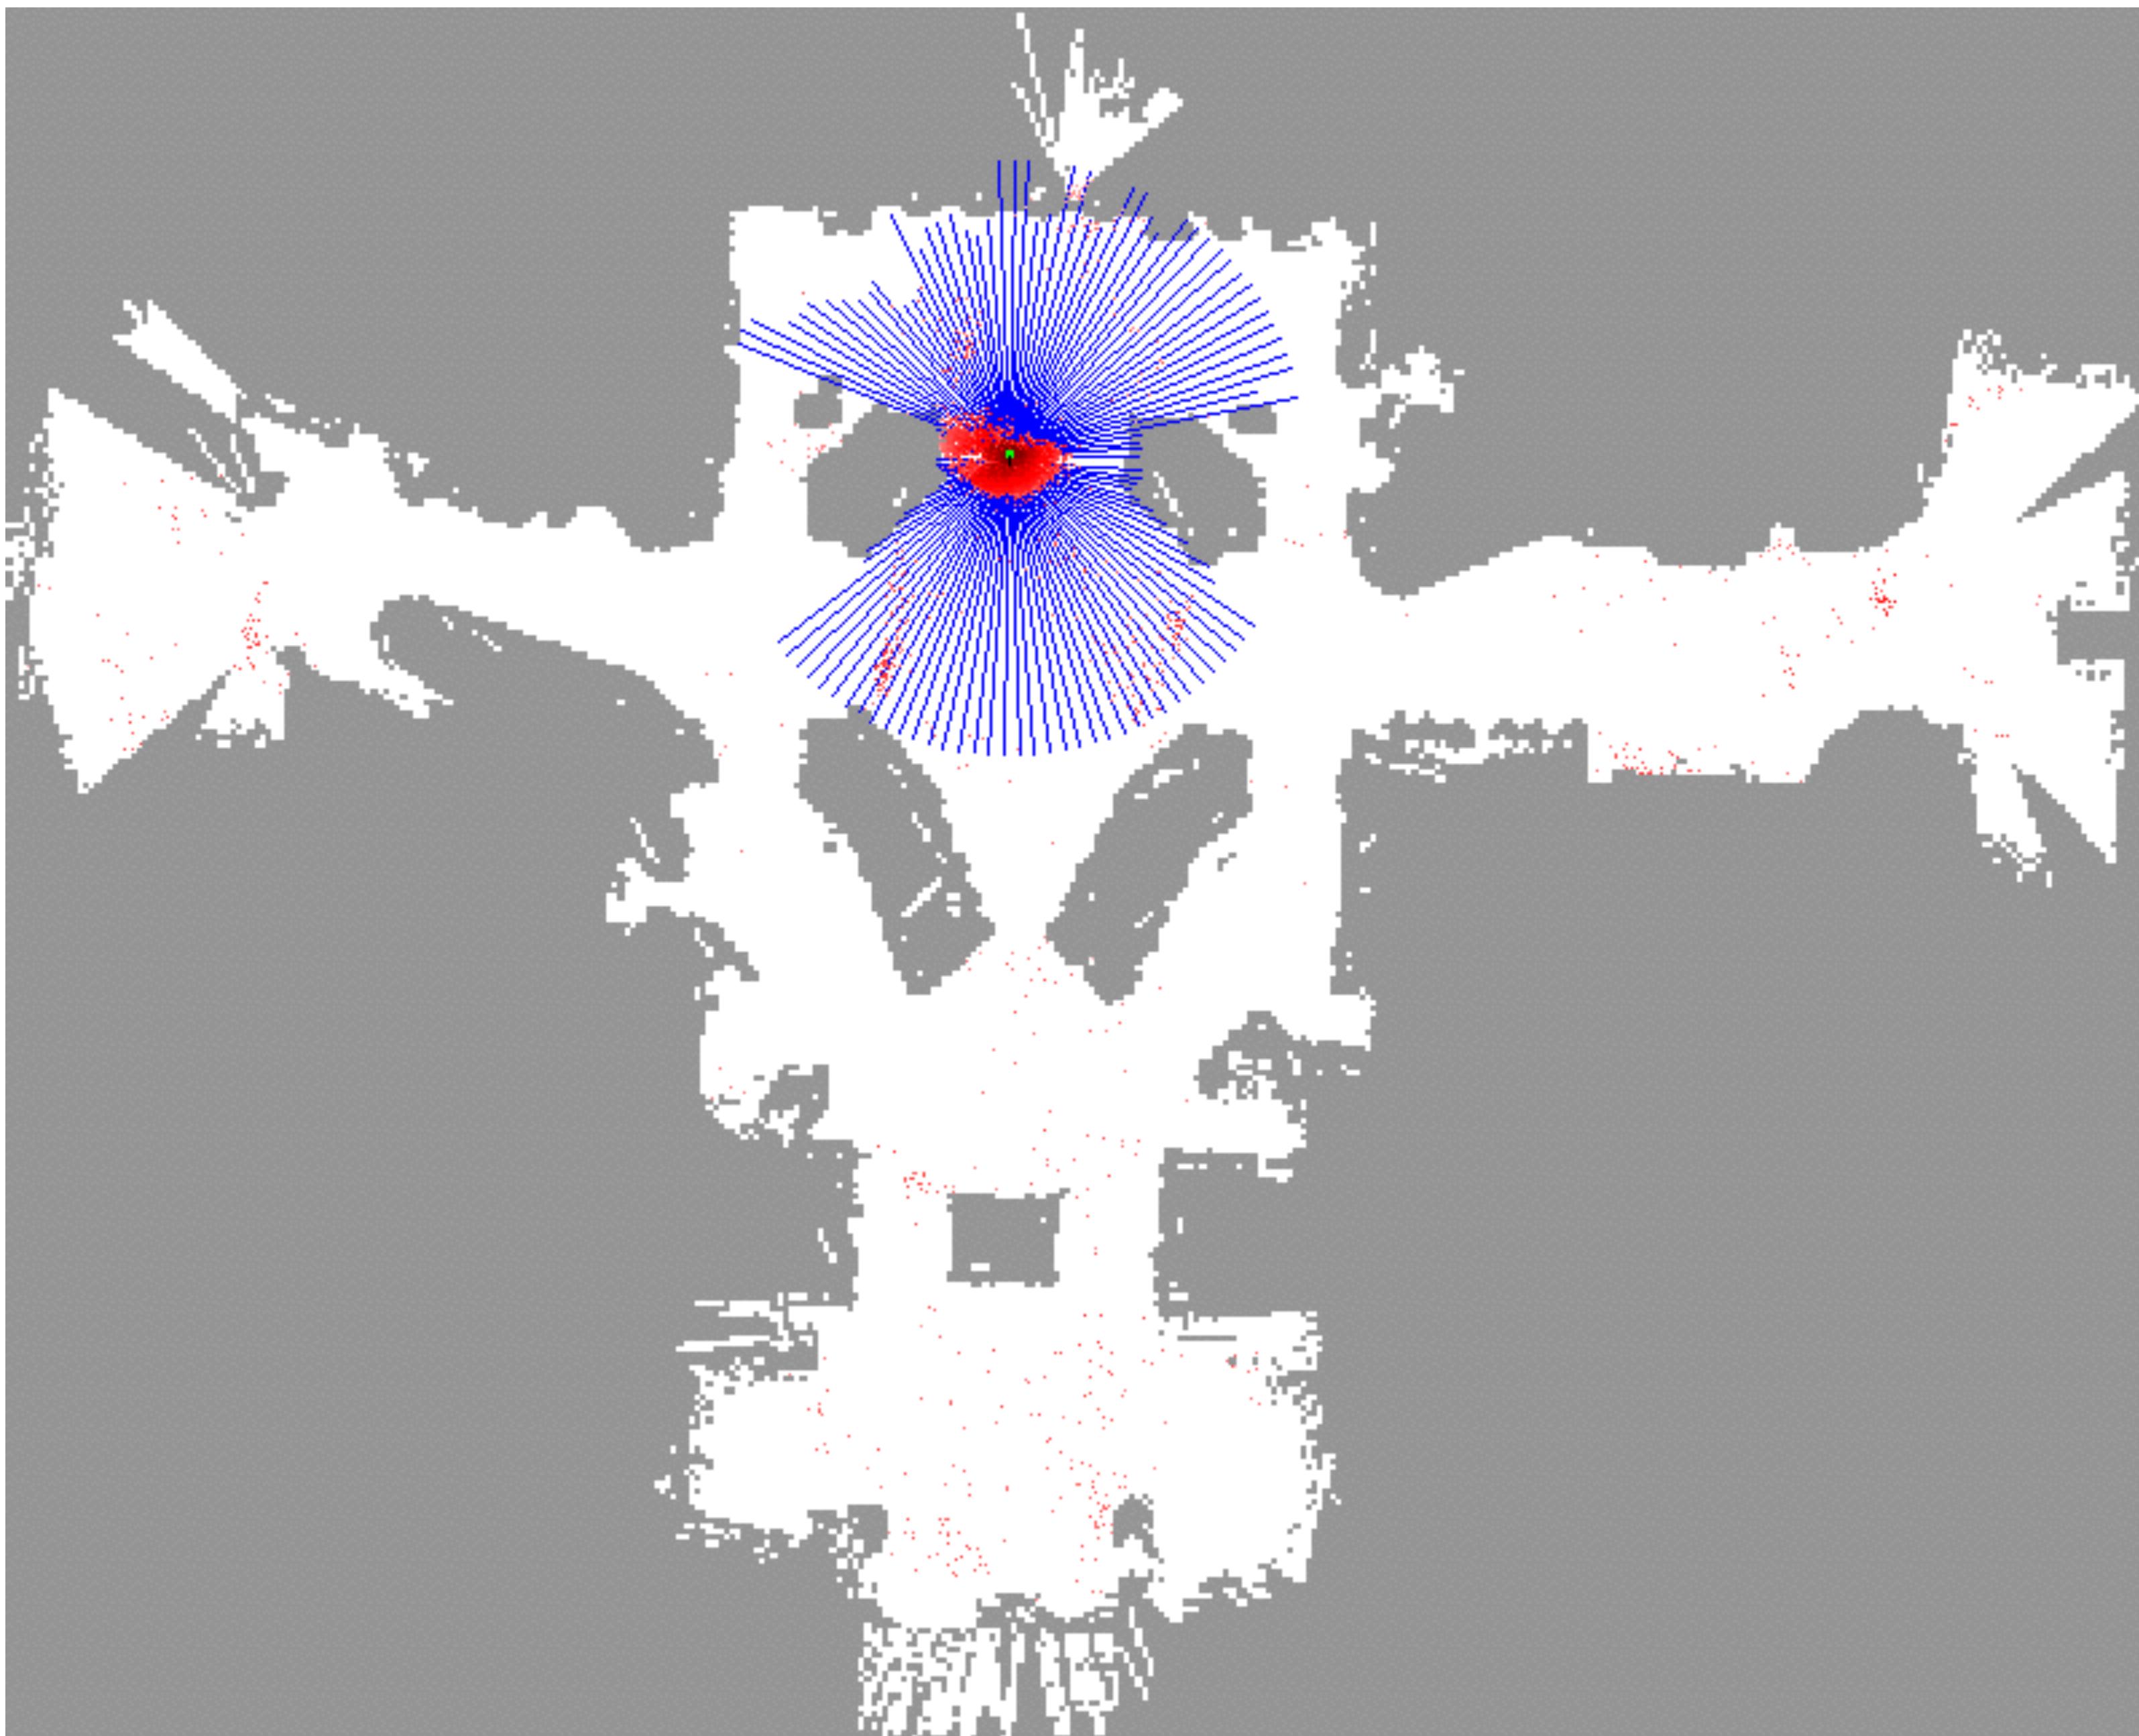




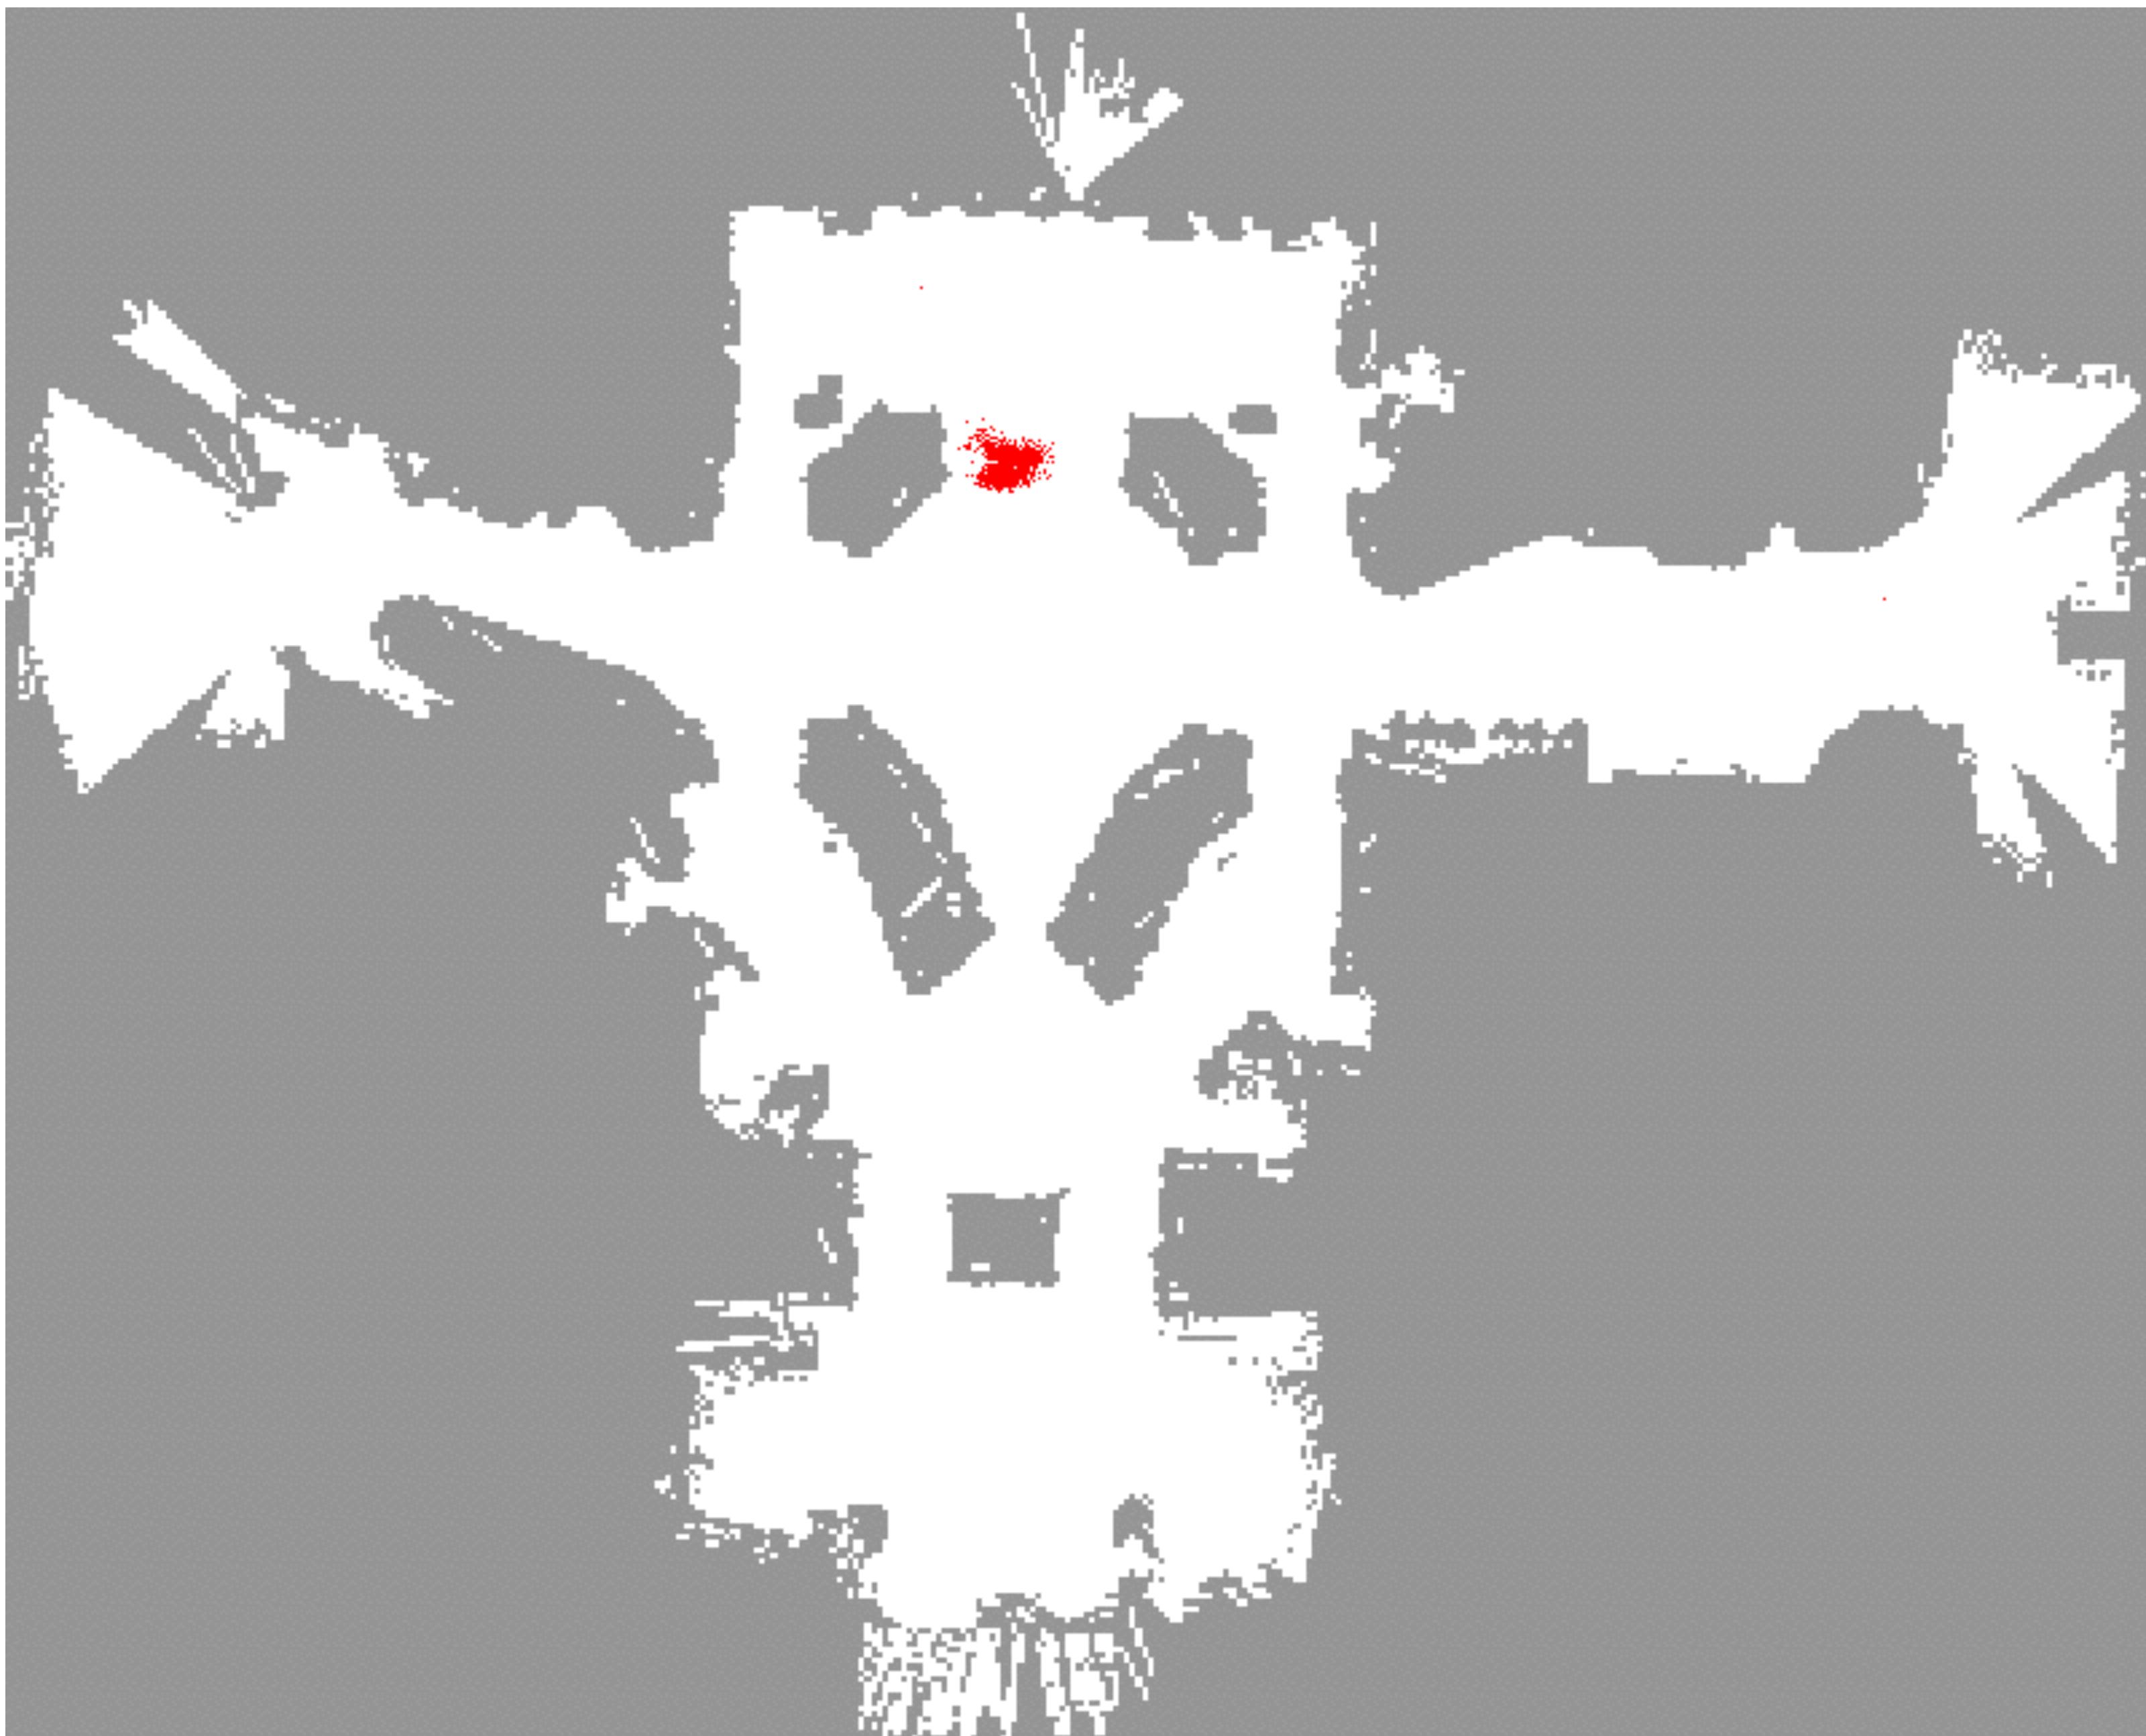


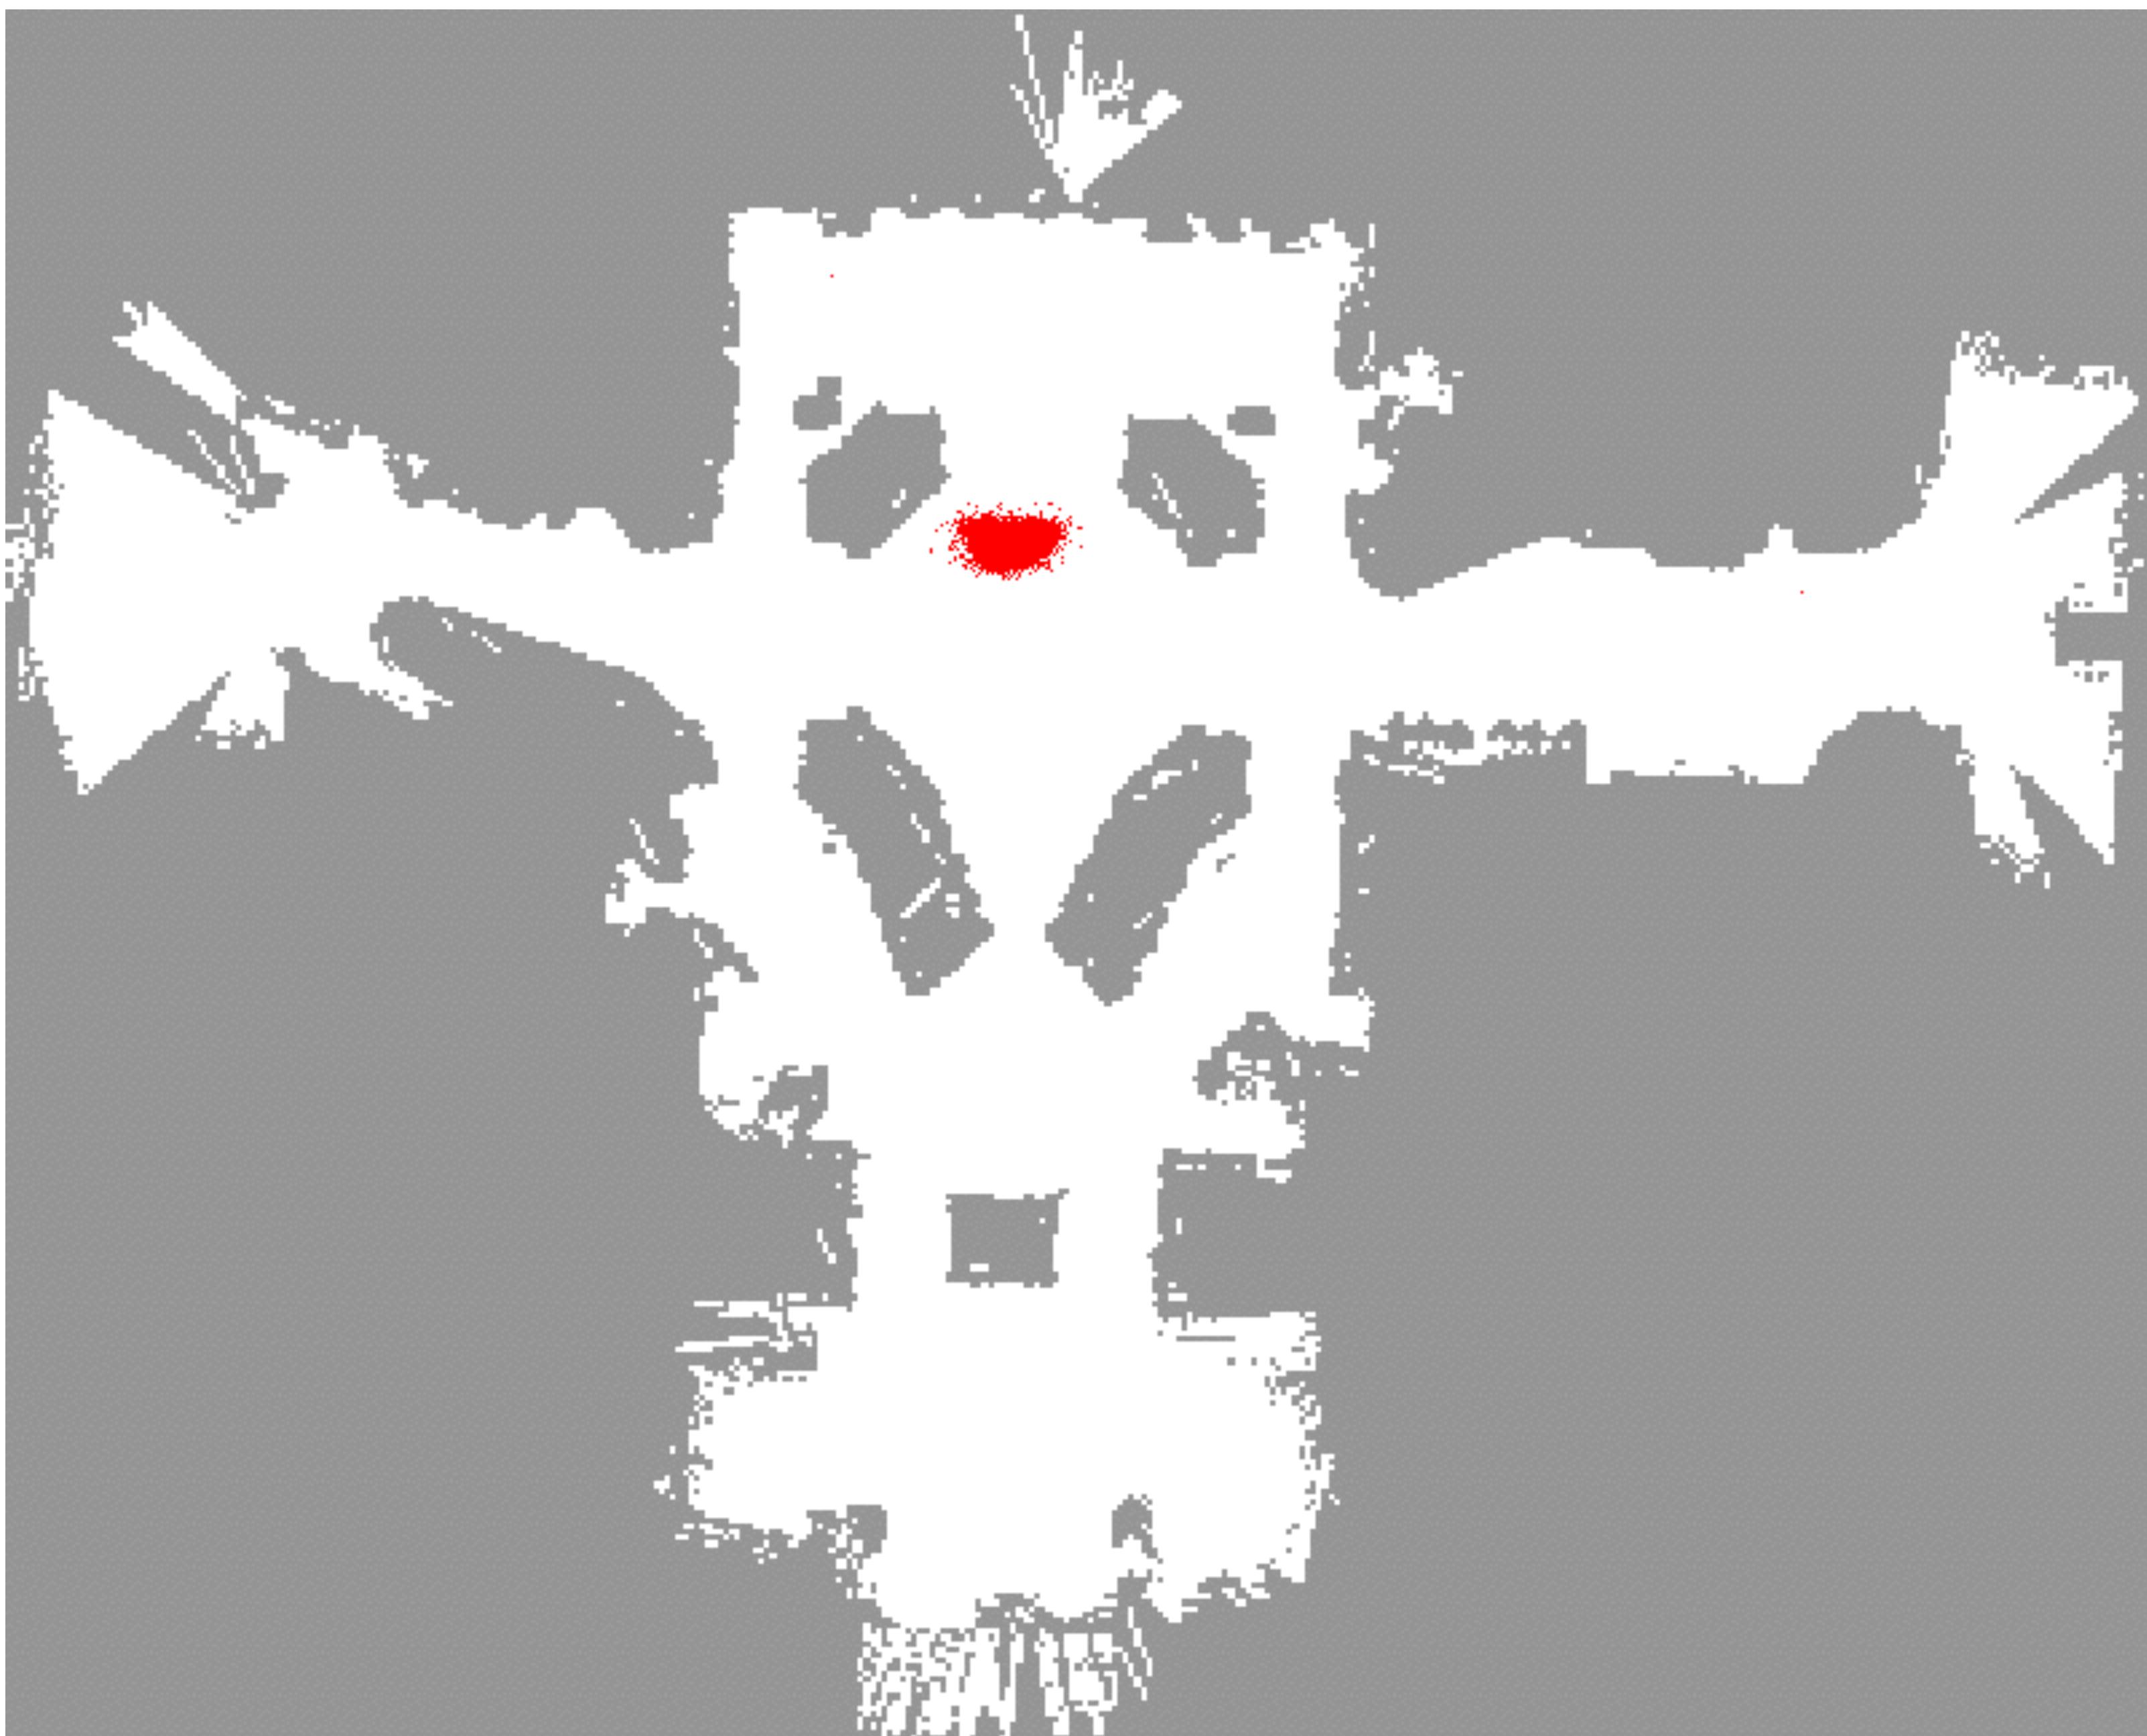


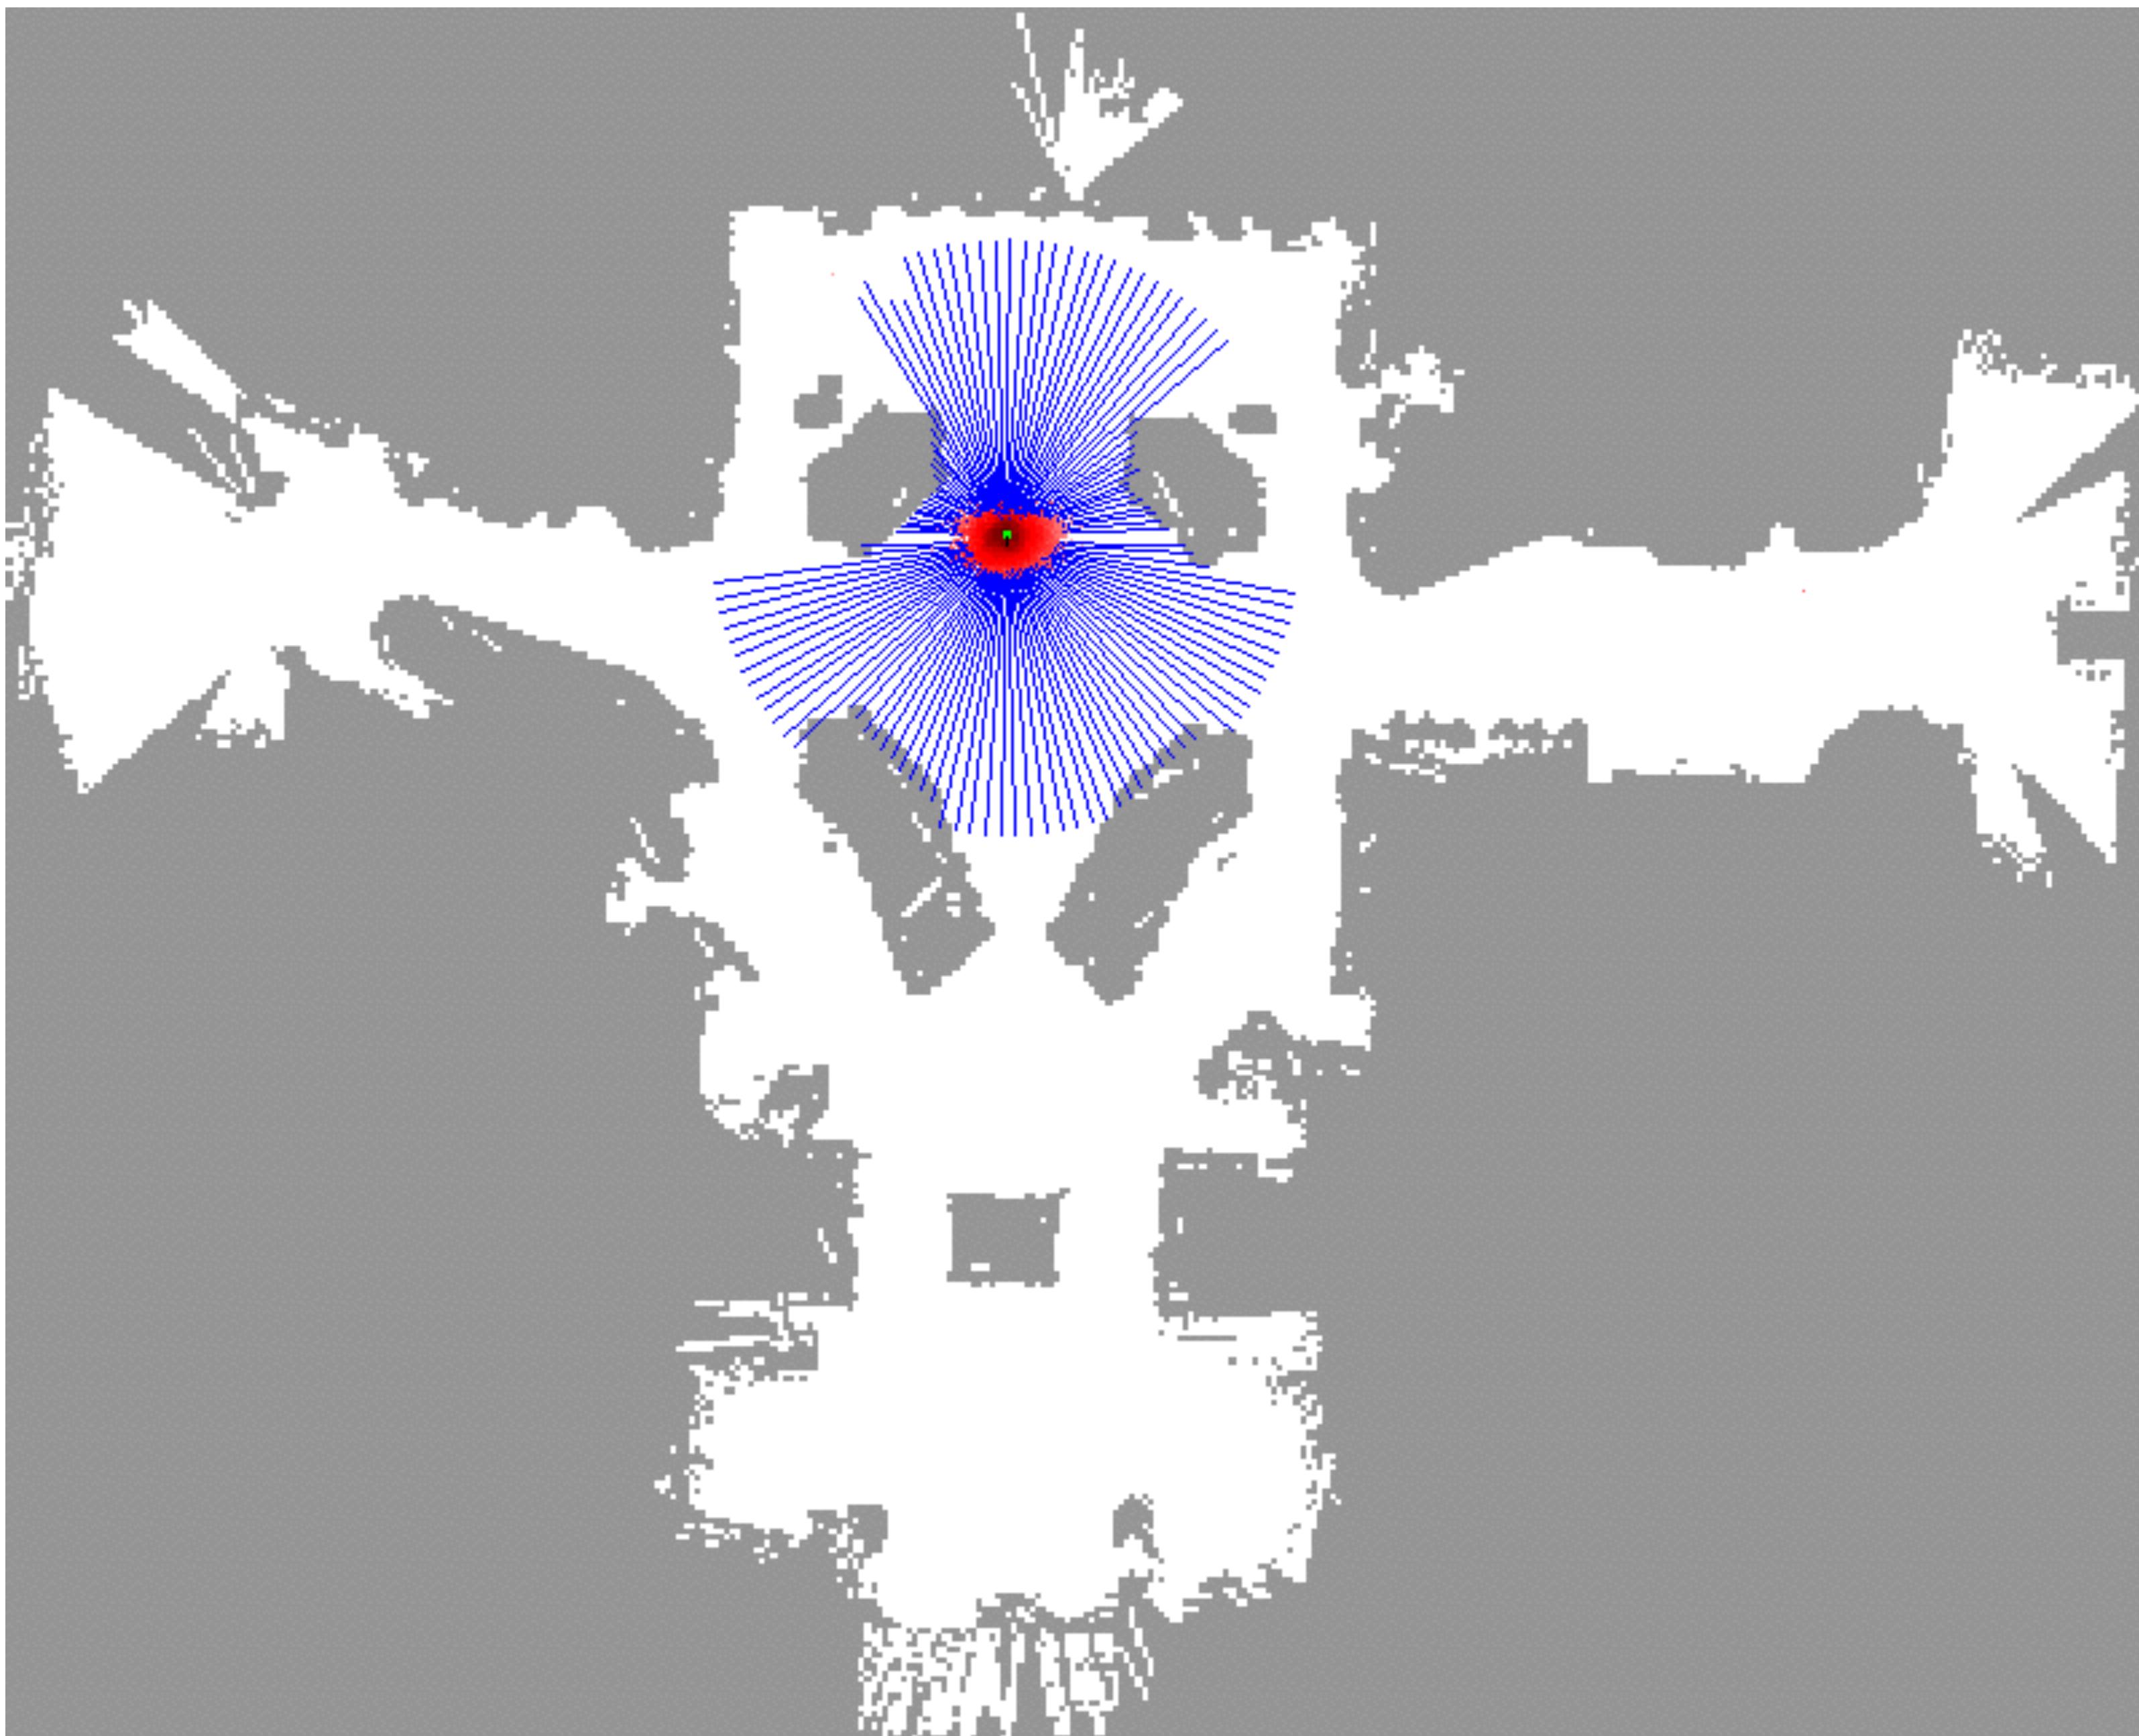


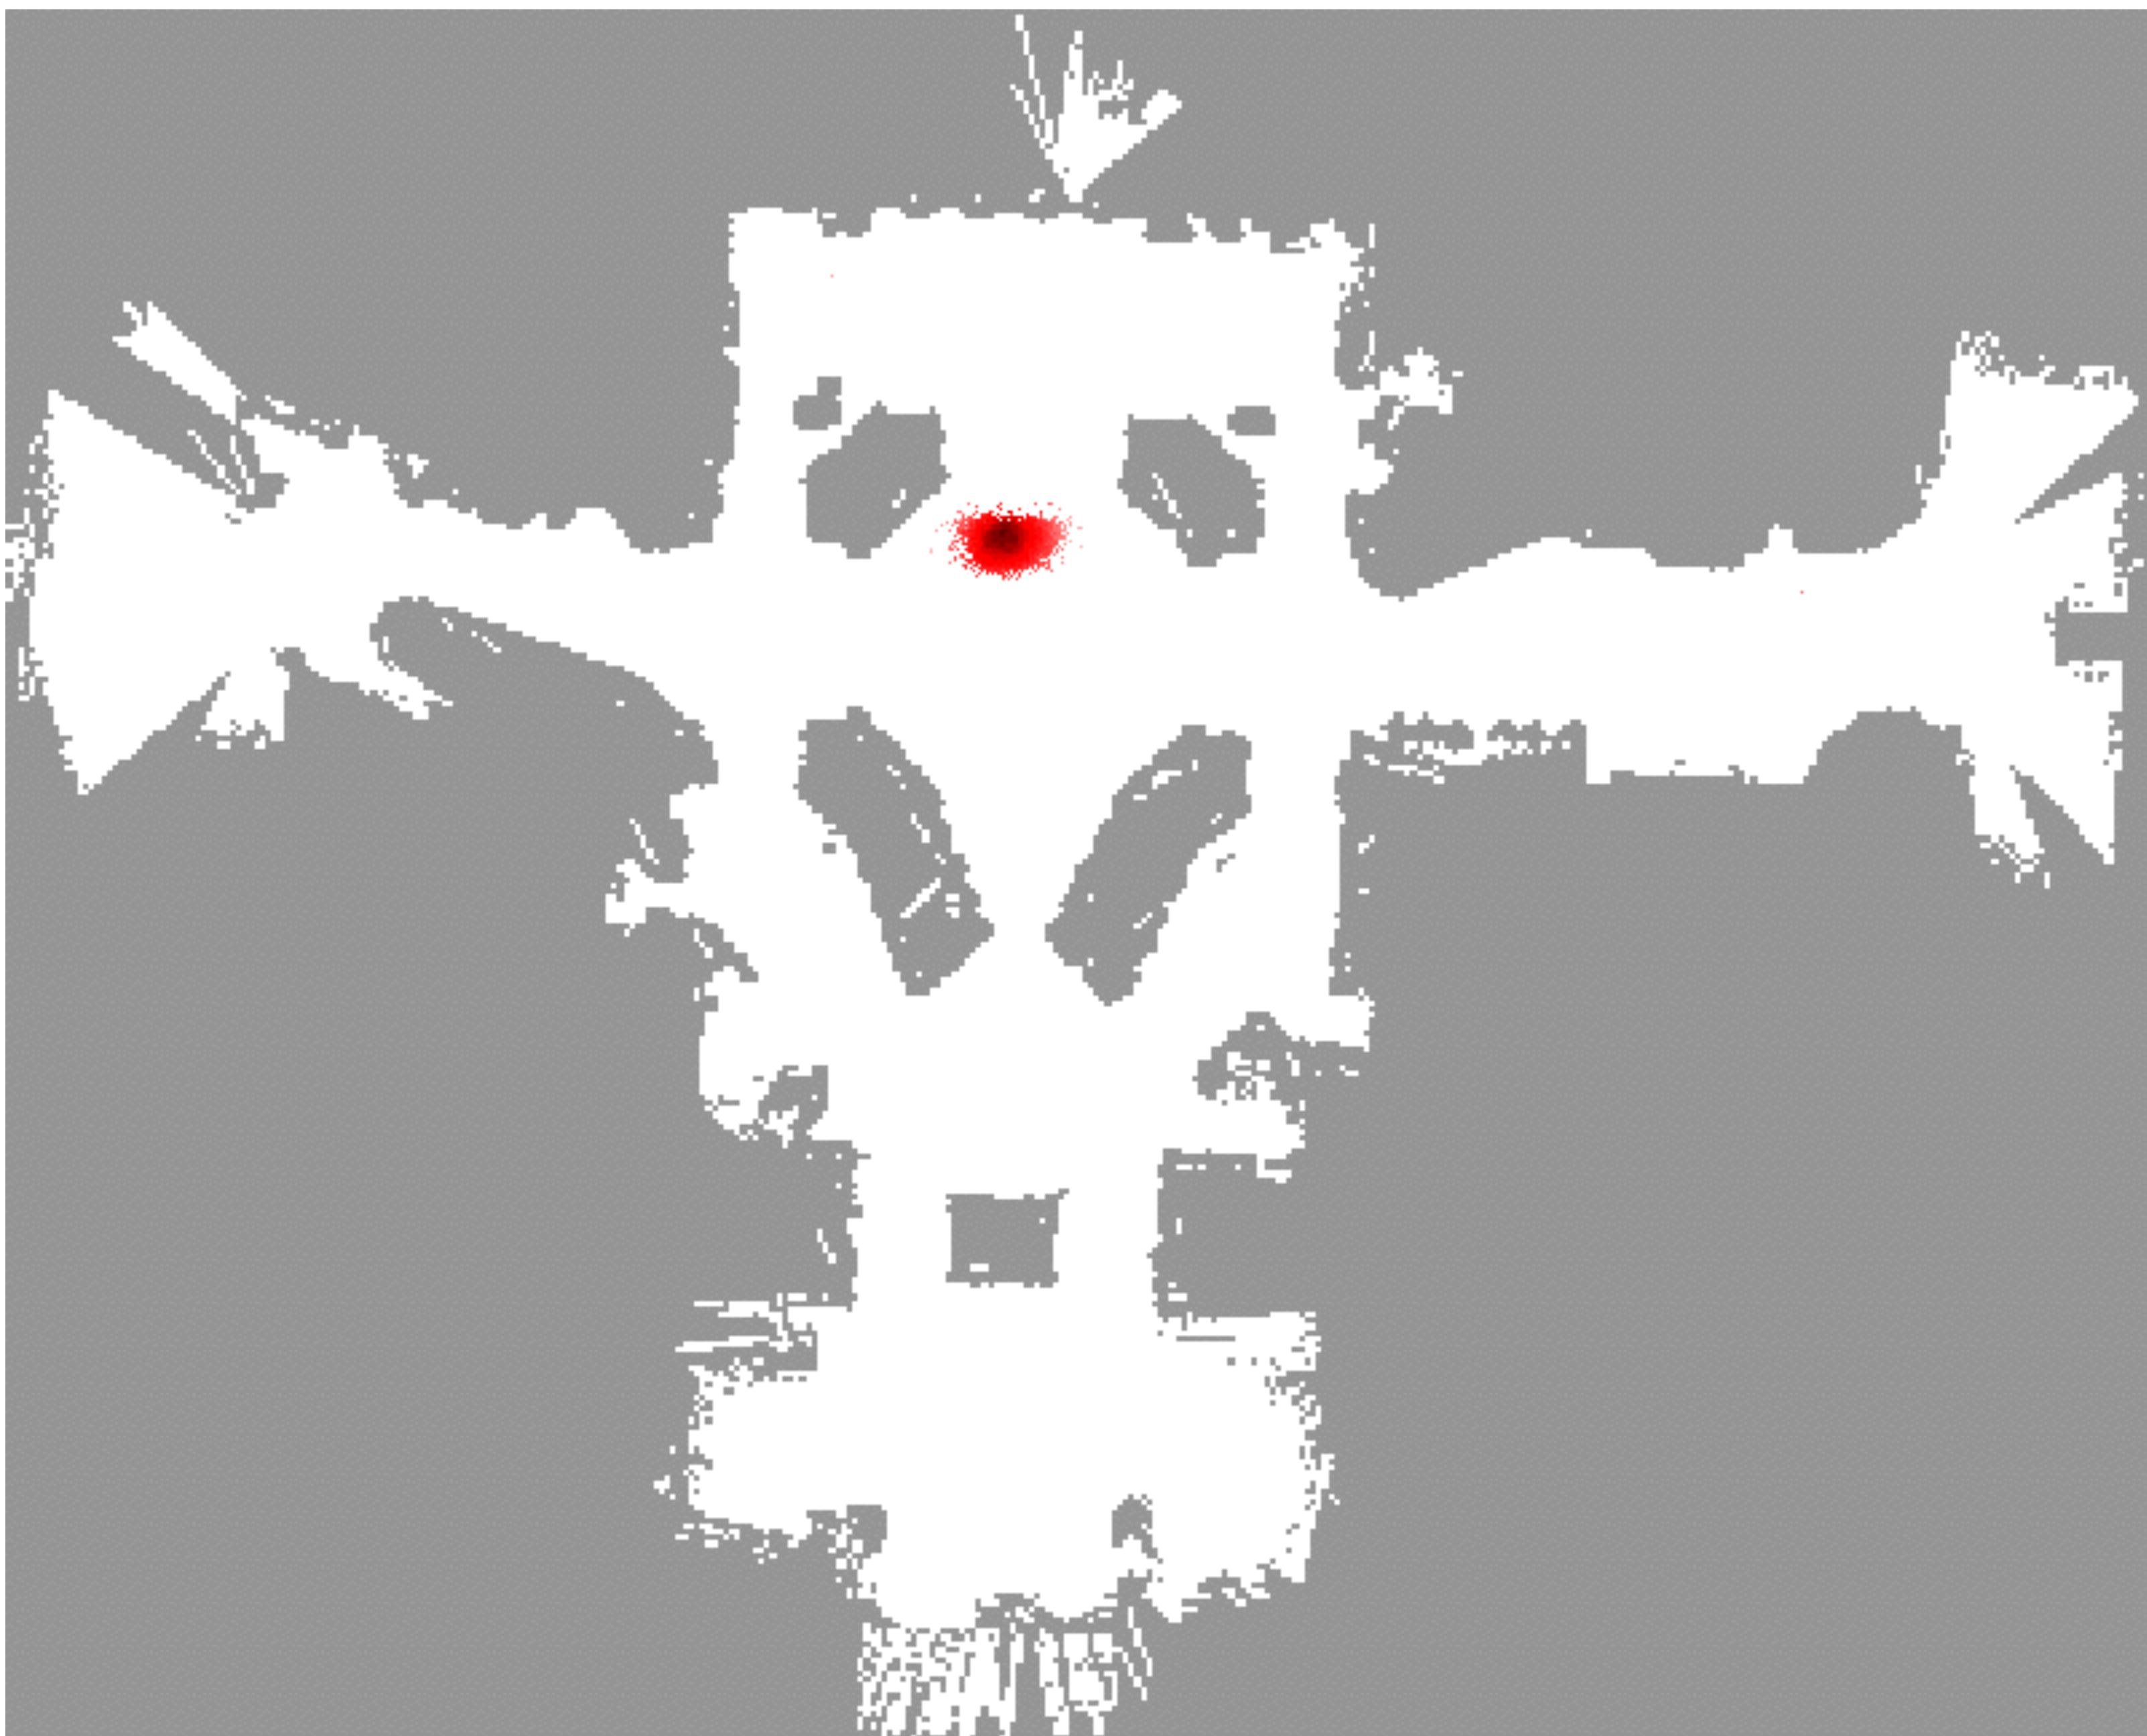


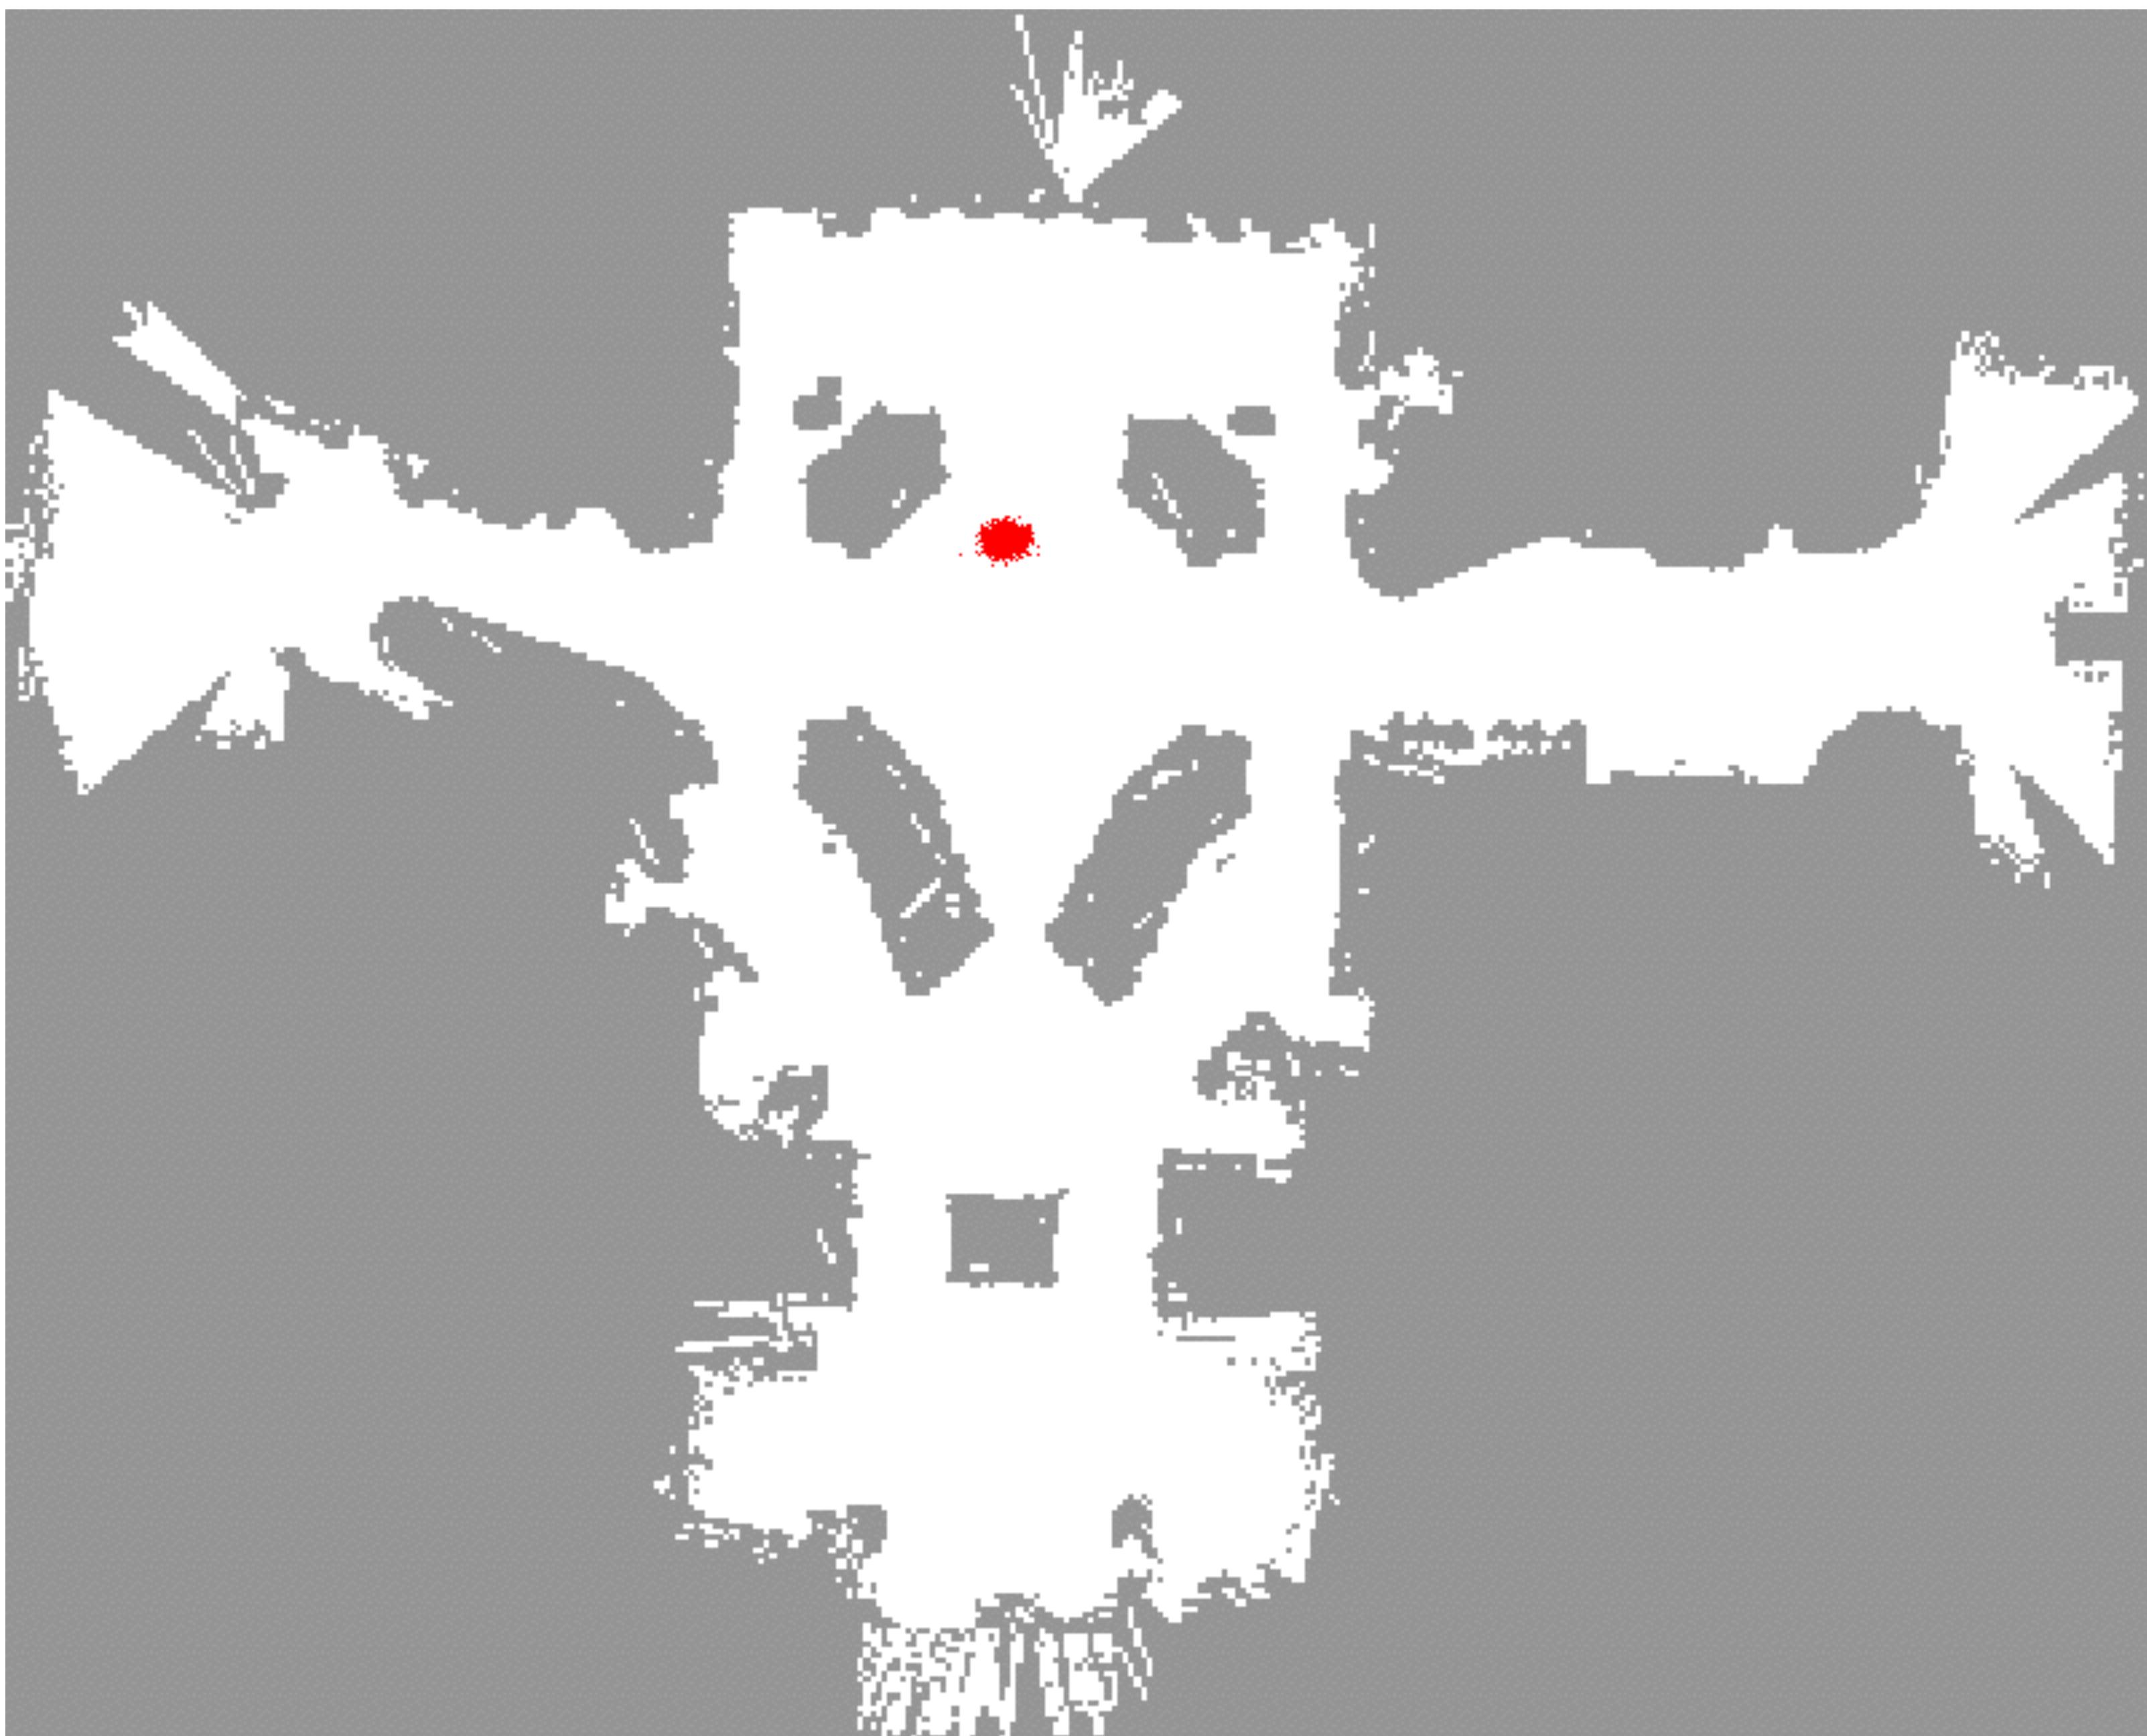


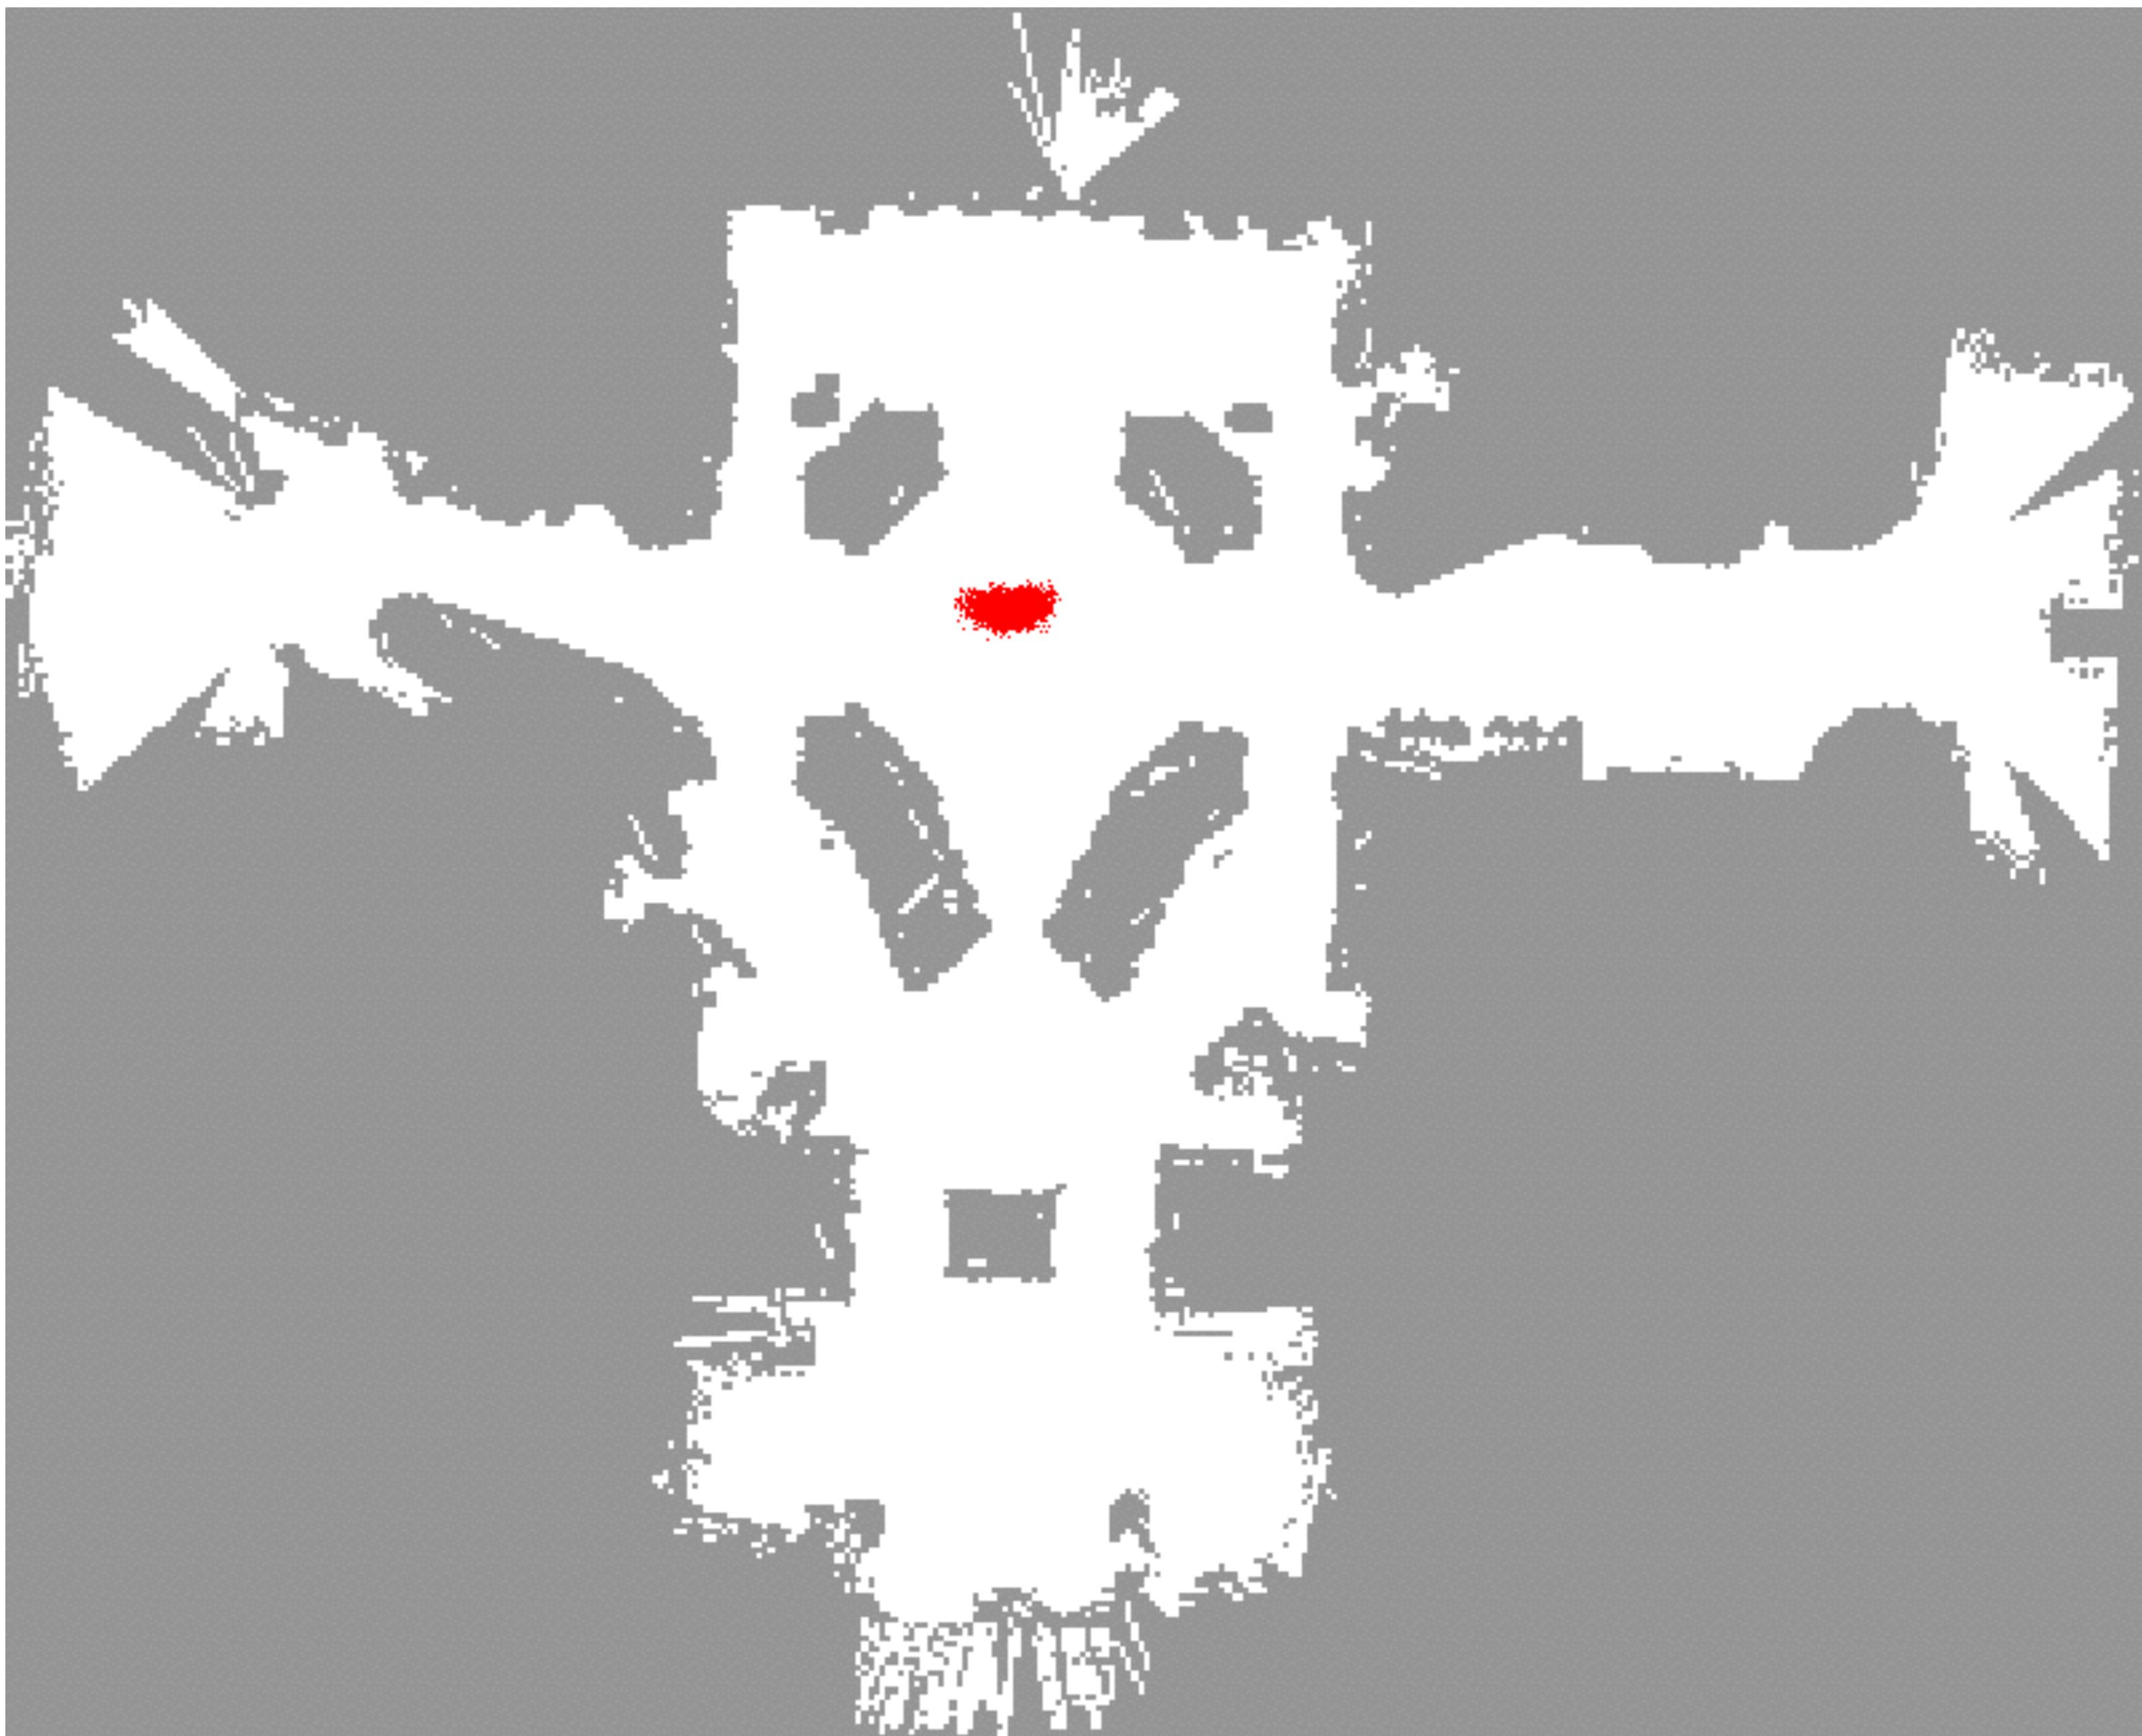


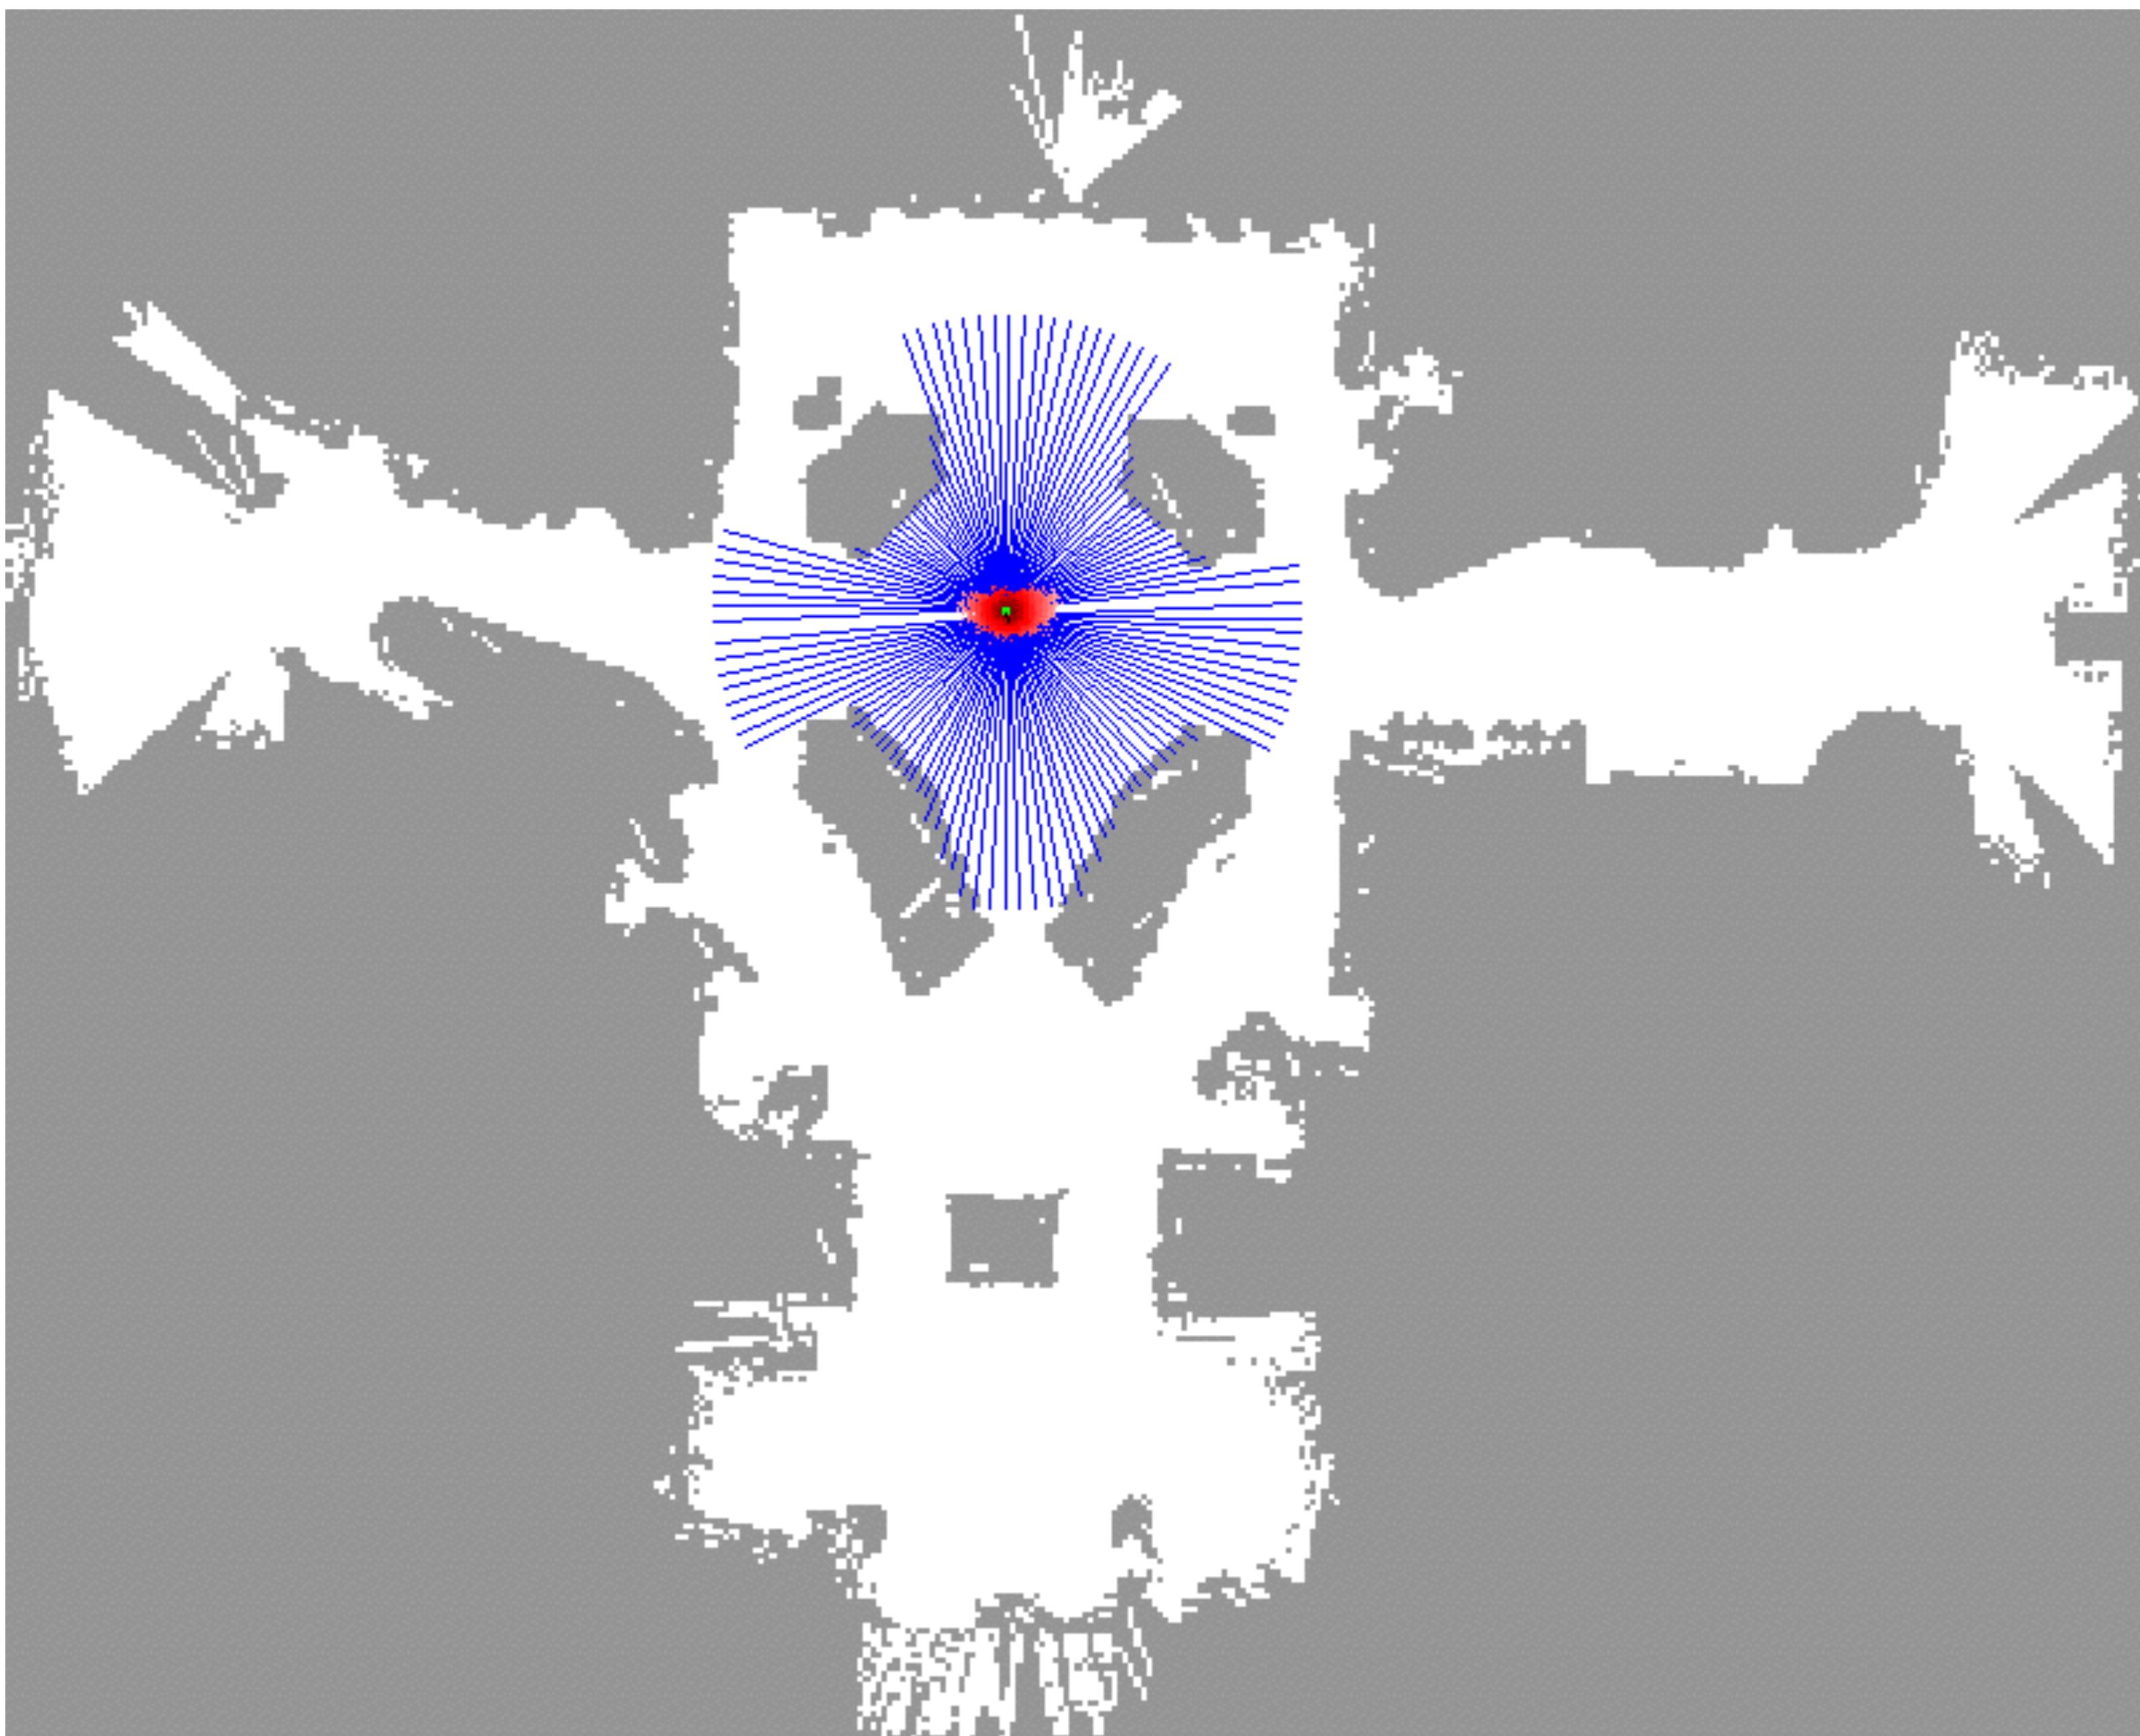


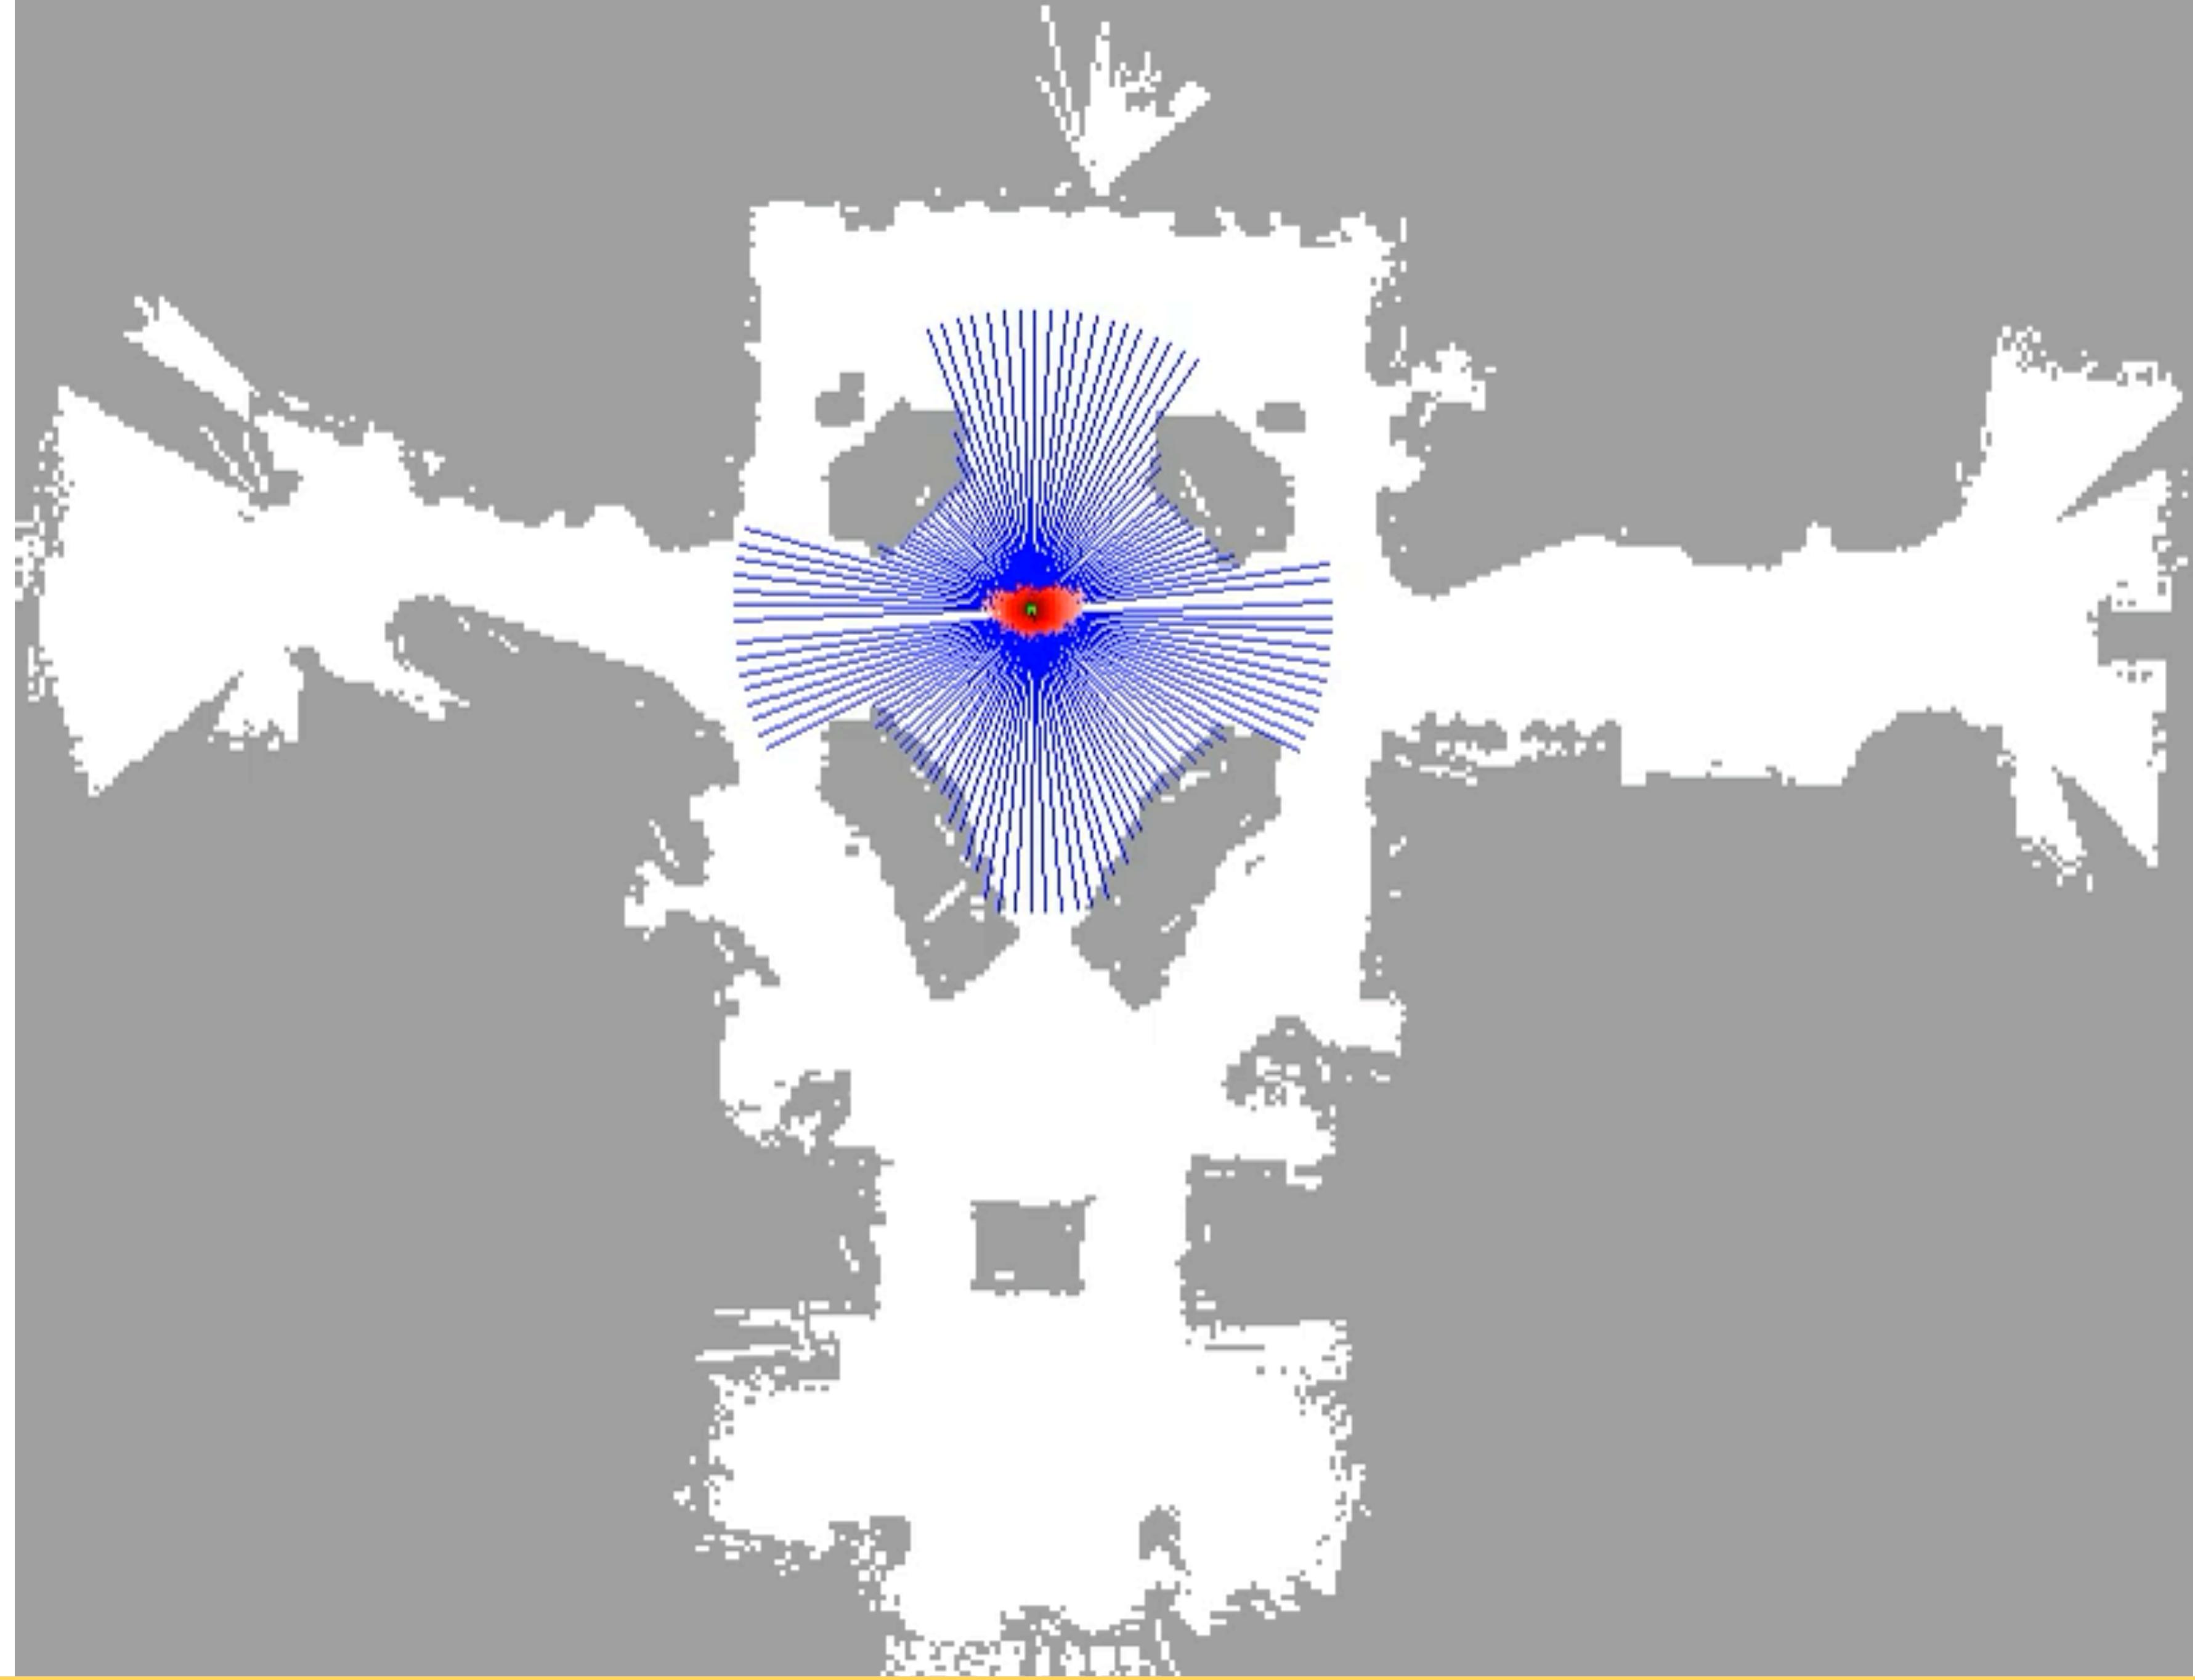










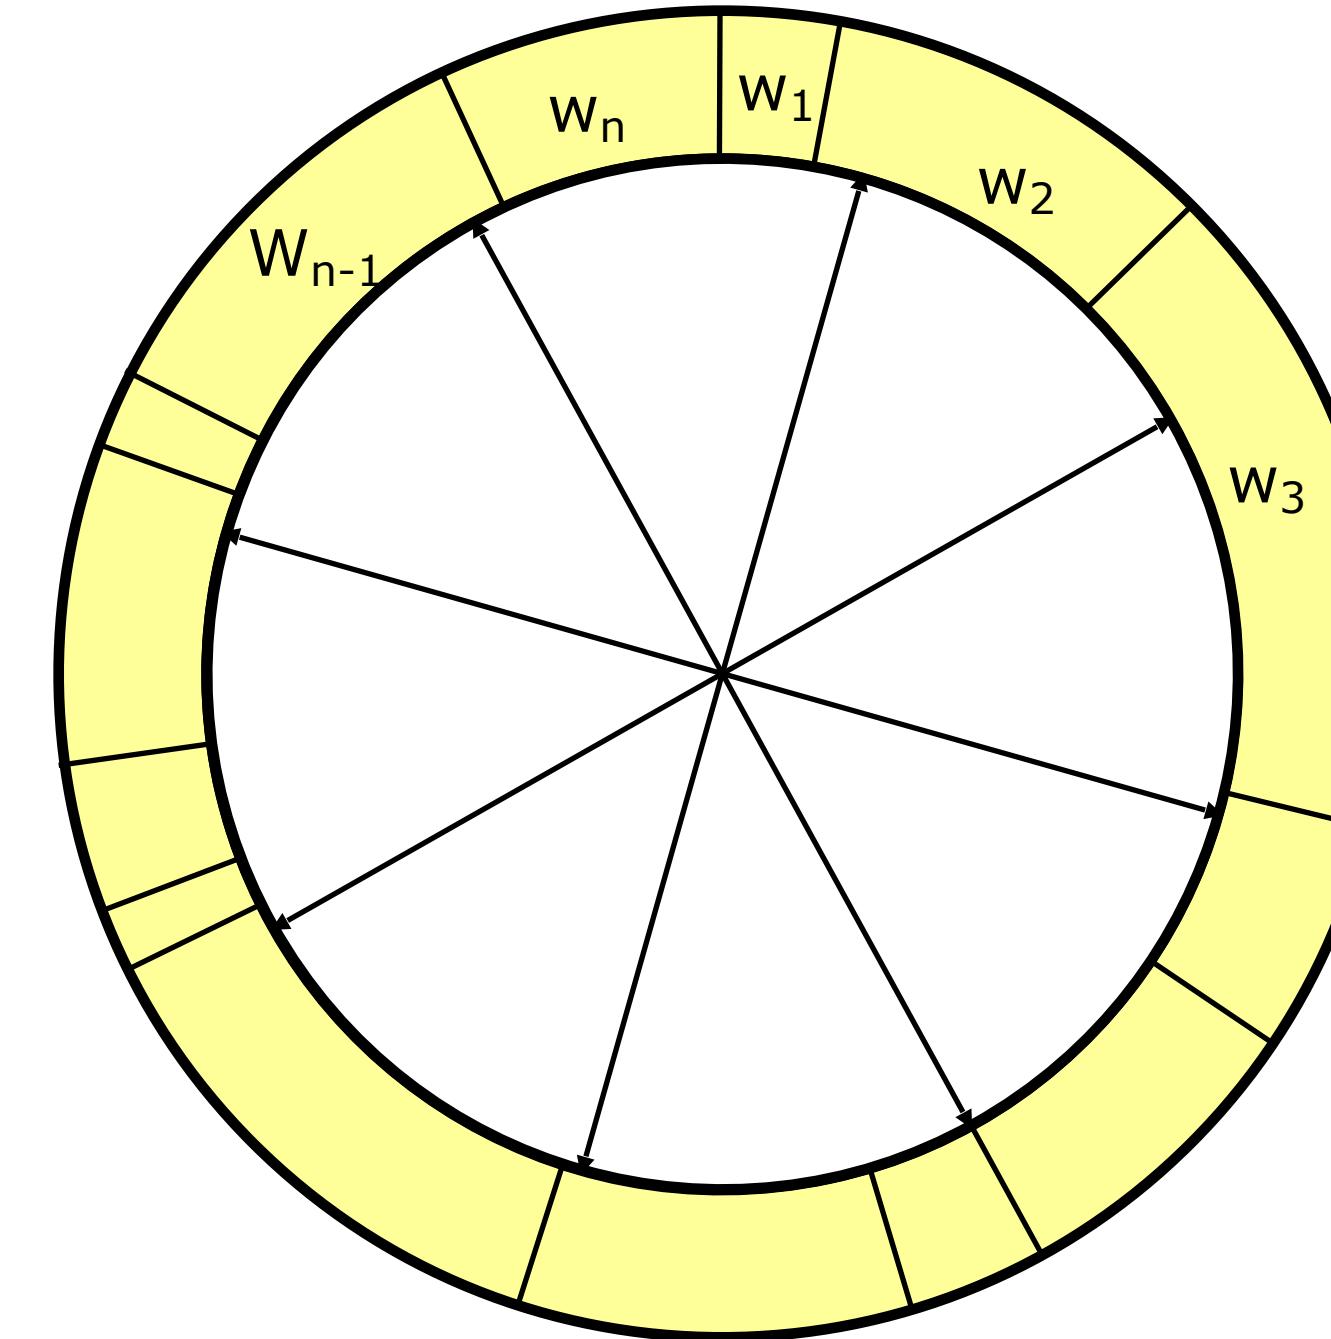
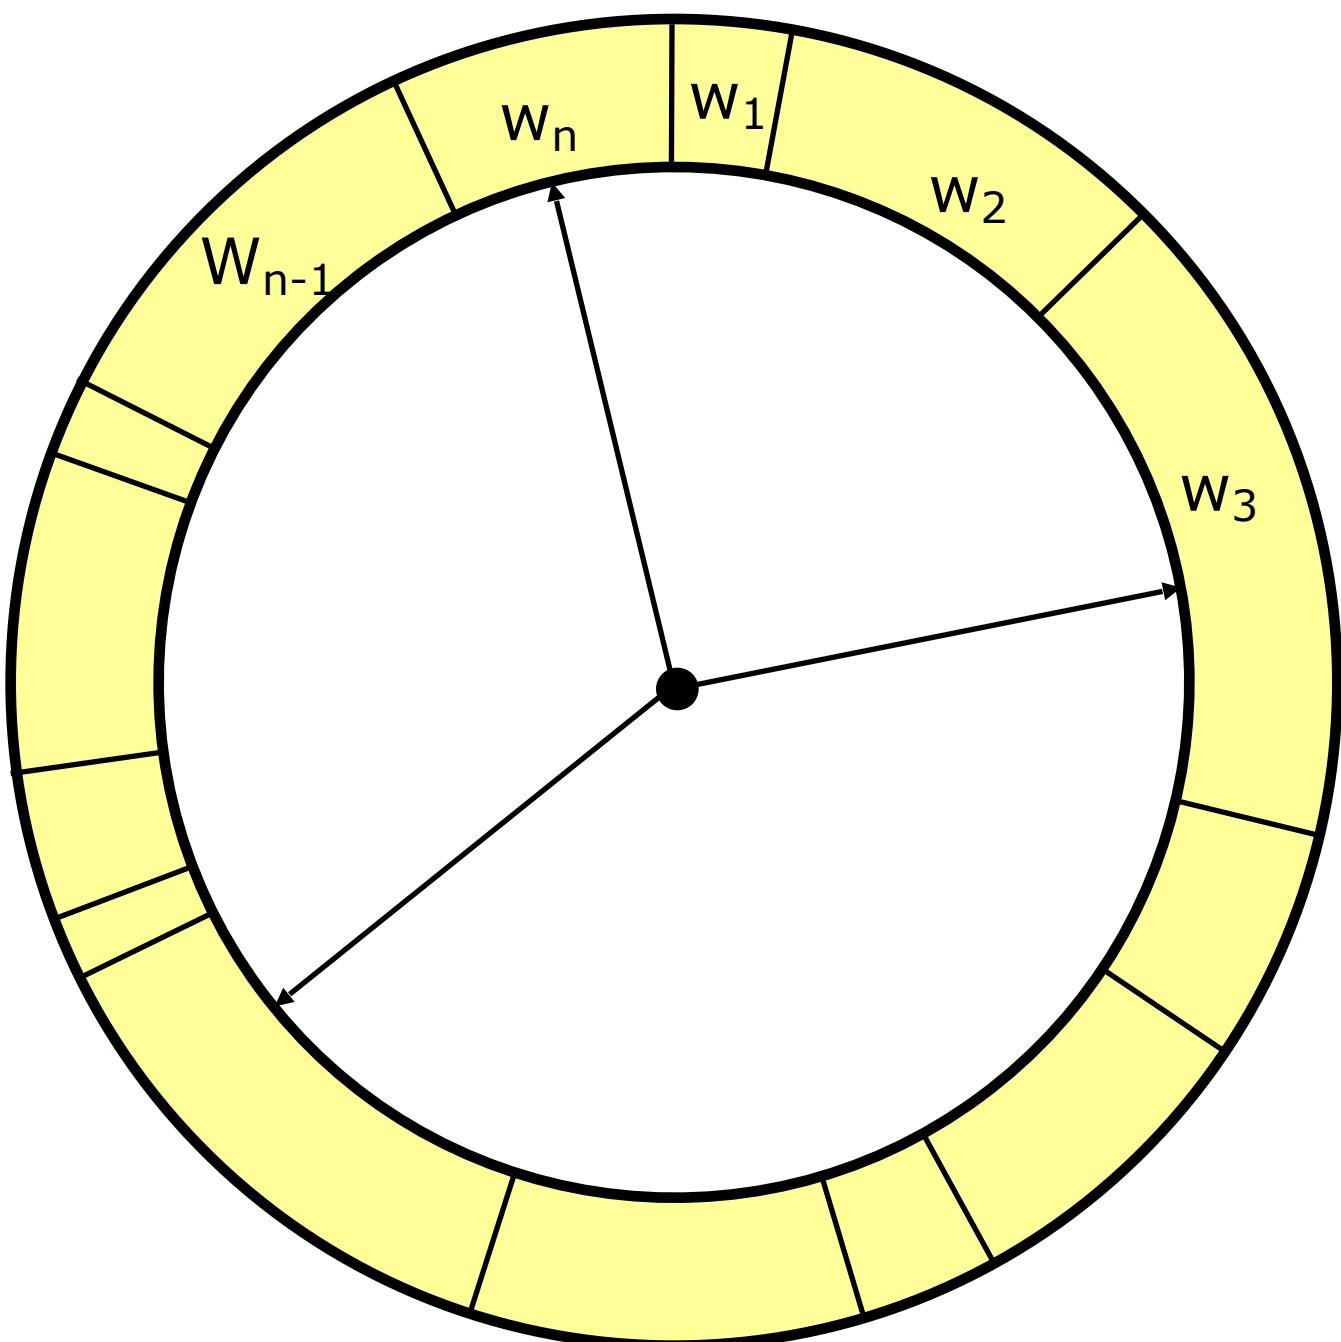


Resampling

- **Given:** Set S of weighted samples.
- **Wanted :** Random sample, where the probability of drawing x_i is given by w_i .
- Typically done n times with replacement to generate new sample set S' .

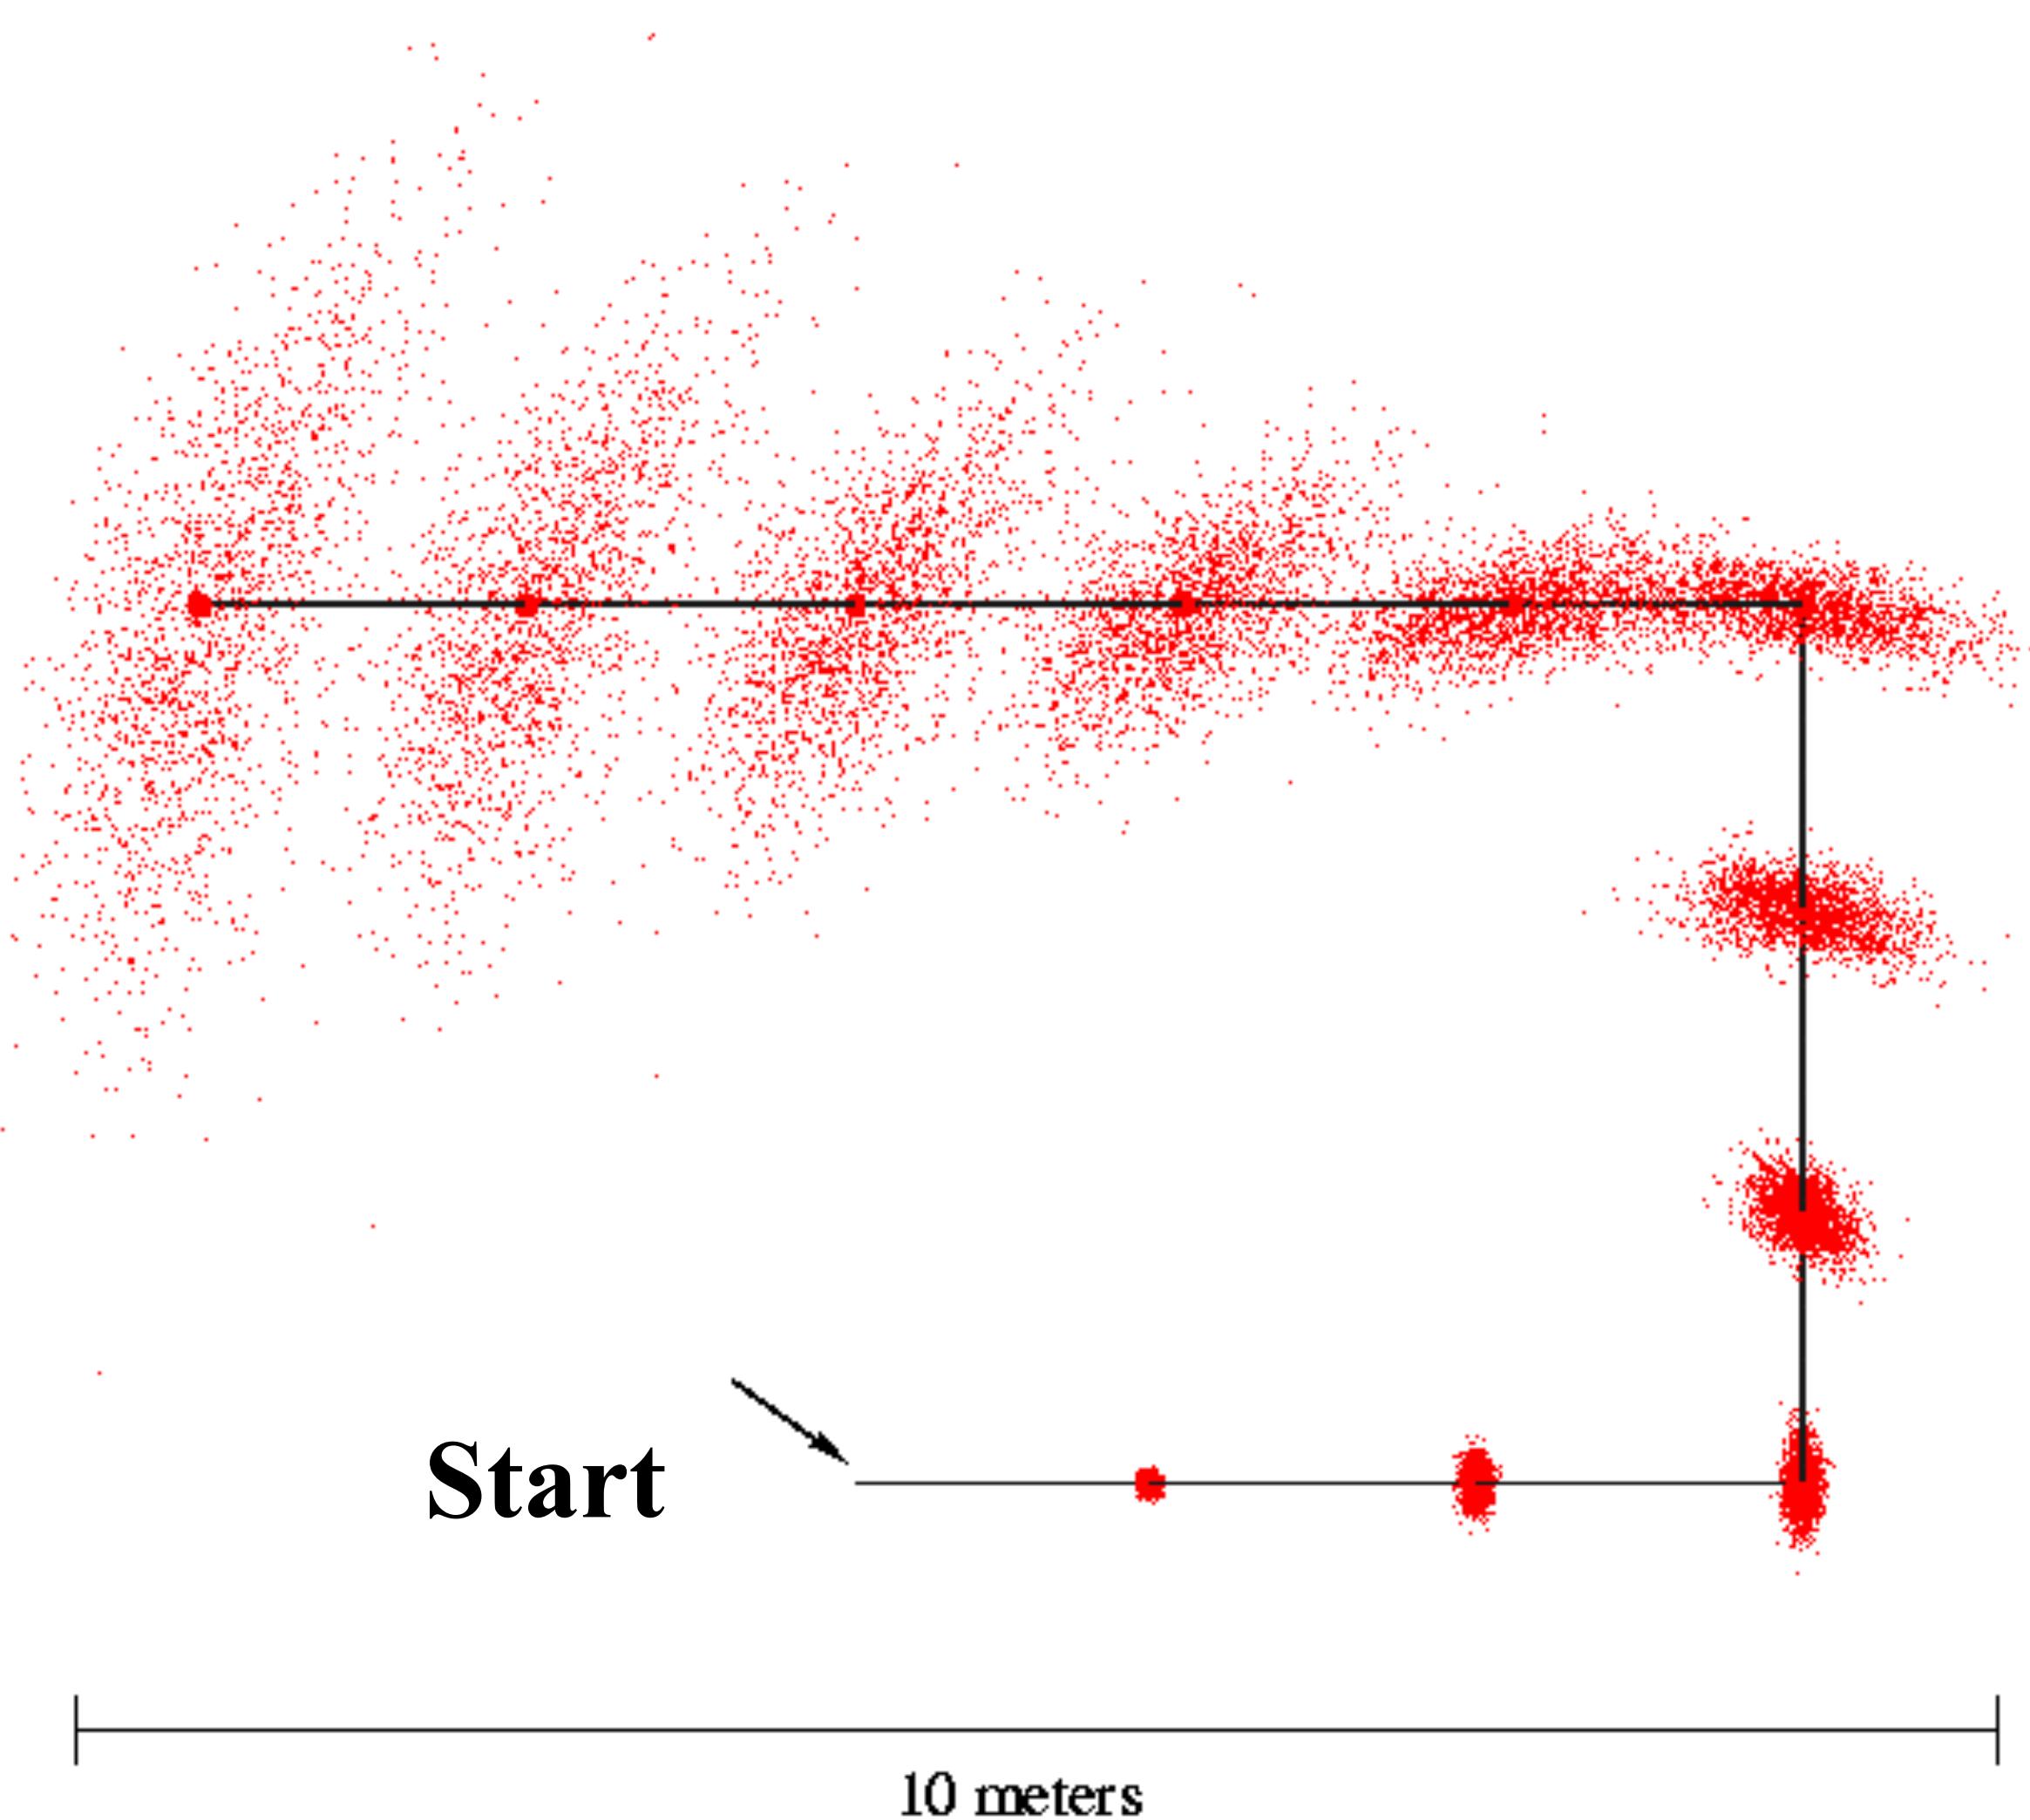


Resampling

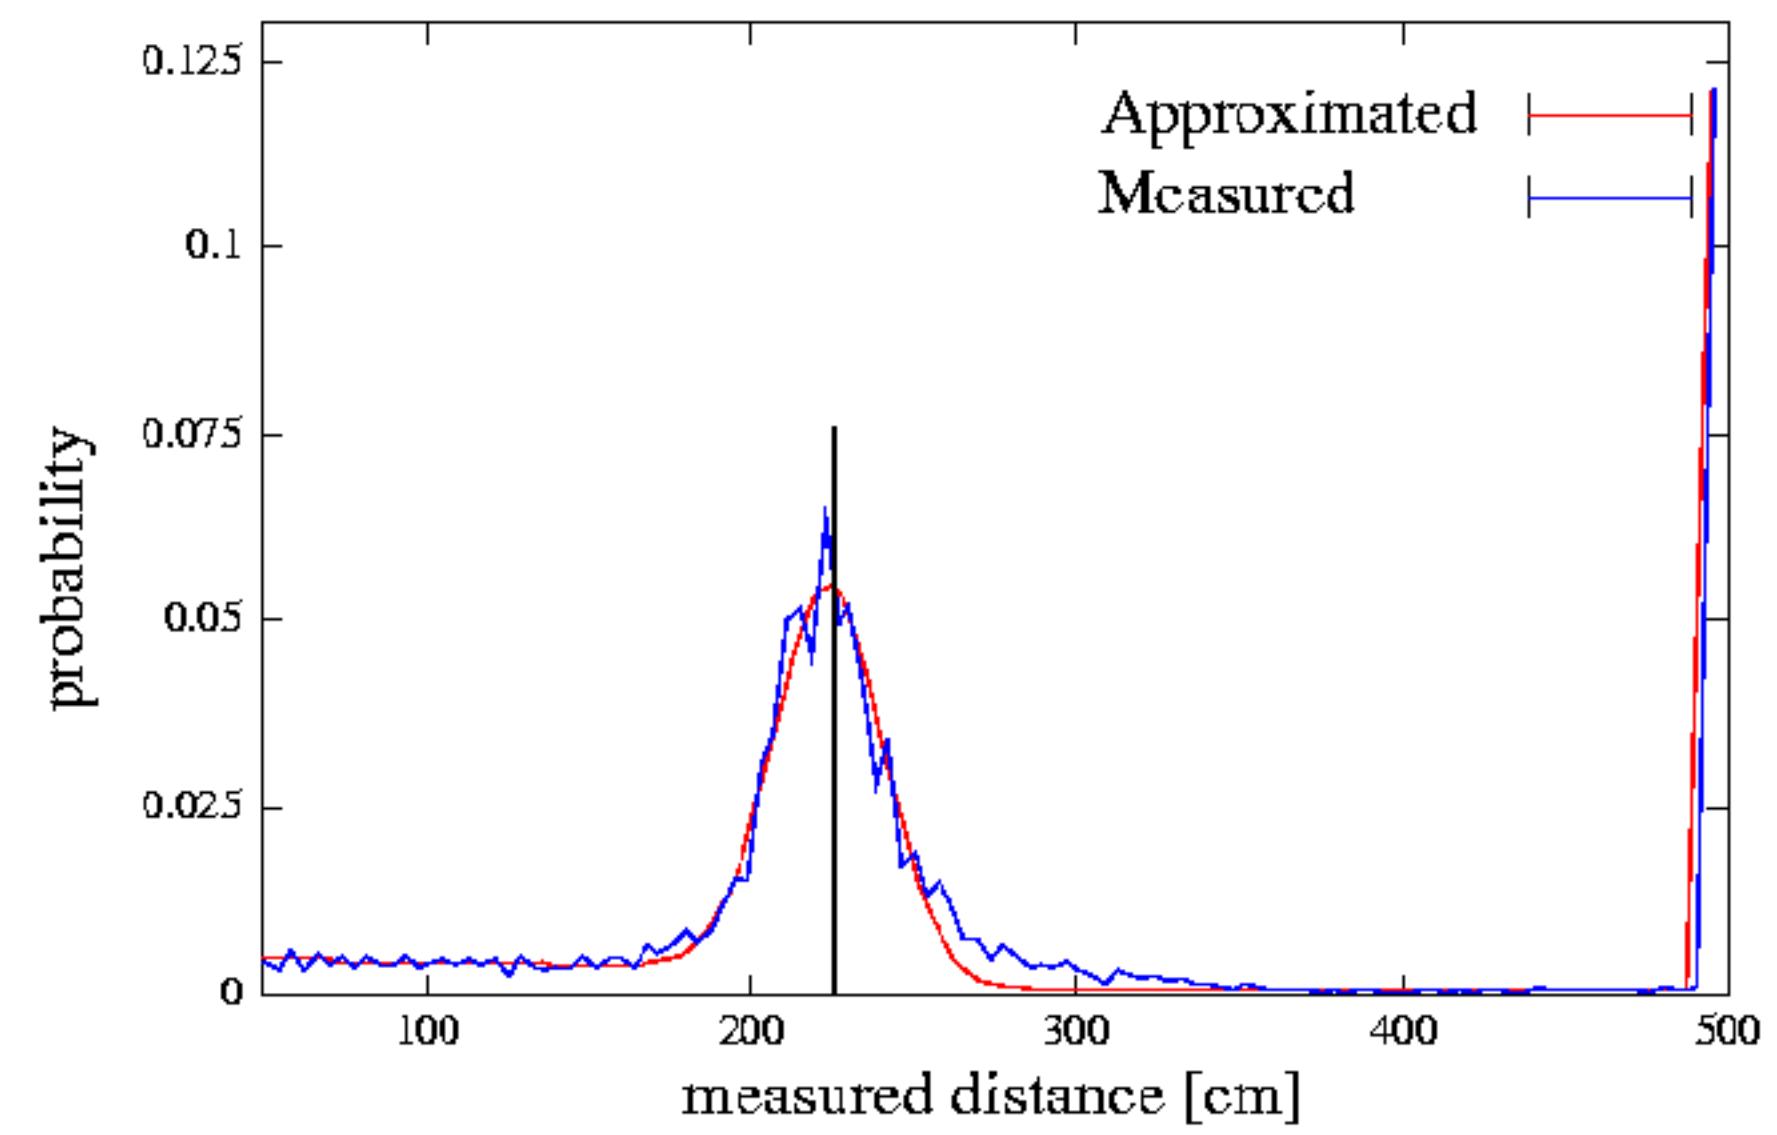


- Roulette wheel
- Binary search, $n \log n$
- Stochastic universal sampling
- Systematic resampling
- Linear time complexity
- Easy to implement, low variance

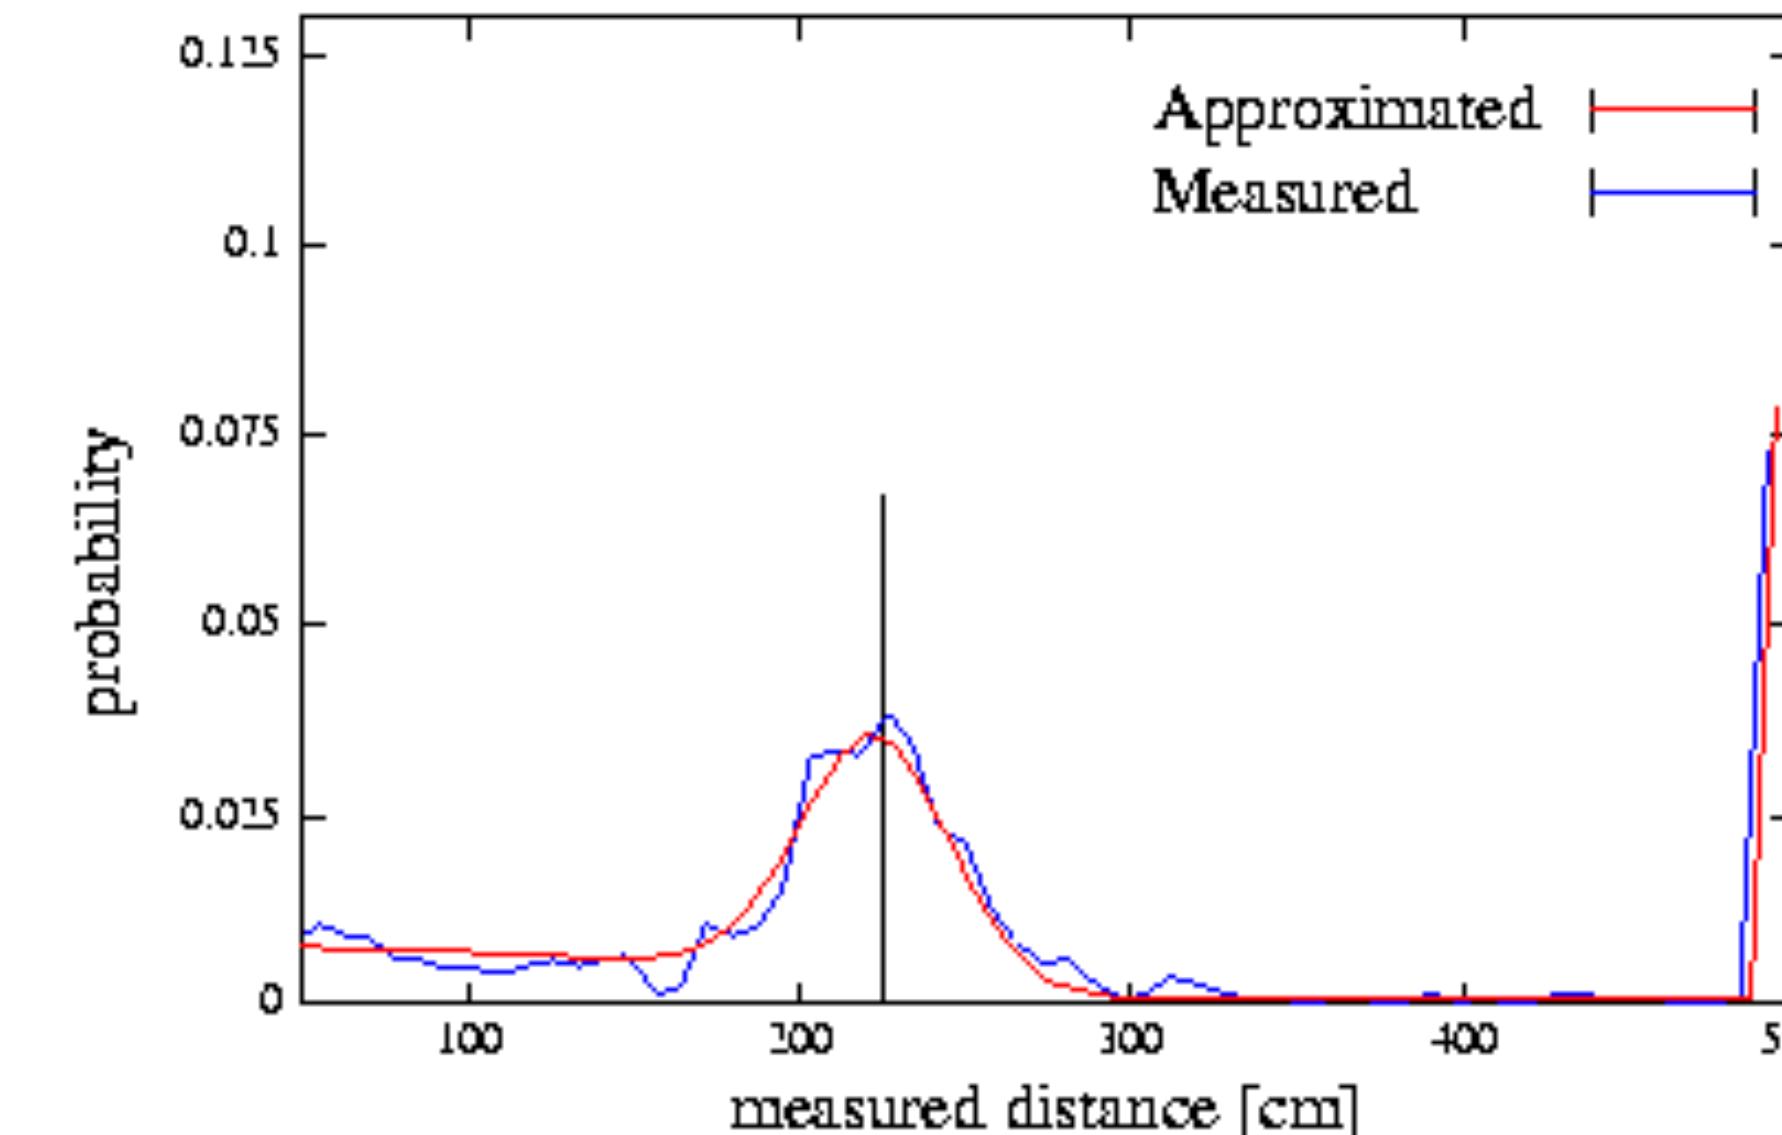
Motion Model Reminder



Proximity Sensor Model Reminder

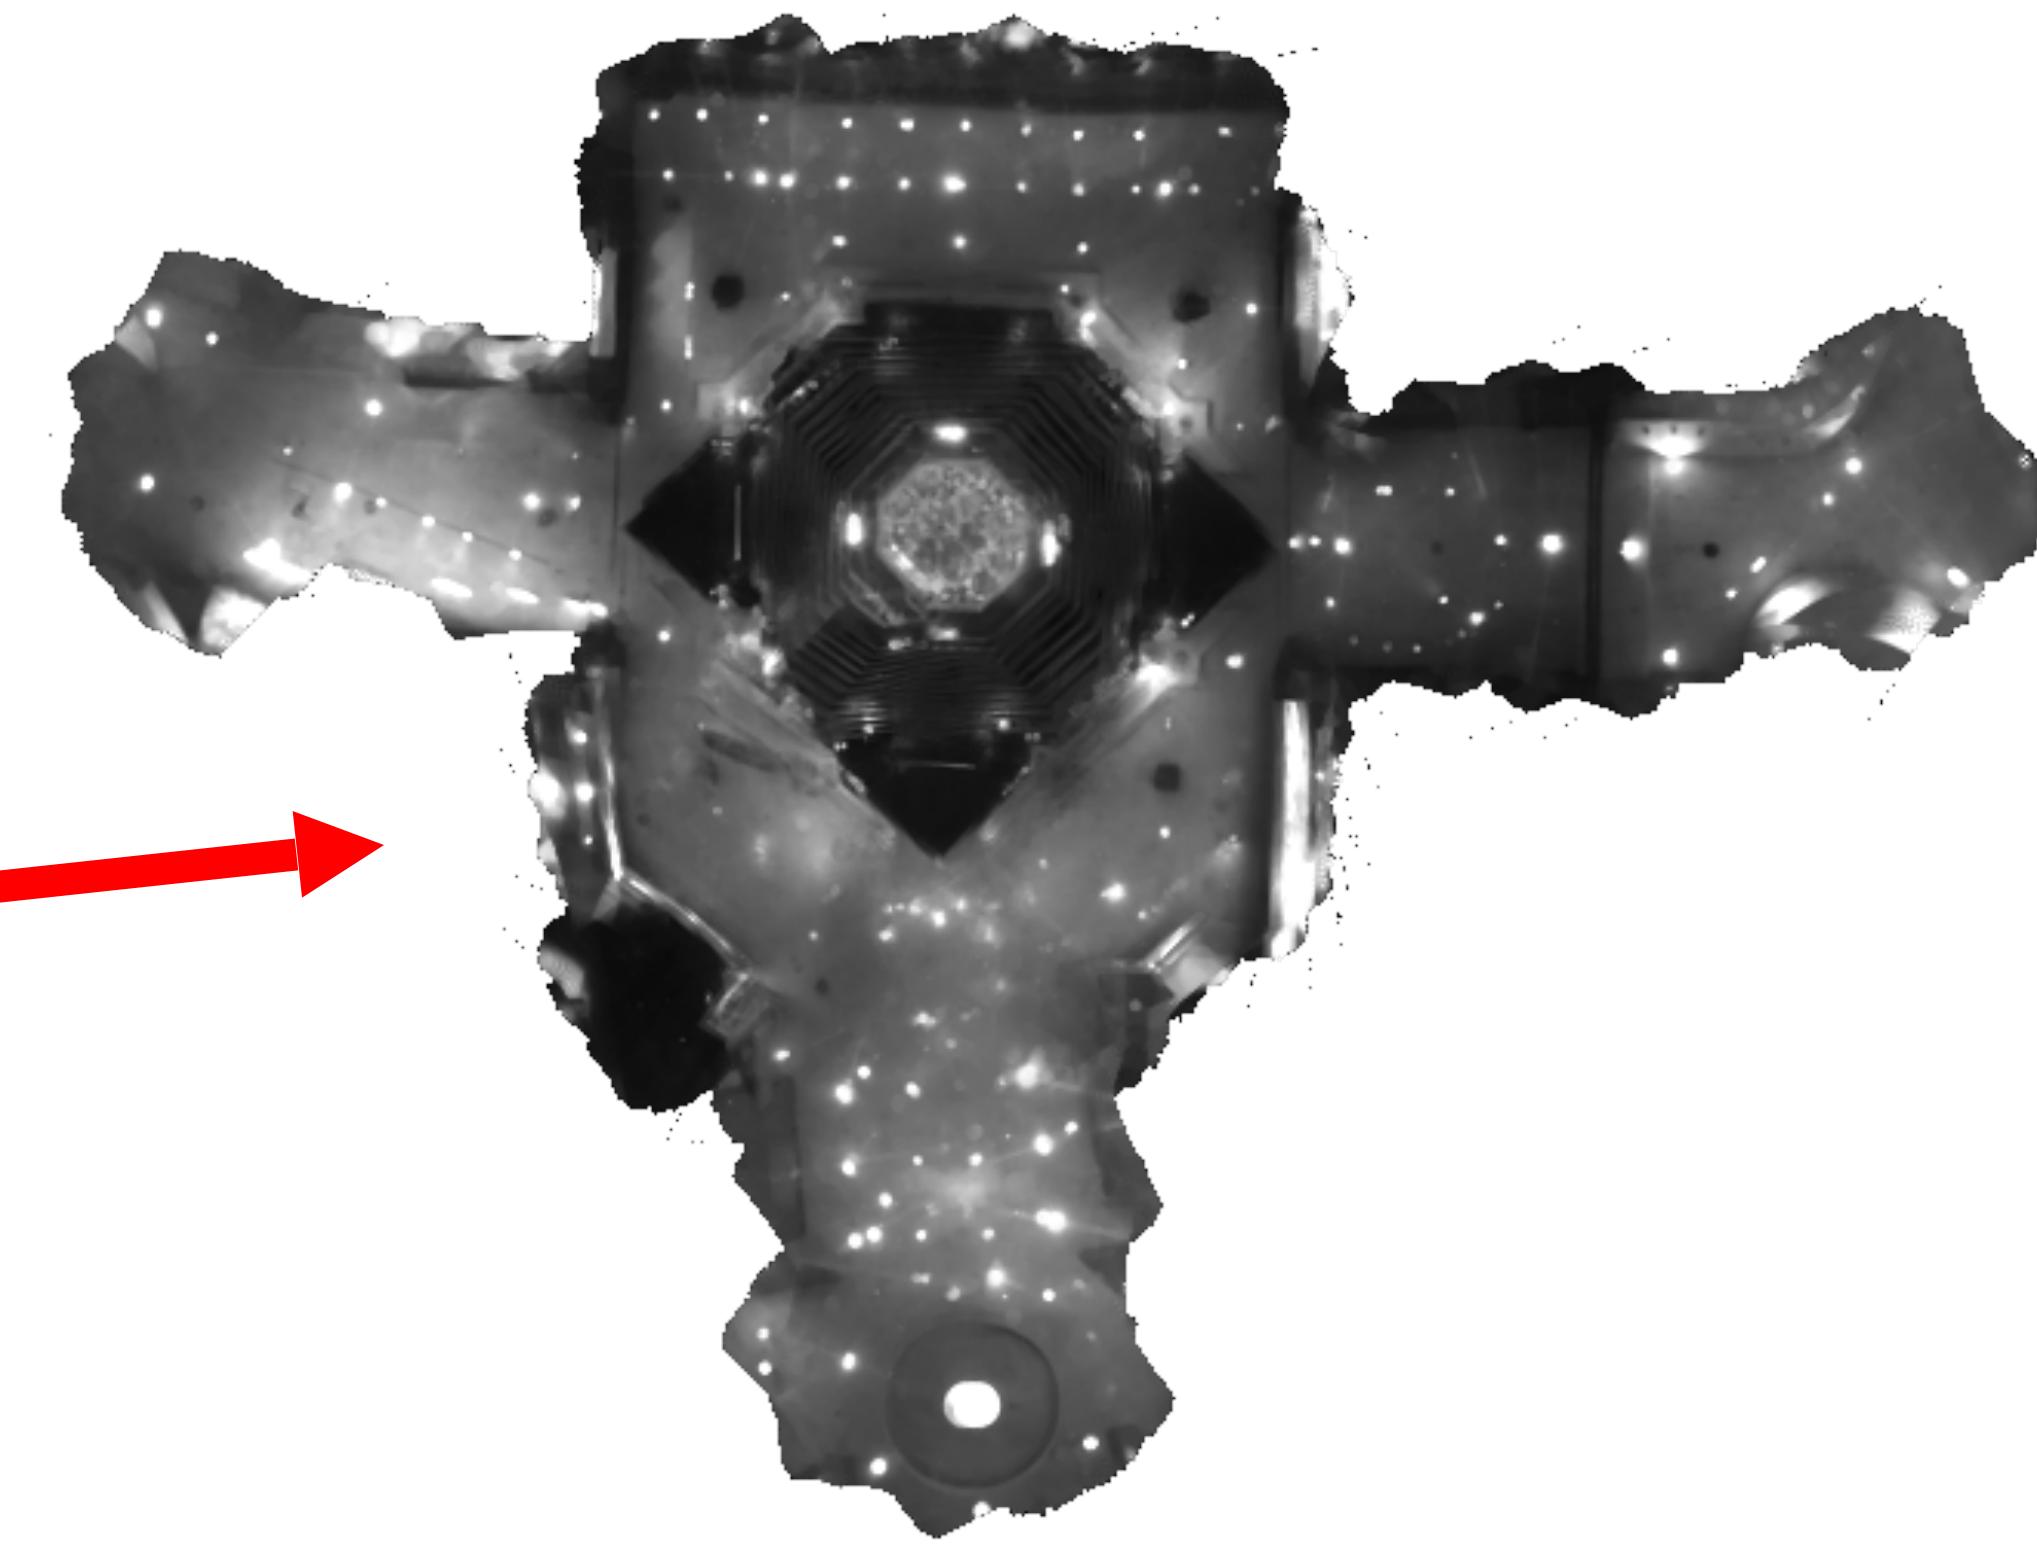


Laser sensor

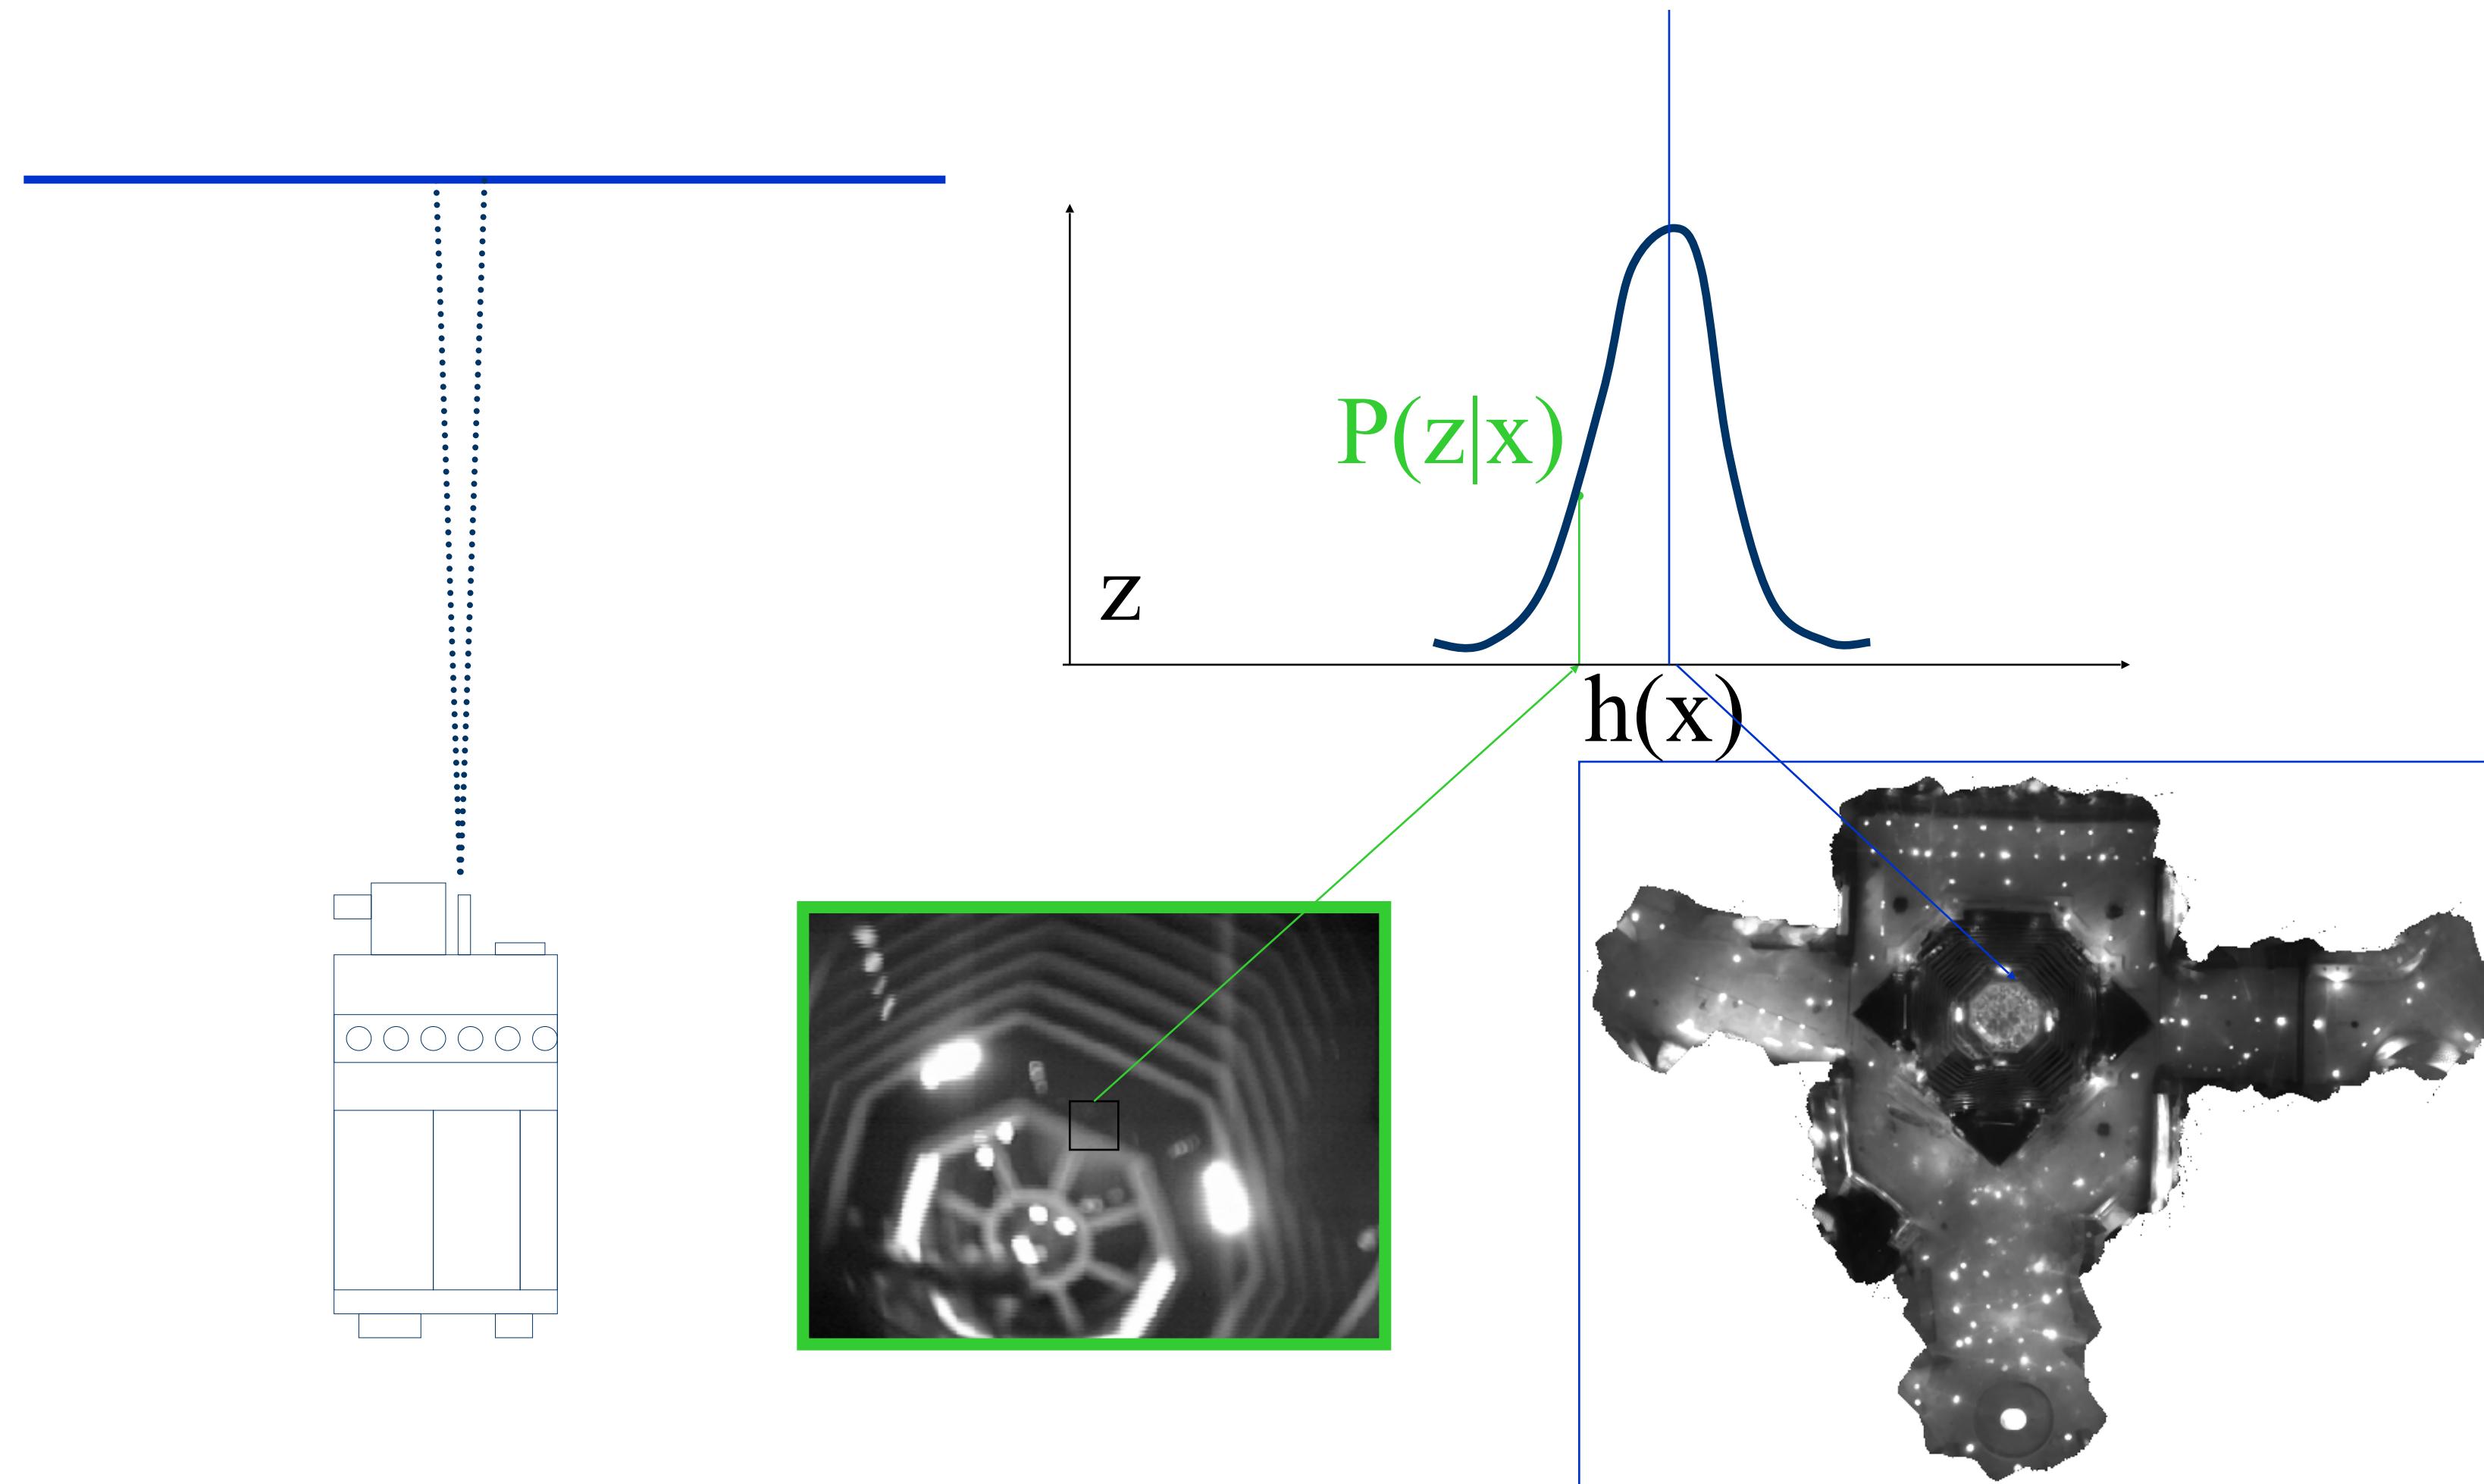


Sonar sensor

Using Ceiling Maps for Localization

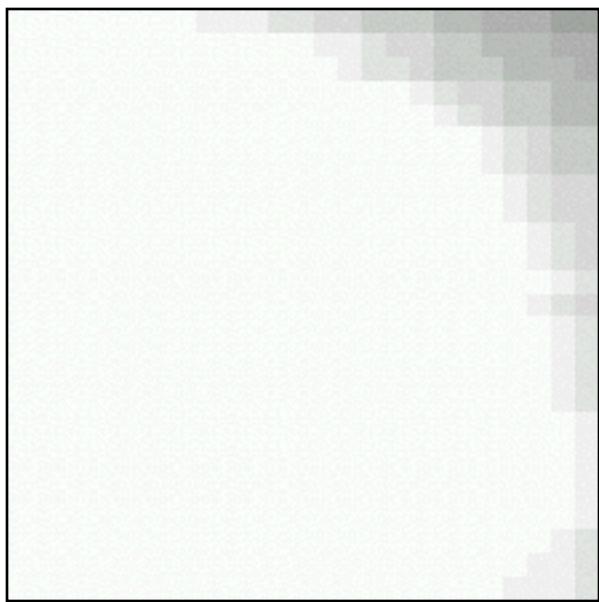


Vision-based Localization

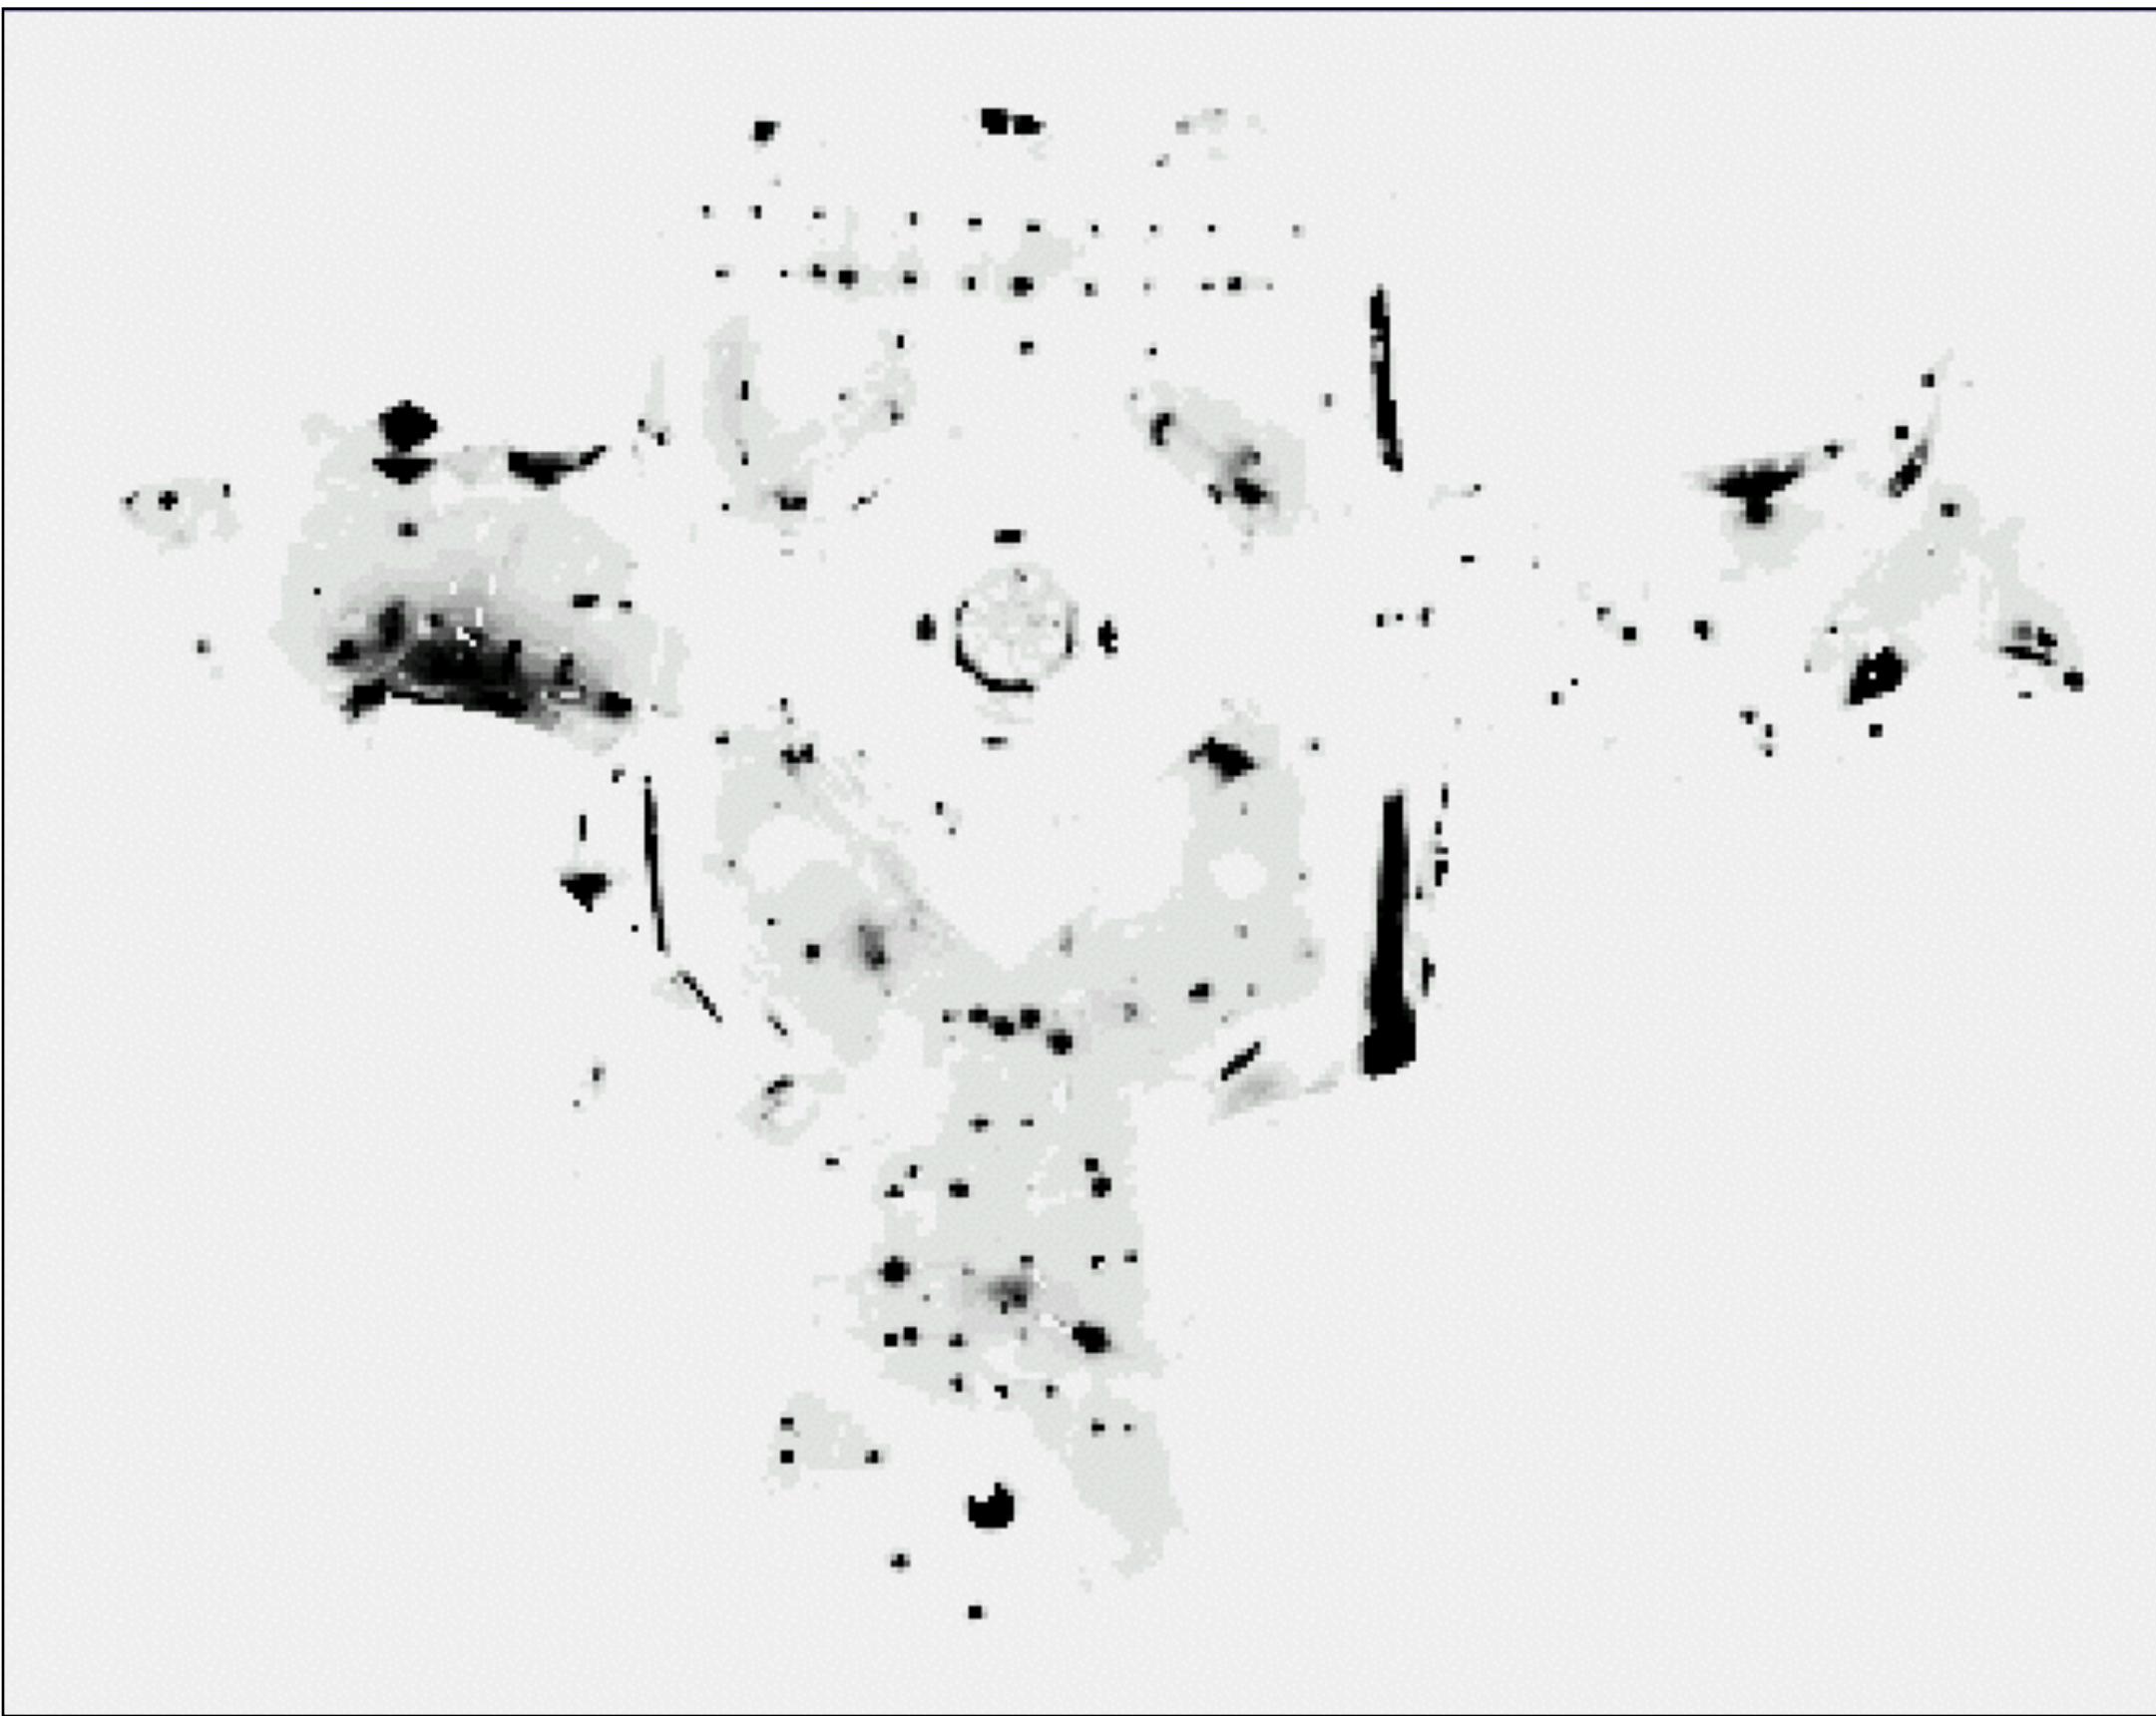


Under a Light

Measurement z:

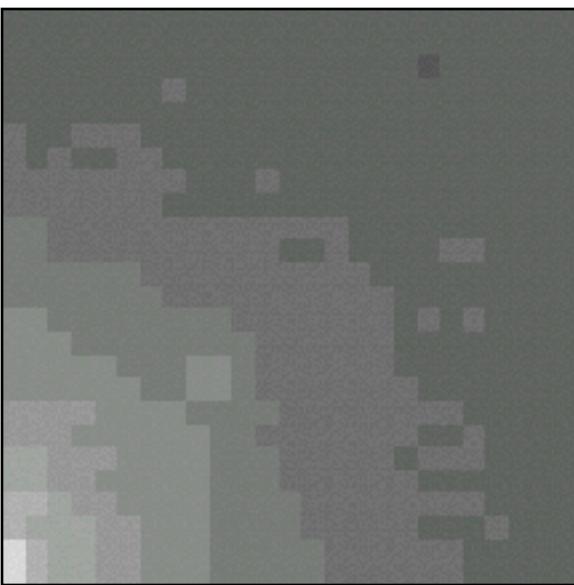


$P(z|x)$:



Next to a Light

Measurement z:



$P(z|x)$:

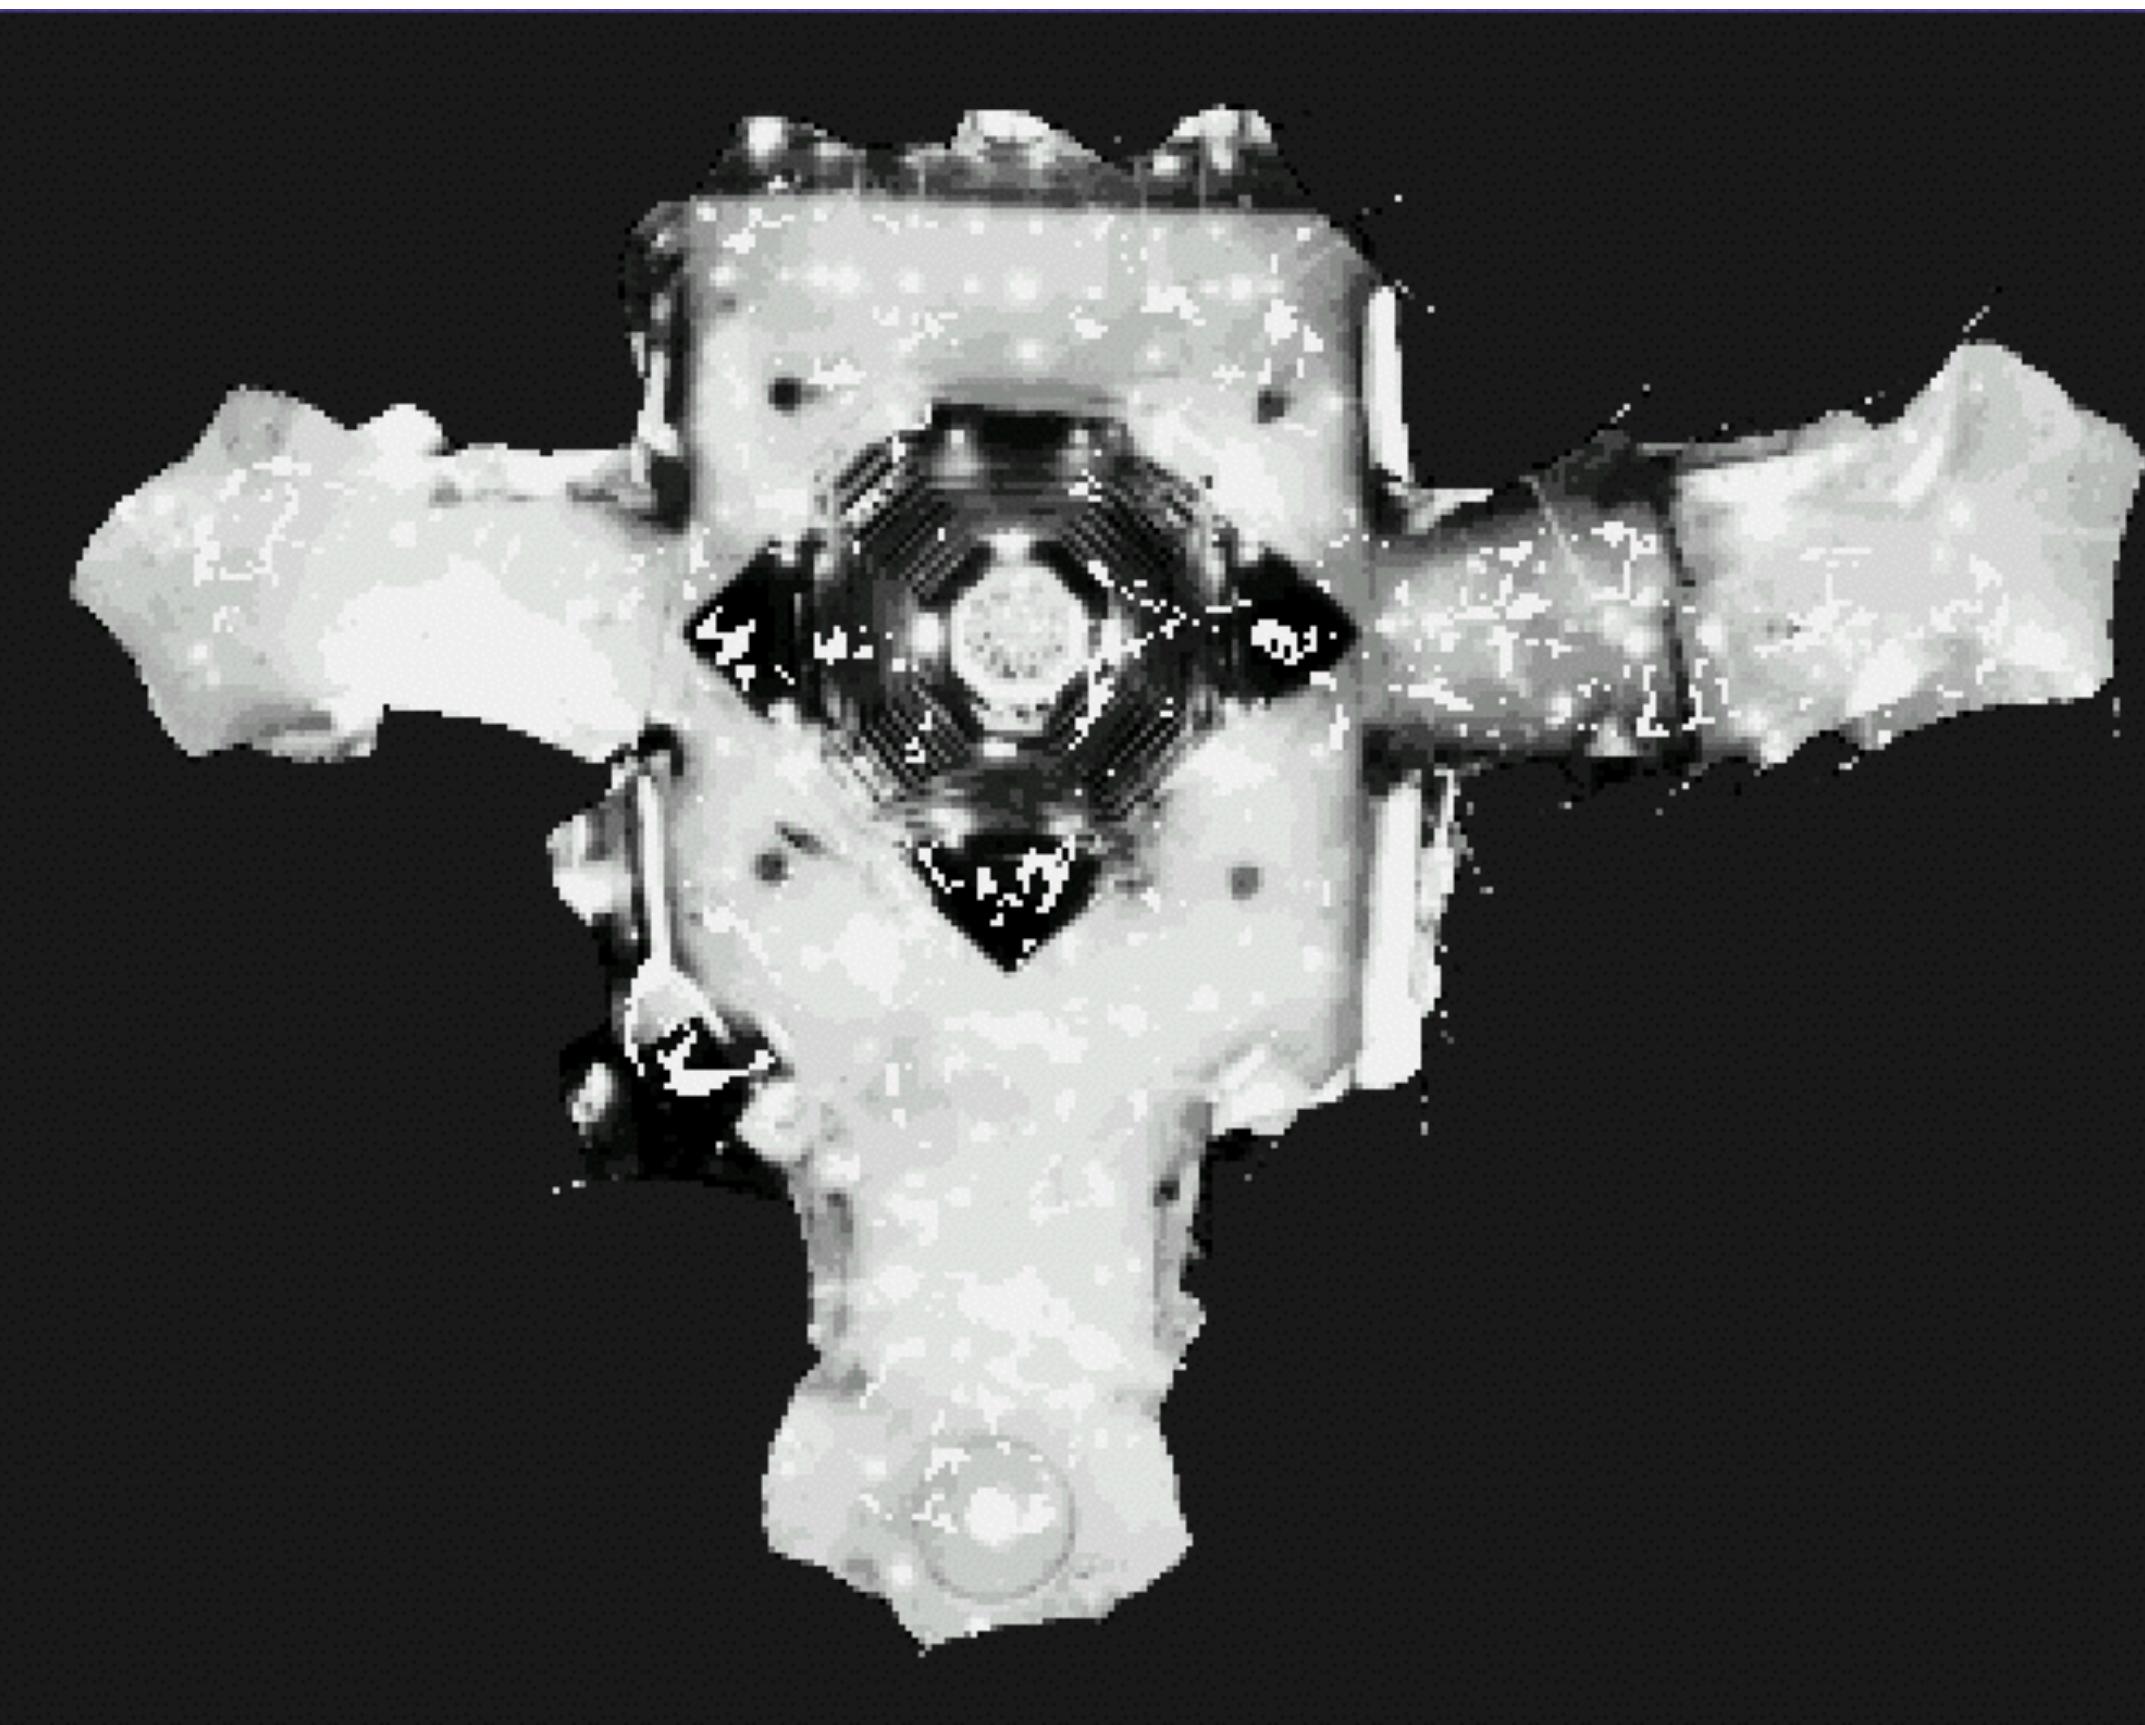


Elsewhere

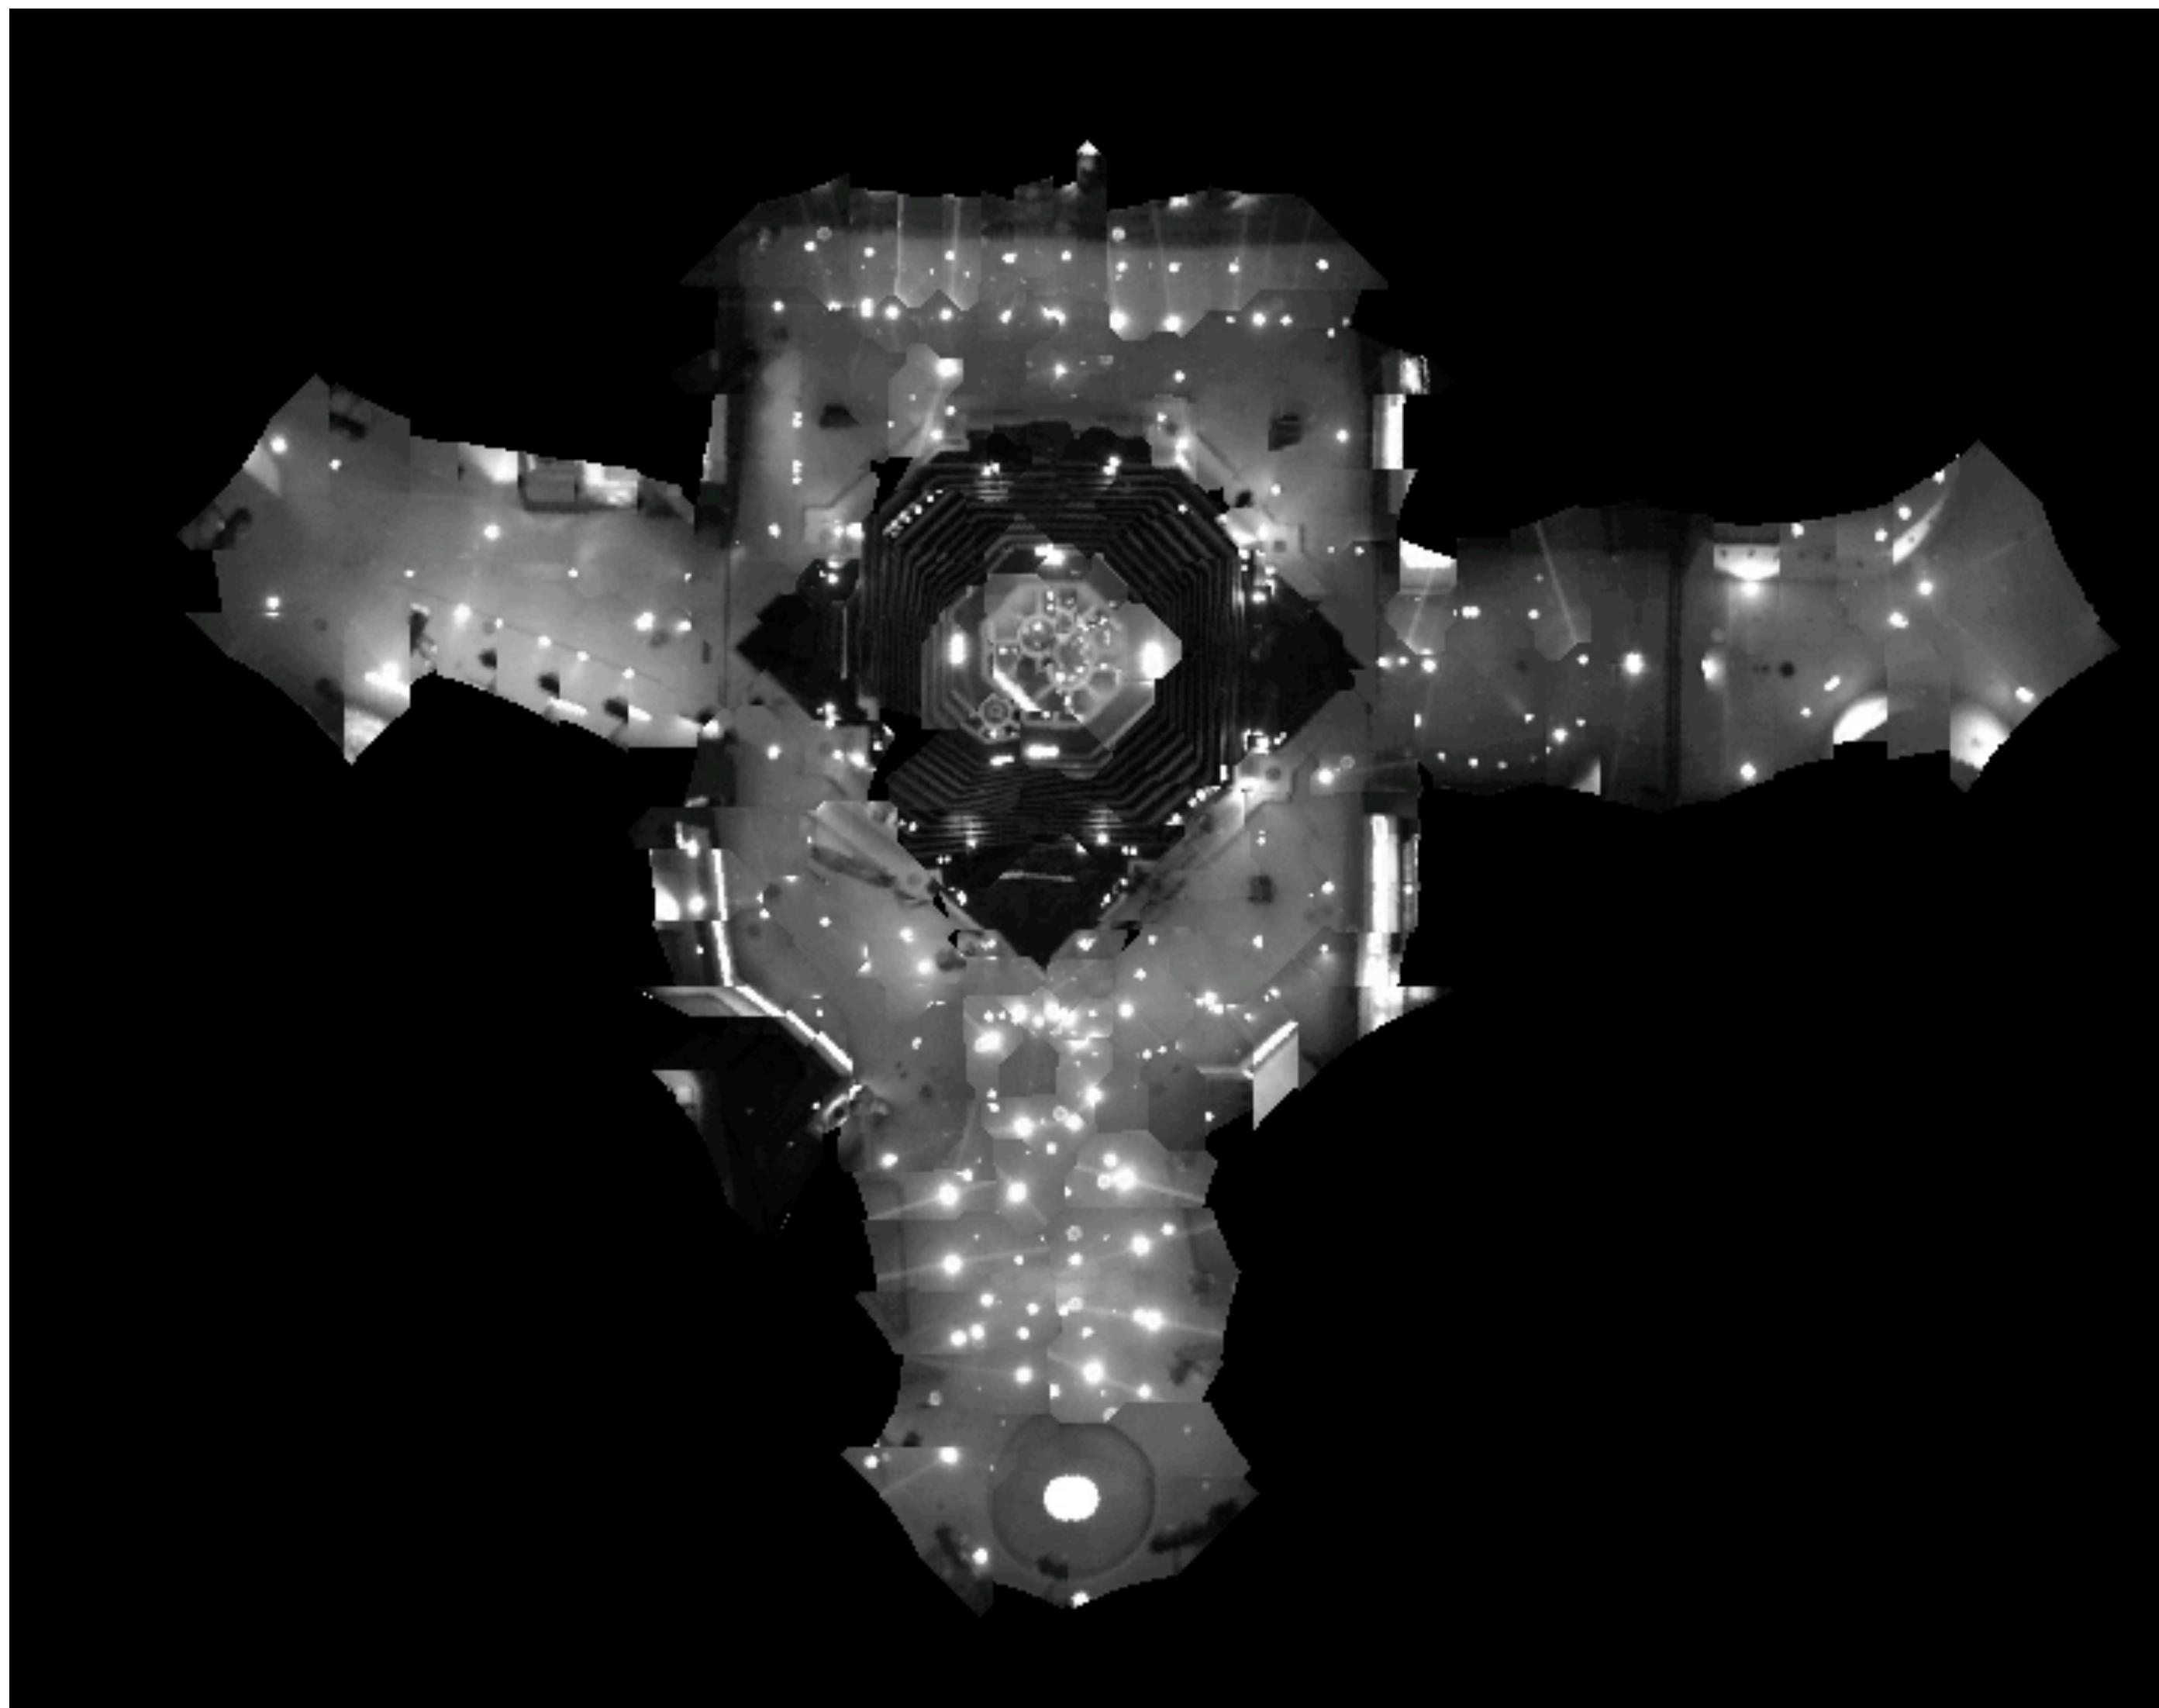
Measurement z:



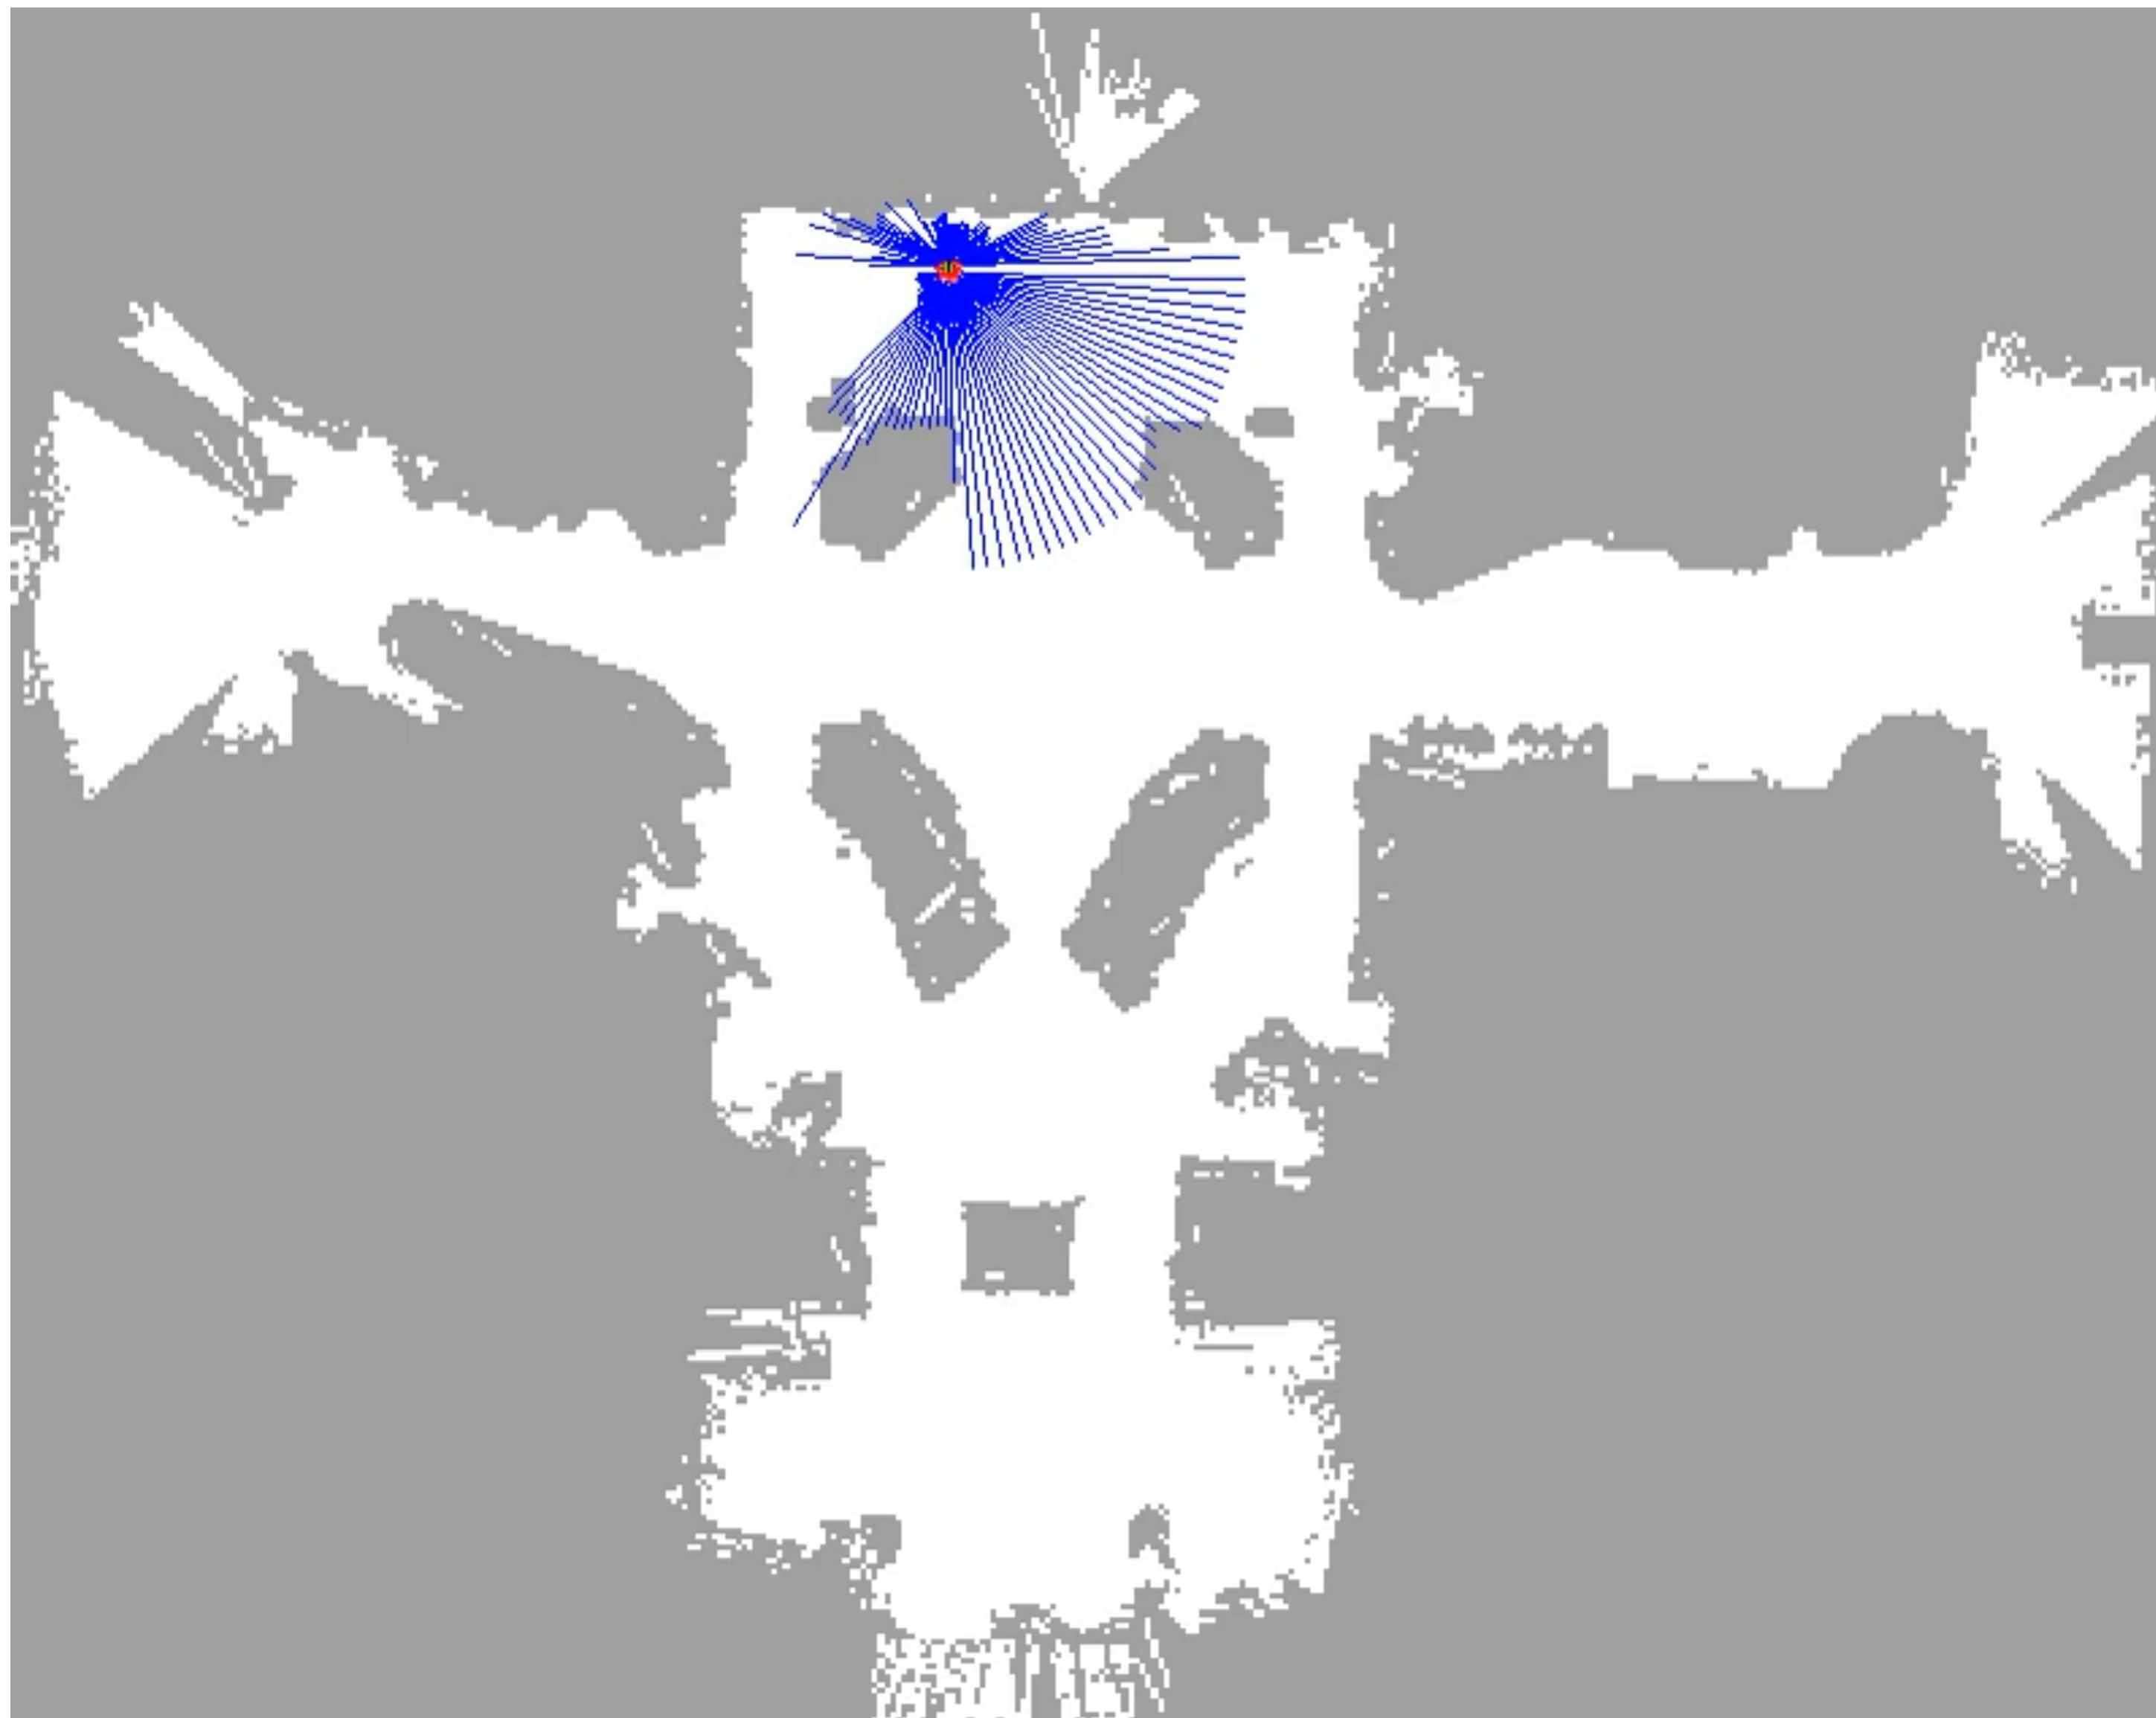
$P(z|x)$:



Global Localization Using Vision



Recovery from Failure



Next Lecture: More PF and Mapping

